

RegressionModelsNotes

Coursera Course by John Hopkins University

INSTRUCTORS: Dr. Brian Caffo, Dr. Roger D. Peng, Dr. Jeff Leek

Contents

Intro	5
GitHub Link for Lectures	5
Course Book	5
Instructor's Note	5
Data Science Specialization Community Site	5
Least Squares and Linear Regression	6
Regression	6
Introduction to Regression	6
Relevant Simply Statistics Post	6
Questions for this Class	7
Introduction to Basic Least Squares	8
Finding the Middle via Least Squares	9
Technical Details	10
Introductory Data Example	10
Comparing Childrens' Heights and Their Parents' Heights	10
Regression Through the Origin	12
Lesson with <code>swirl()</code> : Introduction	14
Linear Least Squares	15
Notation and Background	15
Notation	15
The Empirical Standard Deviation adn Variance	16
Normalization	16
The Empirical Covariance	16
Some Facts About Correlation	16
Linear Least Squares	17
Fitting the Best Line	17
Linear Least Squares Coding Example	17
Adding a Linear Regression to <code>ggplot</code>	18
Technical Details	19
Lesson with <code>swirl()</code> : Least Squares Estimation	19
Regression to the Mean	19
Regression to the Mean	19
Plotting the Regression Implicitly	20
Lesson with <code>swirl()</code> : Residuals	22

Quiz 1	23
Linear Regression & Multivariable Regression	25
Statistical Linear Regression Models	25
Statistical Linear Regression Models	25
Basic Regression Model with Additive Gaussian Errors	25
Interpreting Coefficients	25
Intercept	25
Slope	25
Linear Regression for Prediction	26
Example using diamond Data	26
Changing Units in the Model	28
Estimating a Value	28
Residuals	29
Residuals	29
Properties of the Residuals	29
Residuals, Coding Example	30
Using Residual Plot to Detect a Poorly Fit Model	31
Detecting Heteroskedasticity with a Residual Plot	33
Residual Variance	35
Estimating Residual Variation	35
Diamond Example	35
Summarizing Variation	36
R Squared, the Coefficient of Determination	36
Lesson with <code>swirl()</code> : Residual Variation	37
Inference in Regression	37
Inference in Regression	37
Recall Our Model and Fitted Values	37
Review Some Statistical Inference Concepts	37
Explanation	38
Coding Example	38
Generating Confidence Intervals	39
Prediction	39
Generating Prediction Intervals in Diamond Data Set	39
Lesson with <code>swirl()</code> : Introduction to Multivariable Regression	41
Eliminate Variable Function	41
Lesson with <code>swirl()</code> : MultiVar Examples	41
Reading Multiple Explanatory Variables	42
Quiz 2	44
Multivariable Regression, Residuals, & Diagnostics	46
Multivariable Regression	46
Multivariable Regression Part 1	46
Intro Scenerio Example	46
The Linear Model	46
Multivariable Regression Part 2	46
How to Get Estimates	46
Multivariable Regression Continued	48

Simulation	48
Interpretation of the Coeficients	48
Fitted Values, Residuals and Residual Variation	48
Linear Models Summary	49
Multivariable Regression Tips and Tricks	49
Multivariable Regression Examples Part 1	49
Simulation	50
Back to <code>swiss</code> Data Set	52
Multivariable Regression Examples Part 2	52
Comparing More Than 2 Levels	53
Example in R with InsectSprays	53
Hard Coding the Dummy Variables	55
Summary	56
Multivariable Regression Examples Part 3	57
Fitting the Model	57
Multivariable Regression Examples Part 4	58
Lesson with <code>swirl()</code> : MultiVar Examples2	61
Lesson with <code>swirl()</code> : MultiVar Examples3	61
Adjustment	61
Adjustment Examples	61
Simulation 1	61
Simulation 2	63
Simulation 3	67
Simulation 4	68
Simulation 5	70
Some final thoughts	72
Residuals Again	73
Residuals and Diagnostics Part 1	73
The Linear Model	73
Leverage	74
Residuals and Diagnostics Part 2	74
Residuals and Diagnostics Part 3	75
Simulation 1	75
Simulation 2	77
Using Residuals to See Underlying Patterns	78
Lesson with <code>swirl()</code> : Residuals Diagnostics and Variation	80
Model Selection	80
Model Selection Part 1	80
Multivariable Regression	80
The Rumsfeldian Triplet	81
General Rules	81
Model Selection Part 2	82
Variance Inflation	82
Variance Inflation Factors	83
Variance Inflation with Swiss Data	83
Model Selection Part 3	84
What About Residual Variance Estimation?	84
Nested Model Testing	84

Practice Exercise in Regression Modeling	85
Quiz 3	85
Logistic Regression and Poisson Regression	88
GLMs	88
Intro	88
limitations of Linear Models	88
Example, Linear Models	88
Example, Logistic Regression	89
Example, Poisson Regression	89
Some Things to Note	89
About Variances	89
Odds & Ends	90
Logistic Regression	90
Logistic Regression Part 1	90
Key Ideas	90
Example Baltimore Ravens win/loss	91
Odds	92
Linear vs. Logistic Regression	92
Interpreting Logistic Regression	92
Odds	92
Logistic Regression Part 2	93
Visualizing Fitting Logistic Regression Curves	93
Logistic Regression Part 3	93
Ravens Logistic Regression	93
Odds Ratios & Confidence Intervals	95
ANOVA for Logistic Regression	95
Interpreting Odds Ratios	96
Further Resources	96
Lesson with <code>swirl()</code> : Variance Inflation Factors	96
Lesson with <code>swirl()</code> : Overfitting and Underfitting	97
Lesson with <code>swirl()</code> : Binary Outcomes	97
Poisson Regression	97
Poisson Regression Part 1	97
Overview	97
Poisson Distribution	98
Example: Leek Group Website Traffic	99
Poisson Regression Part 2	102
Linear vs. Poisson Regression	102
Poisson Regression in R	102
More Information	106
Lesson with <code>swirl()</code> : Count Outcomes	106
Hodgepodge	106
Non-Linear Models	106
Fitting Functions Using Linear Models	106
Adding Squared Terms	108
Notes	109
Harmonics Using Linear Models	109

Quiz 4	111
Course Project	116

Intro

This course covers regression analysis, least squares and inference using regression models. Special cases of the regression model, ANOVA and ANCOVA will be covered as well. Analysis of residuals and variability will be investigated. The course will cover modern thinking on model selection and novel uses of regression models including scatterplot smoothing.

GitHub Link for Lectures

[Link to the GitHub for this course](#)

Course Book

Regression Models for Data Science in R, through Leanpub

Further Reading: [Advanced Linear Models for Data Science](#)

Instructor's Note

"We believe that the key word in Data Science is 'science'. Our course track is focused on providing you with three things:

- 1) *An introduction to the key ideas behind working with data in a scientific way that will produce new and reproducible insight*
- 2) *An introduction to the tools that will allow you to execute on a data analytic strategy, from raw data in a database to a completed report with interactive graphics*
- 3) *Giving you plenty of hands on practice so you can learn the techniques for yourself.*

Regression Models represents a both fundamental and foundational component of the series, and it presents the single most practical data analysis toolset. Using only a bare minimum of mathematics, we will attempt to provide you with the fundamentals for the application and practice of regression. We are excited about the opportunity to attempt to scale Data Science education. We intend for the courses to be self-contained, fast-paced, and interactive, and we intend to run them frequently to give people with busy schedules the opportunity to work on material at their own pace.

Brian Caffo and the Data Science Track Team"

Data Science Specialization Community Site

The site is created using GitHub Pages

In addition, Johns Hopkins has a [site on Statistical Methods and Applications for Research in Technology](#) that Dr. Caffo helps manage.

Least Squares and Linear Regression

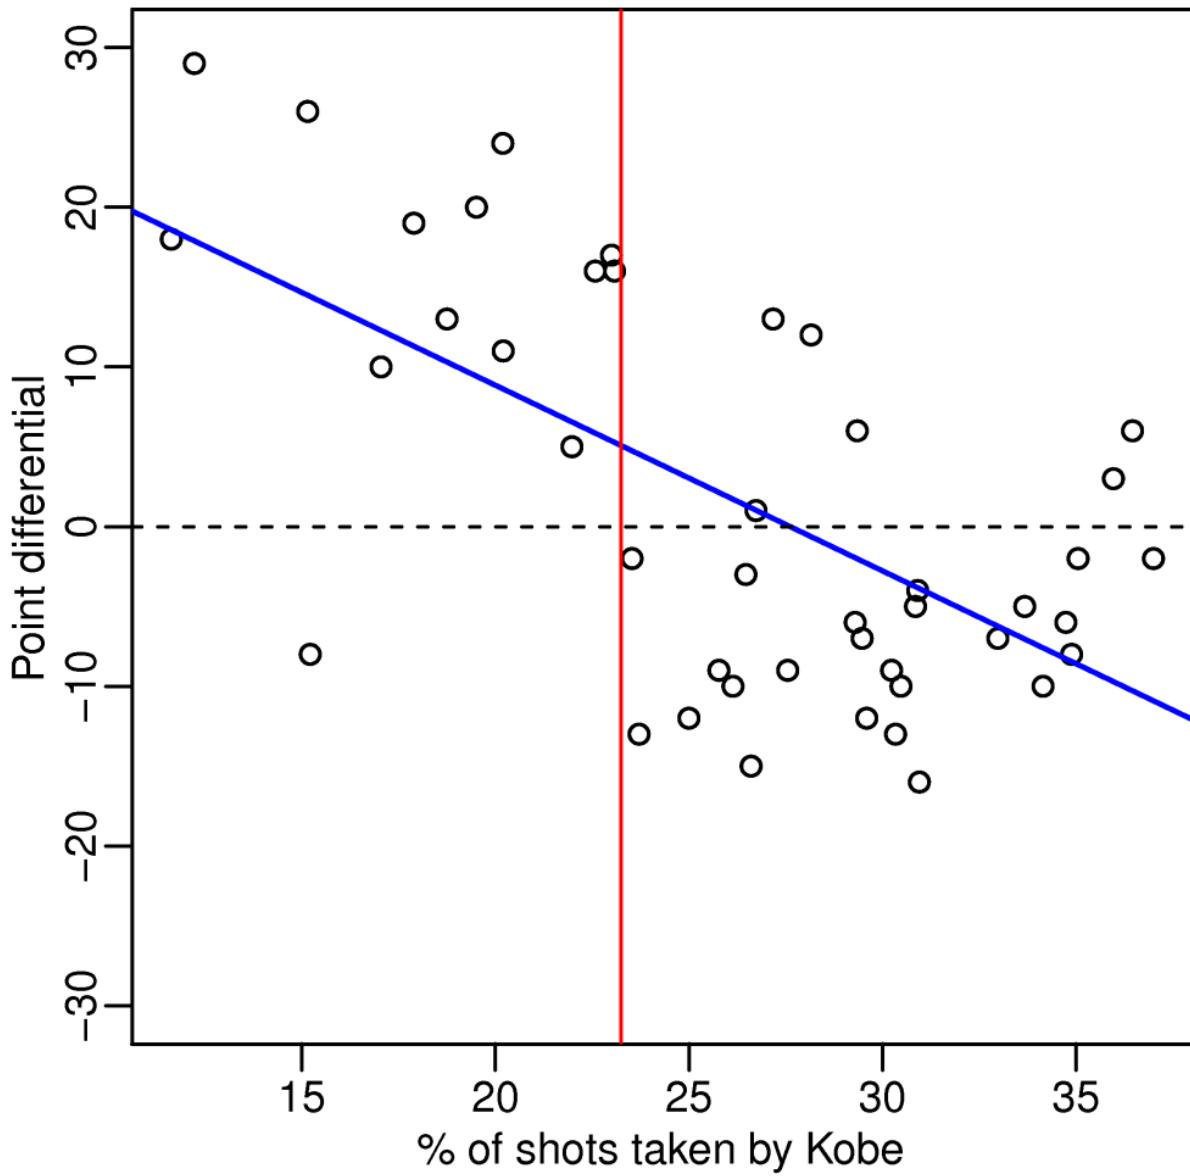
Regression

Introduction to Regression

- The simplicity and intrepretability offered by regression models should make them a first tool of choice for any practical problem.
- First discovered by **Francis Galton** who coined most of the terminology we use today.

Relevant Simply Statistics Post

Simply Statistics is a blog by Jeff Leek, Roger Peng and Rafael Irizarry, who wrote this post



- “Data supports claim that if Kobe stops ball hogging the Lakers will win more”
- “Linear regression suggests that an increase of 1% in percent of shots taken by Kobe results in a drop of 1.16 (+/- 0.22) in score differential.”
+ Standard error given as “+/- 0.22”

Questions for this Class

In reference to Galton’s parent/children height data, which can be accessed from the `galton` dataset in the `UsingR` package.

Consider trying to answer the following kinds of questions:

* To use the parents’ heights to predict childrens’ heights.

- * To try to find a parsimonious (explain the data), easily described mean relationship between parent and children's heights.
- * To investigate the variation in childrens' heights that appears unrelated to parents' heights (residual variation).
- * To quantify what impact genotype information has beyond parental height in explaining child height.
- * To figure out how/whether and what assumptions are needed to generalize findings beyond the data in question.
- * Why do children of very tall parents tend to be tall, but a little shorter than their parents and why children of very short parents tend to be short, but a little taller than their parents? (This is a famous question called "Regression to the mean".)

Introduction to Basic Least Squares

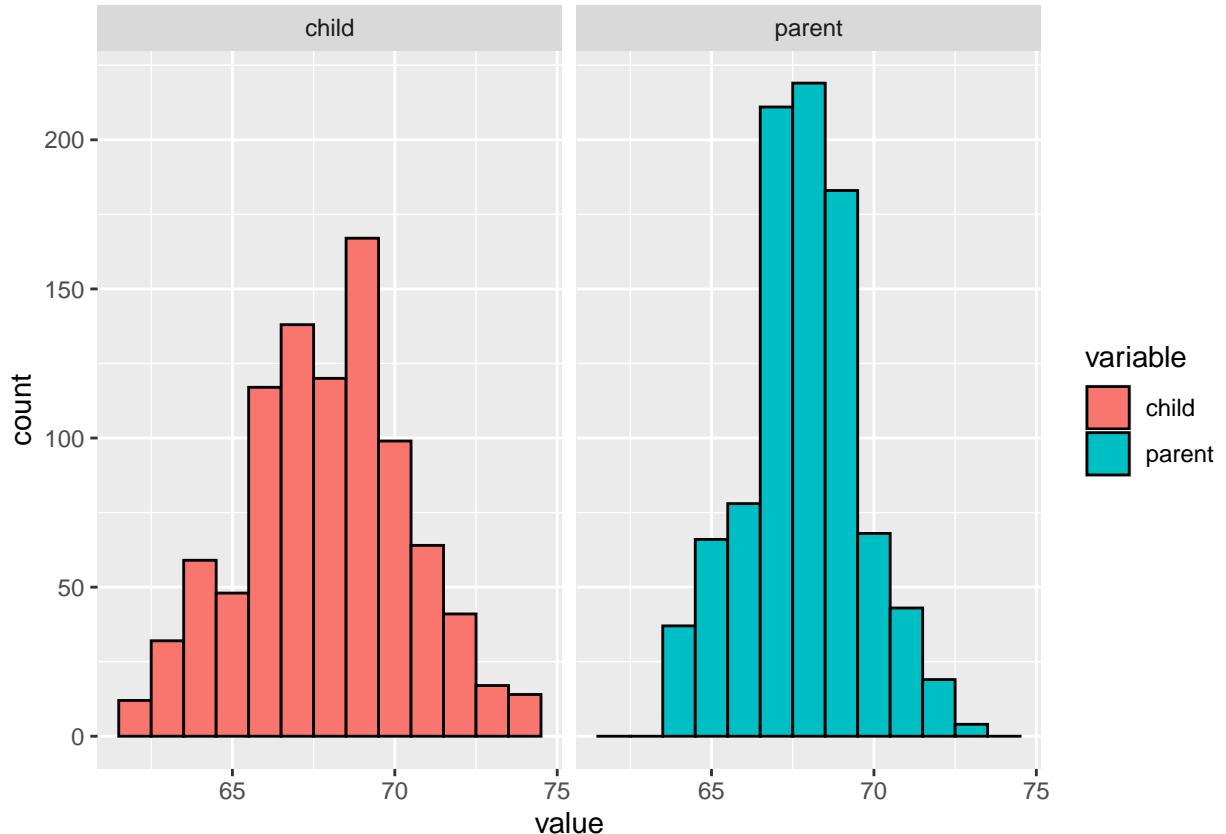
- Let's look at the data first used by Francis Galton in 1885.
- Galton was a statistician who invented the term and concepts of regression and correlation, founded the journal Biometrika, and was the cousin of Charles Darwin.
- Let's look at the marginal (parents disregarding children and children disregarding parents) distributions first.
 - + Parent distribution is all heterosecual couples.
 - + Correction for gender via multiplying female heights by 1.08.
 - + Overplotting is an issue from discretization.

```
library(UsingR); data(galton); library(reshape2); library(tidyverse)

long <- melt(galton)

## No id variables; using all as measure variables

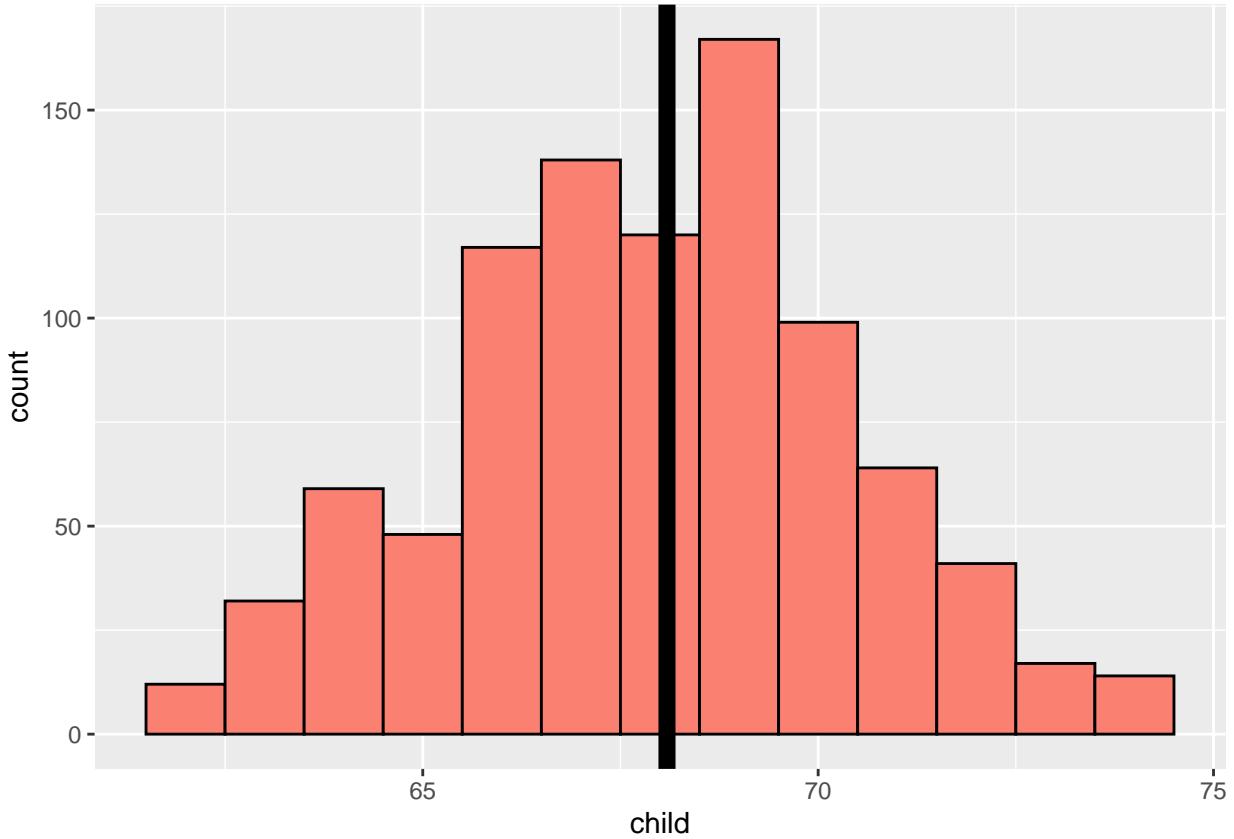
plot <- ggplot(long, aes(x = value, fill = variable)) +
  geom_histogram(colour = "#000000", binwidth = 1)
plot + facet_grid(.~variable)
```



Finding the Middle via Least Squares

- Consider only the children's heights
 - + How could one describe the “middle”?
 - + One definition, let Y_i be the height of child i for $i = 1, \dots, n = 928$, then define the middle as the value of μ that minimizes $\sum_{i=1}^n (Y_i - \mu)^2$
- This is the physical center of mass of the histogram.
- The result of this is that $\mu = \bar{Y}$

```
ggplot(galton, aes(x = child)) +
  geom_histogram(fill = "salmon", colour = "#000000", binwidth = 1) +
  geom_vline(xintercept = mean(galton$child), size = 3)
```



- The above plot of child heights has a mean of 68.0884698

Technical Details

Proof that \bar{Y} is the minimizer for $\sum_{i=1}^n (Y_i - \mu)^2$

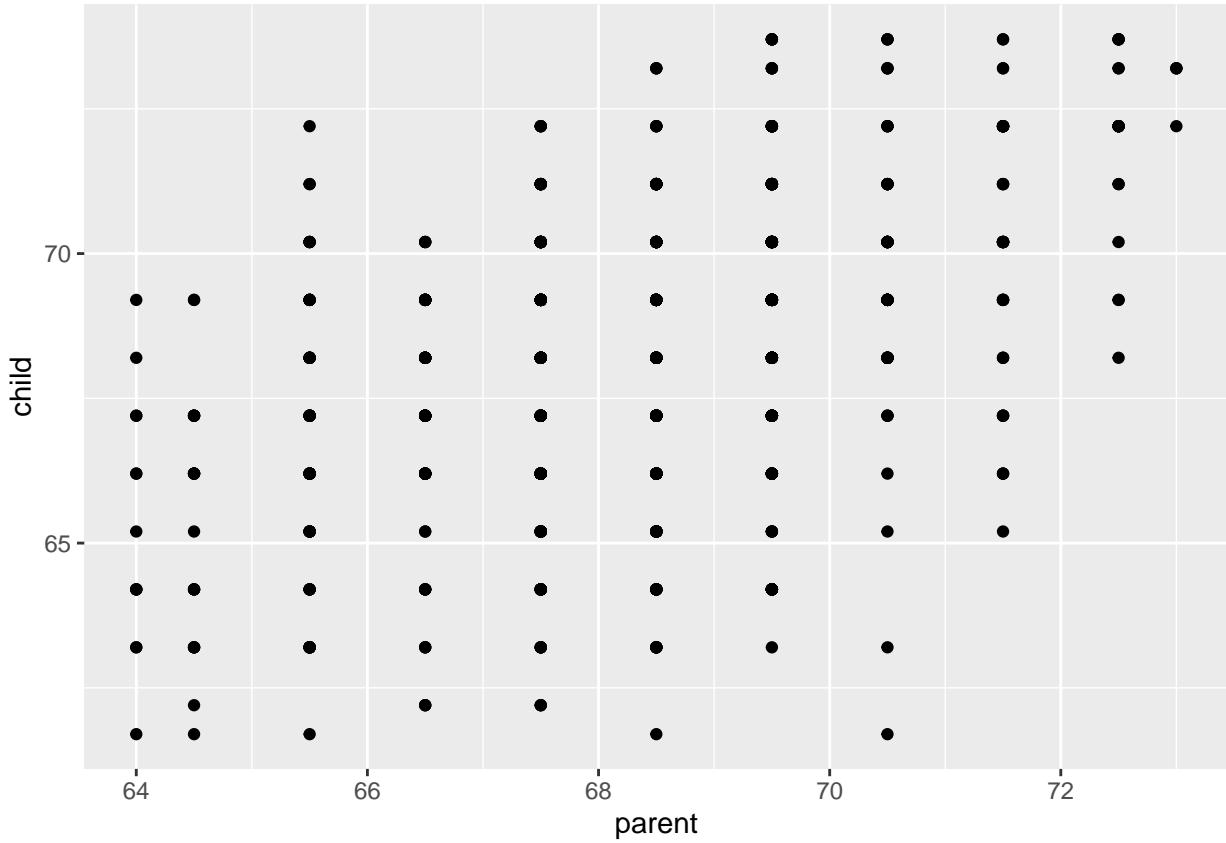
$$\begin{aligned}
 \sum_{i=1}^n (Y_i - \mu)^2 &= \sum_{i=1}^n (Y_i - \bar{Y} + \bar{Y} - \mu)^2 \\
 &= \sum_{i=1}^n (Y_i - \bar{Y})^2 + 2 \sum_{i=1}^n (Y_i - \bar{Y})(\bar{Y} - \mu) + \sum_{i=1}^n (\bar{Y} - \mu)^2 \\
 &= \sum_{i=1}^n (Y_i - \bar{Y})^2 + 2(\bar{Y} - \mu) \sum_{i=1}^n (Y_i - \bar{Y}) + \sum_{i=1}^n (\bar{Y} - \mu)^2 \\
 &= \sum_{i=1}^n (Y_i - \bar{Y})^2 + 2(\bar{Y} - \mu)(\sum_{i=1}^n Y_i - n\bar{Y}) + \sum_{i=1}^n (\bar{Y} - \mu)^2 \\
 &= \sum_{i=1}^n (Y_i - \bar{Y})^2 + 0 + \sum_{i=1}^n (\bar{Y} - \mu)^2 \\
 &\geq \sum_{i=1}^n (Y_i - \bar{Y})^2
 \end{aligned}$$

Therefore, $\sum_{i=1}^n (Y_i - \mu)^2$ is minimized when $\bar{Y} = \mu$

Introductory Data Example

Comparing Childrens' Heights and Their Parents' Heights

```
ggplot(galton, aes(x = parent, y = child)) + geom_point()
```



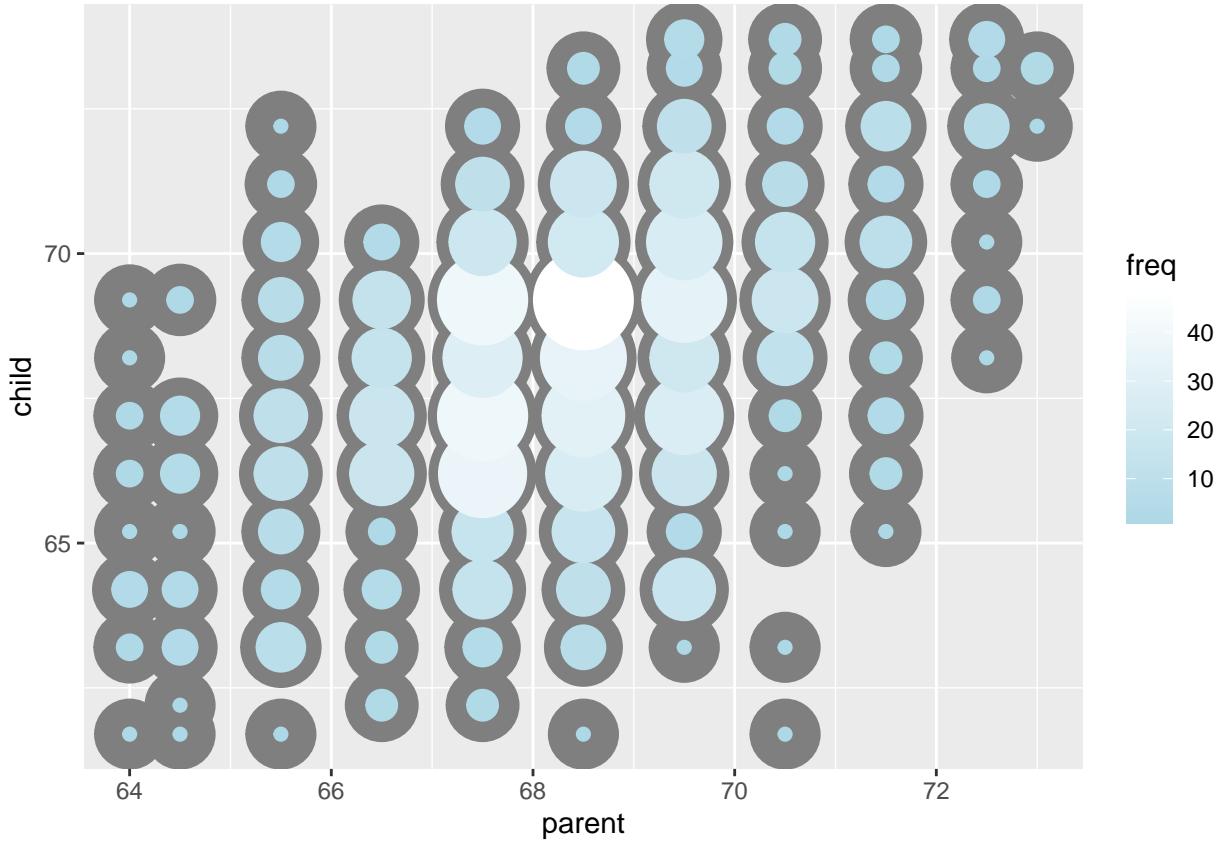
- These points are overplotted, there are multiple overlays at each point, so let's make a better plot

```

freqData <- as.data.frame(table(galton$child, galton$parent))
names(freqData) <- c("child", "parent", "freq")
freqData$child <- as.numeric(as.character(freqData$child))
freqData$parent <- as.numeric(as.character(freqData$parent))
plot <- ggplot(filter(freqData, freq > 0), aes(x = parent, y = child)) +
  scale_size(range = c(2, 20), guide = "none") +
  geom_point(colour = "grey50",
             aes(size = freq + 20, show_guide = FALSE)) +
  geom_point(aes(colour = freq, size = freq)) +
  scale_colour_gradient(low = "lightblue", high = "#FFFFFF")

## Warning: Ignoring unknown aesthetics: show_guide
plot

```



Regression Through the Origin

- Suppose that X_i are the parents' heights
- Consider picking the slope β that minimizes $\sum_{i=1}^n (Y_i - X_i\beta)^2$
- This is exactly using the origin as a pivot point picking the line that minimizes the sum of squared vertical distances of the points to the line
- Subtract the means so that the origin is the mean of the parent and children's heights
+ A plot with a regression line going through true (0,0) often doesn't make sense, so subtracting the means realigns the origin to be in the middle of the data

```

freqData <- as.data.frame(table(galton$parent - mean(galton$parent),
                                galton$child - mean(galton$child)))
names(freqData) <- c("child", "parent", "freq")
freqData$child <- as.numeric(as.character(freqData$child))
freqData$parent <- as.numeric(as.character(freqData$parent))
plot <- ggplot(filter(freqData, freq > 0), aes(x = parent, y = child)) +
  scale_size(range = c(2, 20), guide = "none") +
  geom_point(colour = "grey50",
             aes(size = freq + 20)) +

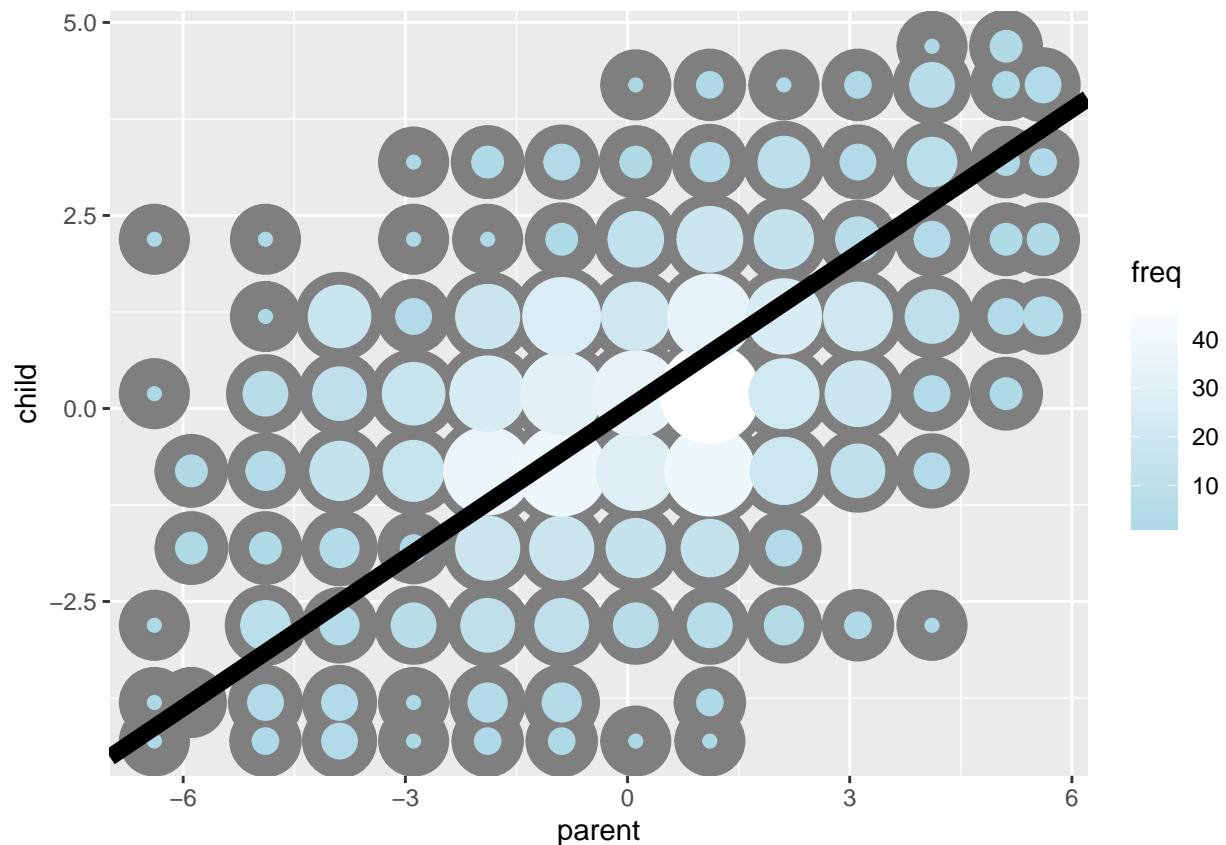
```

```

geom_point(aes(colour = freq, size = freq)) +
  scale_colour_gradient(low = "lightblue", high = "#FFFFFF") +
  geom_abline(intercept = 0,
              slope = lm(
                I(child - mean(child)) ~
                I(parent - mean(parent)) - 1,
                data = galton)$coeff,
              size = 3)

```

plot



- In the next few lectures we'll talk about why this is the solution

```
lm(I(child - mean(child)) ~ I(parent - mean(parent)) - 1, data = galton)
```

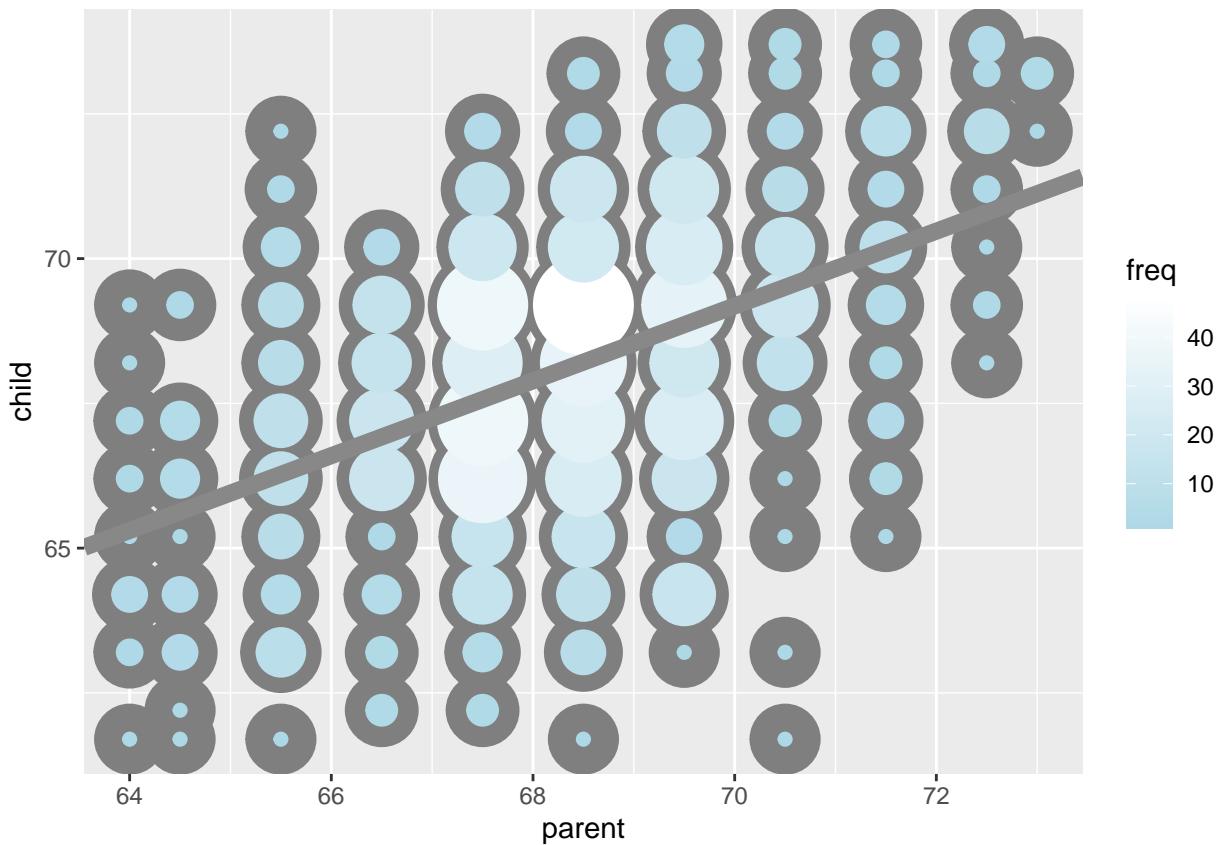
```

## 
## Call:
## lm(formula = I(child - mean(child)) ~ I(parent - mean(parent)) -
##     1, data = galton)
## 
## Coefficients:
## I(parent - mean(parent))
##                 0.6463

```

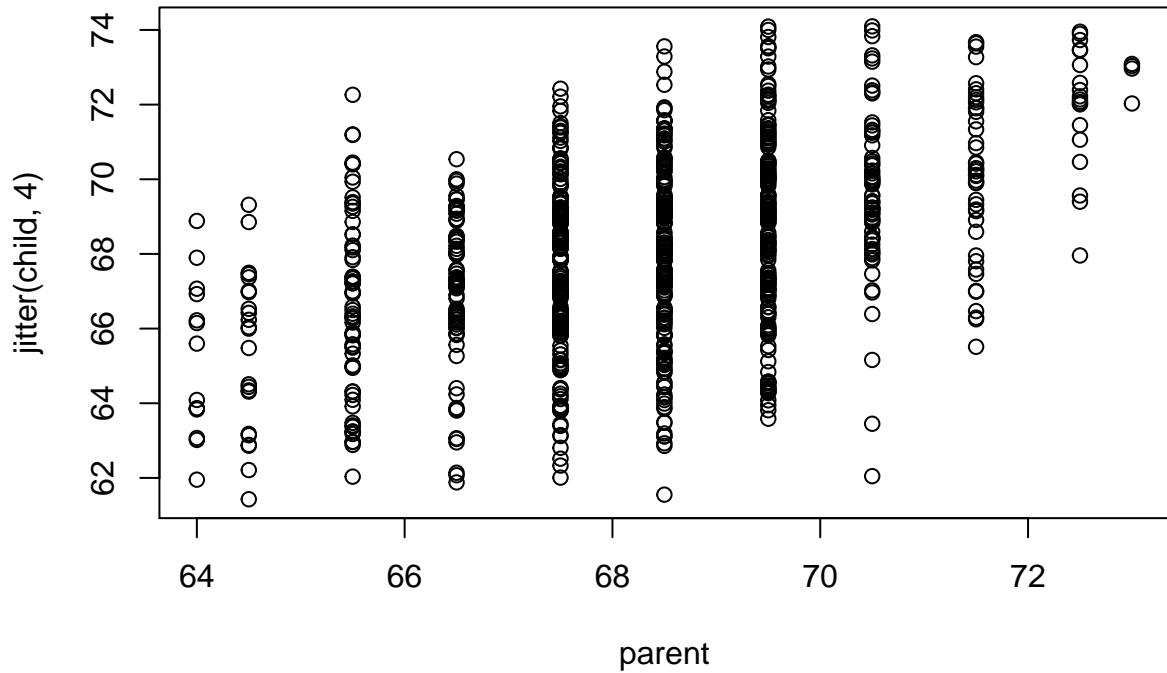
- The `I` function just ignores the intercept, since we already adjusted for that
- We can also fit a line to an un-adjusted model

```
freqData <- as.data.frame(table(galton$child, galton$parent))
names(freqData) <- c("child", "parent", "freq")
freqData$child <- as.numeric(as.character(freqData$child))
freqData$parent <- as.numeric(as.character(freqData$parent))
plot <- ggplot(filter(freqData, freq > 0), aes(x = parent, y = child)) +
  scale_size(range = c(2, 20), guide = "none") +
  geom_point(colour = "grey50", aes(size = freq + 20)) +
  geom_point(aes(colour = freq, size = freq)) +
  scale_colour_gradient(low = "lightblue", high = "#FFFFFF")
lm1 <- lm(galton$child ~ galton$parent)
plot + geom_abline(intercept = coef(lm1)[1], slope = coef(lm1)[2],
                    size = 3, colour = "#888888")
```



Lesson with `swirl()`: Introduction

- Another way we could have gotten past overlapping plot points is to use the `jitter` function
- ```
plot(jitter(child, 4) ~ parent, galton)
```



## Linear Least Squares

- Also called **Ordinary Least Squares (OLS)**; it fits a line through some data.

## Notation and Background

### Notation

- The empirical mean is defined as  
 $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$
- If we subtract the mean from data points, we get data that has a mean of 0. That is, if we define:  
 $\tilde{X}_i = X_i - \bar{X}$ .  
+ The mean of  $\tilde{X}_i$  is 0
- This process is called “**centering**” the random variables
- Recall from the previous lecture that the mean is the least squares solution for minimizing  
 $\sum_{i=1}^n (X_i - \mu)^2$

## The Empirical Standard Deviation adn Variance

- Define the empirical variance as  

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2 = \frac{1}{n-1} (\sum_{i=1}^n X_i^2 - n\bar{X}^2)$$
- The empirical standard deviation is defined as  $S = \sqrt{S^2}$ .  
 + Notice that the standard deviation has the same units as the data.
- The data defined by  $\frac{X_i}{s}$  have an empirical standard deviation of 1. + This is called “**scaling**” the data.

## Normalization

- The data defined by  

$$Z_i = \frac{X_i - \bar{X}}{s}$$
  
 have an empirical mean of 0 and an empirical standard deviation of 1.
- The process of centering then scaling the data is called “**normalizing**” the data.
- Normalized data are centered at 0 and have units equal to standard deviations of the original data.
- For example, a value of 2 from normalized data is saying that data point was two standard deviations larger than the mean.

## The Empirical Covariance

- Consider now when we have pairs of data,  $(X_i, Y_i)$
- Their empirical covariance is  

$$\begin{aligned} Cov(X, Y) &= \frac{1}{n-1} \sum_{i=1}^n n(X_i - \bar{X})(Y_i - \bar{Y}) \\ &= \frac{1}{n-1} (\sum_{i=1}^n X_i Y_i - n\bar{X}\bar{Y}) \end{aligned}$$
- The correlation is defined as  

$$Cor(X, Y) = \frac{Cov(X, Y)}{S_x S_y}$$
  
 + Where  $S_x$  and  $S_y$  are the estimates of standard deviaitons for the X observations and Y observations, respectively.

## Some Facts About Correlation

- $Cor(X, Y) = Cor(Y, X)$
- $-1 \leq Cor(X, Y) \leq 1$
- $Cor(X, Y) = 1$  and  $Cor(X, Y) = -1$  only when the  $X$  or  $Y$  observations fall perfectly on a positive or negative sloped line, repectively.

- $Cor(X, Y)$  measures the strength of the linear relationship between the  $X$  and  $Y$  data, with stronger relationships as  $Cor(X, Y)$  heads towards either -1 or 1 {
- $Cor(X, Y) = 0$  implies no linear relationship

## Linear Least Squares

### Fitting the Best Line

- Let  $Y_i$  be the  $i^{th}$  child's height and  $X_i$  be the  $i^{th}$  (average over the pair of) parents' heights.
- Consider finding the best line  
+ Child's Height =  $\beta_0 + \text{Parent's Height} * \beta_1$   
 $\sum_{i=1}^n Y_i - (\beta_0 + \beta_1 X_i)^2$
- the least squares model fit to the line  $Y = \beta_0 + \beta_1 X$  through the data pairs  $(X_i, Y_i)$  with  $Y_i$  as the outcome obtains the line  $Y = \hat{\beta}_0 + \hat{\beta}_1 X$  where  
 $\hat{\beta}_1 = Cor(Y, X) \frac{Sd(Y)}{Sd(X)}$   
 $\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$
- $\hat{\beta}_1$  has the units of  $Y/X$ ,  $\hat{\beta}_0$  has the units of  $Y$ .
- The line passes through the point  $(\bar{X}, \bar{Y})$
- The slope of the regression line with  $X$  as the outcome and  $Y$  as the predictor is  $\frac{Cor(Y, X) Sd(X)}{Sd(Y)}$
- The slope si the same one you would get if you centered the data,  $(X_i - \bar{X}, Y_i - \bar{Y})$ , and made a regression through the orgin
- If you normalized the data,  $(\frac{X_i - \bar{X}}{Sd(X)}, \frac{Y_i - \bar{Y}}{Sd(Y)})$ , the slope is  $Cor(Y, X)$

### Linear Least Squares Coding Example

```

y <- galton$child
x <- galton$parent
beta1 <- cor(y,x) * sd(y) / sd(x)
beta0 <- mean(y) - beta1 * mean(x)

#Showing the computations by hand are the same as coef from lm function
rbind(c(beta0, beta1), coef(lm(y~x)))

(Intercept) x
[1,] 23.94153 0.6462906
[2,] 23.94153 0.6462906

```

- `lm` stands for *linear model*

```
#The slope is the same in centered data
yc <- y - mean(y)
xc <- x - mean(x)
beta1 <- sum(yc * xc) / sum(xc^2)
c(beta1, coef(lm(y ~ x))[2])

x
0.6462906 0.6462906

lm(yc ~ xc - 1)$coef #minus 1 gets rid of intercept

xc
0.6462906

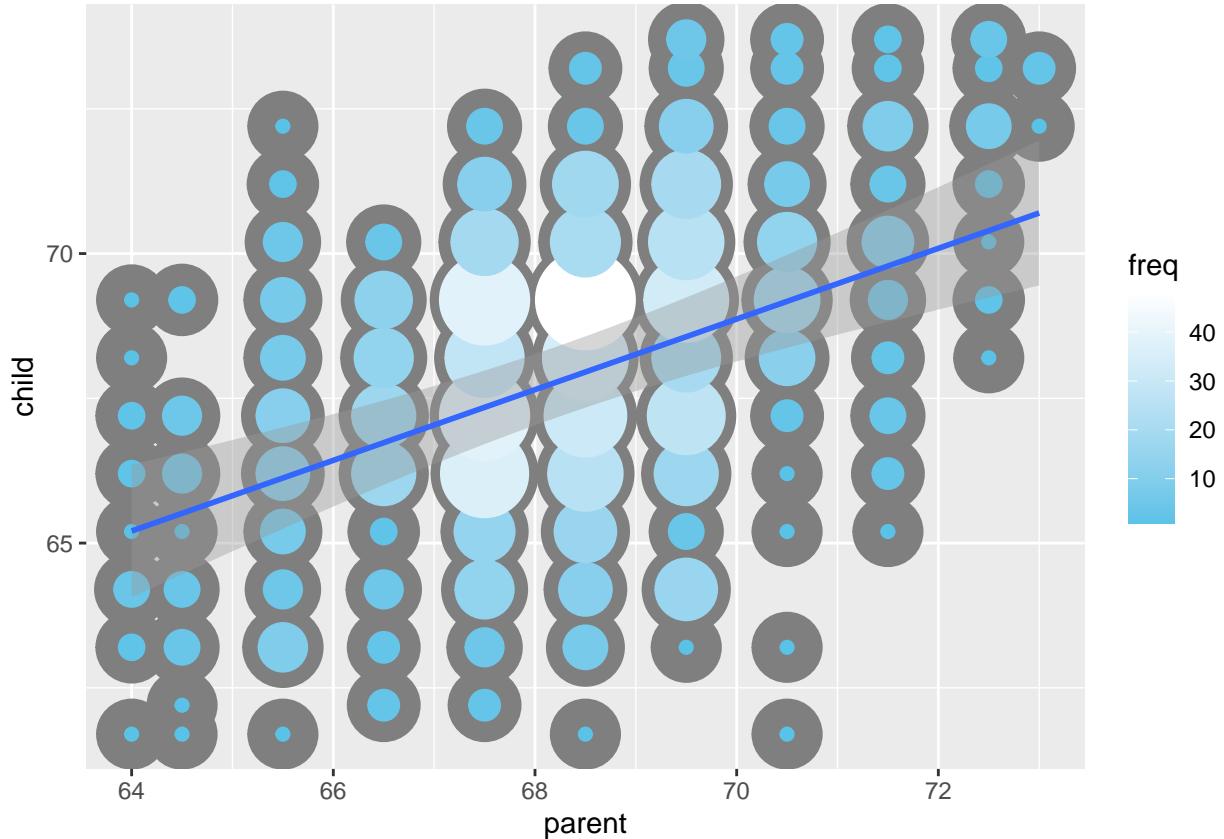
#Normalizing variables results in the slope being the correlation
yn <- (y - mean(y))/sd(y)
xn <- (x - mean(x))/sd(x)
results <- cbind(cor(y,x), lm(yn ~ xn)$coef[2], cor(yn, xn))
colnames(results) <- c("cor(y,x)", "Slope(yn ~ xn)", "cor(yn, xn)")
results

cor(y,x) Slope(yn ~ xn) cor(yn, xn)
xn 0.4587624 0.4587624 0.4587624
```

## Adding a Linear Regression to ggplot

```
plot <- ggplot(filter(freqData, freq > 0), aes(parent, child)) +
 scale_size(range = c(2, 20), guide = "none") +
 geom_point(colour = "grey50", aes(size = freq + 20)) +
 geom_point(aes(colour = freq, size = freq)) +
 scale_colour_gradient(low = "#5BC2E7", high = "#FFFFFF")

#Adding smoother
#y ~ x is assumed if not given
plot + geom_smooth(method = "lm", formula = y ~ x)
```



- A confidence interval is also given around the line automatically

### Technical Details

Brian Caffo discusses the proof for least squares regression beta\_1 value in this video

### Lesson with swirl(): Least Squares Estimation

(No new content)

### Regression to the Mean

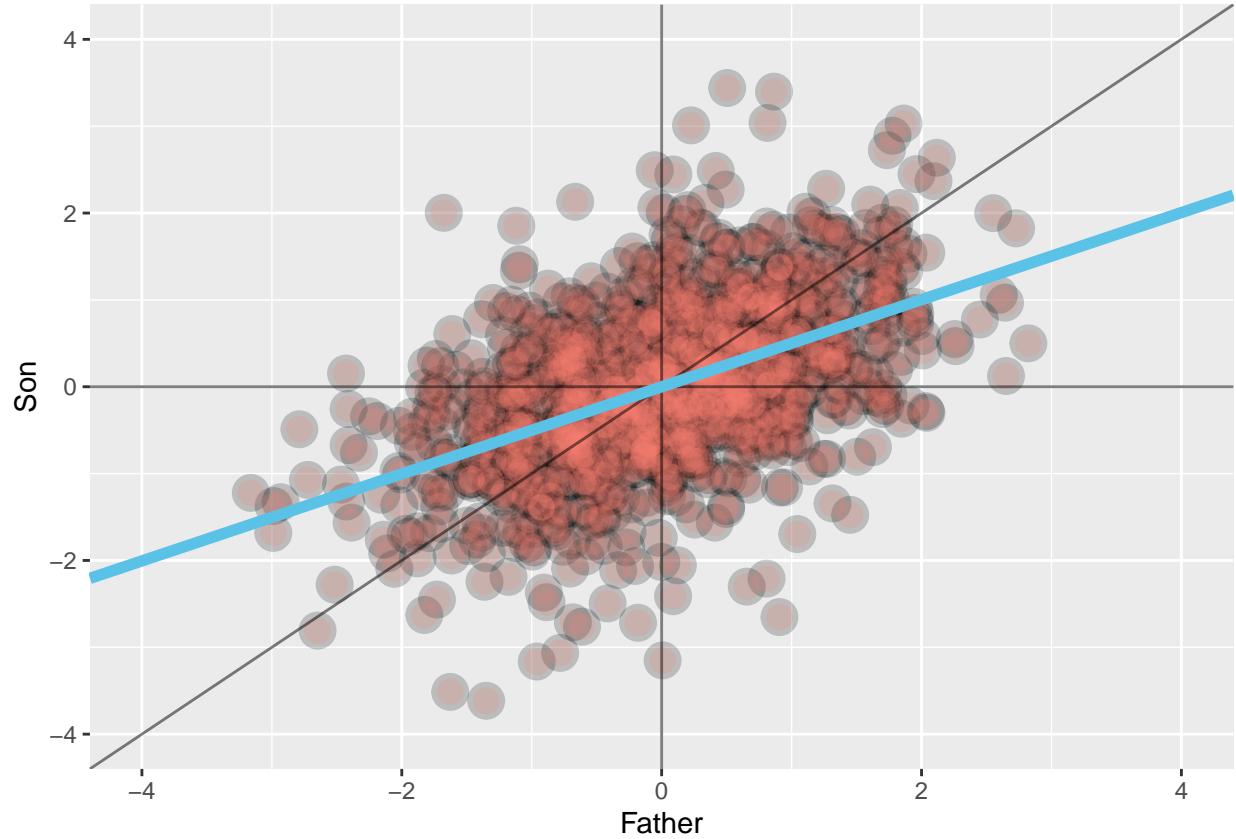
#### Regression to the Mean

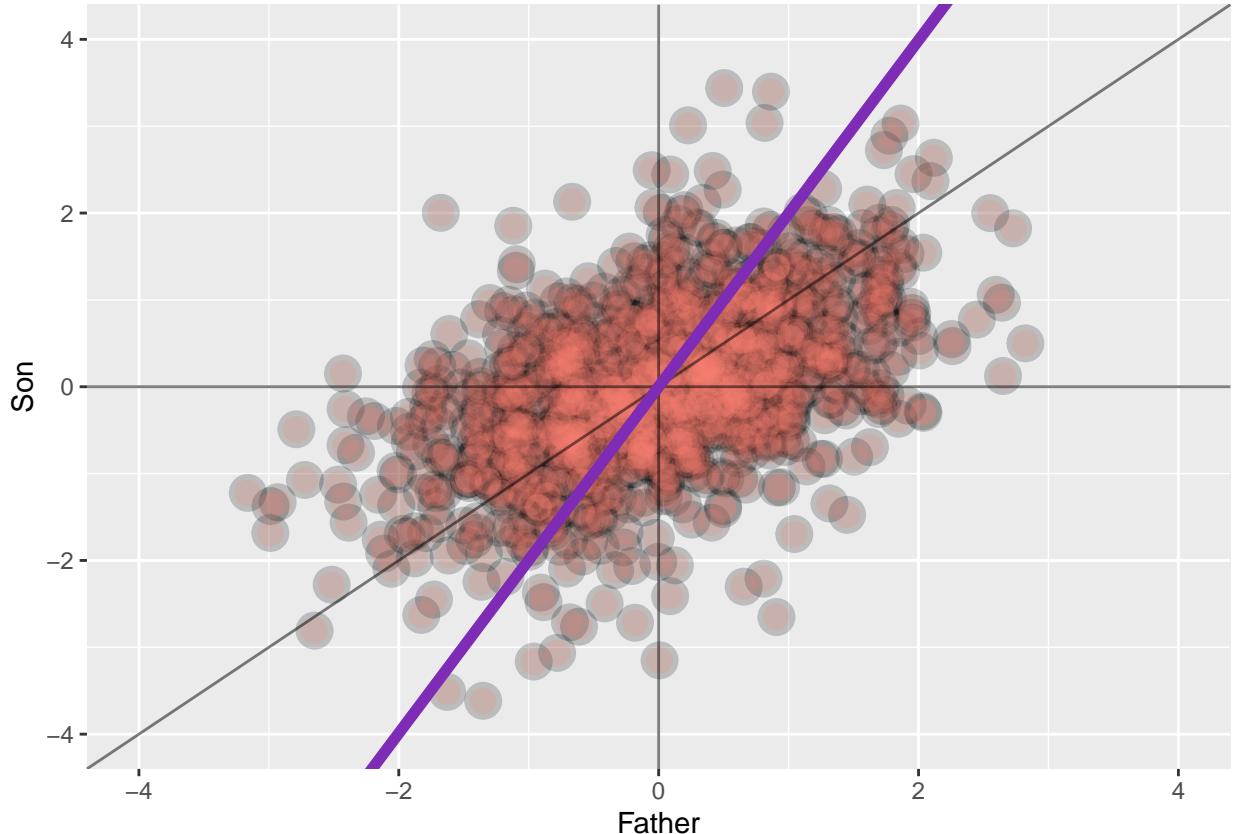
- $P(Y < x|X = x)$  gets bigger as  $x$  tends towards very large values.
  - + Similarly  $P(Y > x|X = x)$  gets bigger as  $x$  tends towards very small values.
- Regression line is like the intrinsic part of this relation
  - + Unless  $Cor(Y, X) = 1$  the intrinsic part isn't perfect

- Suppose we center  $X$  (child's height) and  $Y$  (parent's height) so that they both have a mean of 0
  - + Then, recall, our regression line passes through  $(0, 0)$
- We then normalize the data points too
  - + The slope of the regression line is  $\text{Cor}(Y, X)$ , regardless of which variable is the outcome (since both  $sds$  are 1)
- If the outcome is plotted on the horizontal axis the slope of the least squares line will be  $\frac{1}{\text{Cor}(Y, X)}$

### Plotting the Regression Implicitly

```
library(UsingR); data(father.son)
y <- father.son$sheight
x <- father.son$fheight
y <- (y - mean(y)) / sd(y)
x <- (x - mean(x)) / sd(x)
rho <- cor(x, y) #rho is std greek letter for correlations
plot <- ggplot(data.frame(Father = x, Son = y), aes(Father, Son)) +
 geom_point(size = 6, colour = "#000000", alpha = 0.2) +
 geom_point(size = 4, colour = "salmon", alpha = 0.2) +
 xlim(-4,4) +
 ylim(-4,4) + #Std. norm being +/- 4 is very unlikely
 geom_abline(intercept = 0, slope = 1, alpha = 0.5) +
 geom_vline(xintercept = 0, alpha = 0.5) +
 geom_hline(yintercept = 0, alpha = 0.5)
plot + geom_abline(intercept = 0, slope = rho, size = 2, colour = "#5BC2E7")
```





- \* The blue line is where the Father's height is the predictor and the Son's height is the outcome
- \* The purple line is where the Son's height is the predictor and the Father's height is the outcome ( $1/\rho$  because the outcome is on the horizontal axis)

### Lesson with `swirl()`: Residuals

- A residual is the distance between the actual data point and the regression line.
  - + I've previously heard it also called the "Unexplained Variation" since the distance from the mean value to data point is the "Total Variation (from the mean)", then the distance from the mean to reg. line is the "Explained Variation".
- You can get some info on a data set's residuals by calling `summary` on the results of `lm` as seen below

```
summary(lm(child ~ parent, galton))
```

```
##
Call:
lm(formula = child ~ parent, data = galton)
##
Residuals:
Min 1Q Median 3Q Max
-7.8050 -1.3661 0.0487 1.6339 5.9264
##
```

```

Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 23.94153 2.81088 8.517 <2e-16 ***
parent 0.64629 0.04114 15.711 <2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
Residual standard error: 2.239 on 926 degrees of freedom
Multiple R-squared: 0.2105, Adjusted R-squared: 0.2096
F-statistic: 246.8 on 1 and 926 DF, p-value: < 2.2e-16

```

- `est` will return the estimate,  $\hat{y}$
- `sqe` will calculate the sum of the squared residuals, also called the Residual Sum of Squares
- $\text{var}(\text{residuals}) = \text{var}(\text{data}) - \text{var}(\text{estimate})$   
+ As such the variance of residuals is always less than the variance of data
- The residuals shouldn't be correlated to either factor, if it did this may imply a different relationship is present

## Quiz 1

- Given...

```

x <- c(0.18, -1.54, 0.42, 0.95)
w <- c(2, 1, 3, 1)

```

Give the value of  $\mu$  that minimizes the least squares equation  $\sum_{i=1}^n w_i(x_i - \mu)^2$

```
sum(w * x) / sum(w)
```

```
[1] 0.1471429
```

- Given...

```

x <- c(0.8, 0.47, 0.51, 0.73, 0.36, 0.58, 0.57, 0.85, 0.44, 0.42)
y <- c(1.39, 0.72, 1.55, 0.48, 1.19, -1.59, 1.23, -0.65, 1.49, 0.05)

```

Fit the regression through the origin and get the slope treating  $y$  as the outcome and  $x$  as the regressor.

```
lm(y ~ x - 1)$coef
```

```

x
0.8262517

```

- Do `data(mtcars)` from the datasets package and fit the regression model with mpg as the outcome and weight as the predictor. Give the slope coefficient.

```

data(mtcars)
lm(mpg ~ wt, mtcars)$coef

```

```
(Intercept) wt
37.285126 -5.344472
```

4. Consider data with an outcome (Y) and a predictor (X). The standard deviation of the predictor is one half that of the outcome. The correlation between the two variables is 0.5. What value would the slope coefficient for the regression model with Y as the outcome and X as the predictor?

```
0.5 * 2/1
```

```
[1] 1
```

5. Students were given two hard tests and scores were normalized to have empirical mean 0 and variance 1. The correlation between the scores on the two tests was 0.4. What would be the expected score on Quiz 2 for a student who had a normalized score of 1.5 on Quiz 1?

```
beta1 <- 0.4 * 1/1
beta0 <- 0 - beta1*0
yhat <- beta0 + beta1*1.5
yhat
```

```
[1] 0.6
```

6. Given...

```
x <- c(8.58, 10.46, 9.01, 9.64, 8.86)
```

What is the value of the first measurement if x were normalized?

```
xn <- (x-mean(x))/sd(x)
xn[1]
```

```
[1] -0.9718658
```

7. Given...

```
x <- c(0.8, 0.47, 0.51, 0.73, 0.36, 0.58, 0.57, 0.85, 0.44, 0.42)
y <- c(1.39, 0.72, 1.55, 0.48, 1.19, -1.59, 1.23, -0.65, 1.49, 0.05)
```

What is the intercept for fitting the model with x as the predictor and y as the outcome?

```
lm(y ~ x)$coef
```

```
(Intercept) x
1.567461 -1.712846
```

8. You know that both the predictor and response have mean 0. What can be said about the intercept when you fit a linear regression?

- The intercept is the origin

9. Given...

```
x <- c(0.8, 0.47, 0.51, 0.73, 0.36, 0.58, 0.57, 0.85, 0.44, 0.42)
```

What value minimizes the sum of the squared distances between these points and itself?

```
mean(x)
```

```
[1] 0.573
```

10. Let the slope having fit Y as the outcome and X as the predictor be denoted as  $\beta_1$ . Let the slope from fitting X as the outcome and Y as the predictor be denoted as  $\gamma_1$ . Suppose that you divide  $\beta_1$  by  $\gamma_1$  What is this ratio always equal to?

- $\beta_1 = \text{Cor}(Y, X) \frac{\text{sd}(Y)}{\text{sd}(X)}$
- $\gamma_1 = \text{Cor}(Y, X) \frac{\text{sd}(X)}{\text{sd}(Y)}$
- $\frac{\beta_1}{\gamma_1} = \frac{\text{Cor}(Y, X) * \text{sd}(Y) / \text{sd}(X)}{\text{Cor}(Y, X) * \text{sd}(X) / \text{sd}(Y)} = \frac{\text{sd}(Y) * \text{sd}(Y)}{\text{sd}(X) * \text{sd}(X)} = \frac{\text{Var}(Y)}{\text{Var}(X)}$

## Linear Regression & Multivariable Regression

### Statistical Linear Regression Models

#### Statistical Linear Regression Models

##### Basic Regression Model with Additive Gaussian Errors

- Consider developing a probabilistic model for linear regression  
$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$$
  - + Here the  $\epsilon_i$  are assumed iid  $N(0, \sigma^2)$
  - Can be thought of as accumulated variables that aren't modeled by act on the response as iid gaussian errors  $+ E[Y_i | X_i = x_i] = \mu_i = \beta_0 + \beta_1 x_i$
  - +  $\text{Var}(Y_i | X_i = x_i) = \sigma^2$

### Interpreting Coefficients

#### Intercept

- $\beta_0$  is the expected value of the response when the predictor is 0  
$$E[Y | X = 0] = \beta_0 + \beta_1 \times 0 = \beta_0$$
  - + This isn't always a value of interest, for example when  $X = 0$  is impossible (x represents weight) or far outside of the range of data.
- A solution to non-interpretable intercepts is to shift the equation by some value,  $a$  then define a new intercept,  $\tilde{\beta}_0$ .  
$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i = \beta_0 + a\beta_1 + \beta_1(X_i - a) + \epsilon_i = \tilde{\beta}_0 + \beta_1(X_i - a) + \epsilon_i$$
  - + Shifting your  $X$  values by value  $a$  changes the intercept, but not the slope.
  - + Often  $a$  is set to  $\bar{X}$  so that the intercept is interpreted as the expected response at the average X value.

#### Slope

- $\beta_1$  is the expected change in response for a 1 unit change in the predictor
- Consider the impact of changing the units of  $X$ .  

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i = \beta_0 + \frac{\beta_1}{a} (X_i a) + \epsilon_i = \beta_0 + \tilde{\beta}_1 (X_i a) + \epsilon_i$$
+ Since  $\beta_1$  is in units of Y/X we divide by the factor,  $a$ , that we're multiplying with  $X_i$ .
- Example:  $X$  is height in  $m$  and  $Y$  is weight in  $kg$ . Then  $\beta_1$  is  $kg/m$ . Converting  $X$  to  $cm$  implies multiplying  $X$  by  $100cm/m$ . To get  $\beta_1$  in the right units, we have to divide by  $100cm/m$  to get it to have the right units.  

$$Xm \times \frac{100cm}{m} = (100X)cm \text{ and } \beta_1 \frac{kg}{m} \times \frac{1m}{100cm} = (\frac{\beta_1}{100}) \frac{kg}{cm}$$

## Linear Regression for Prediction

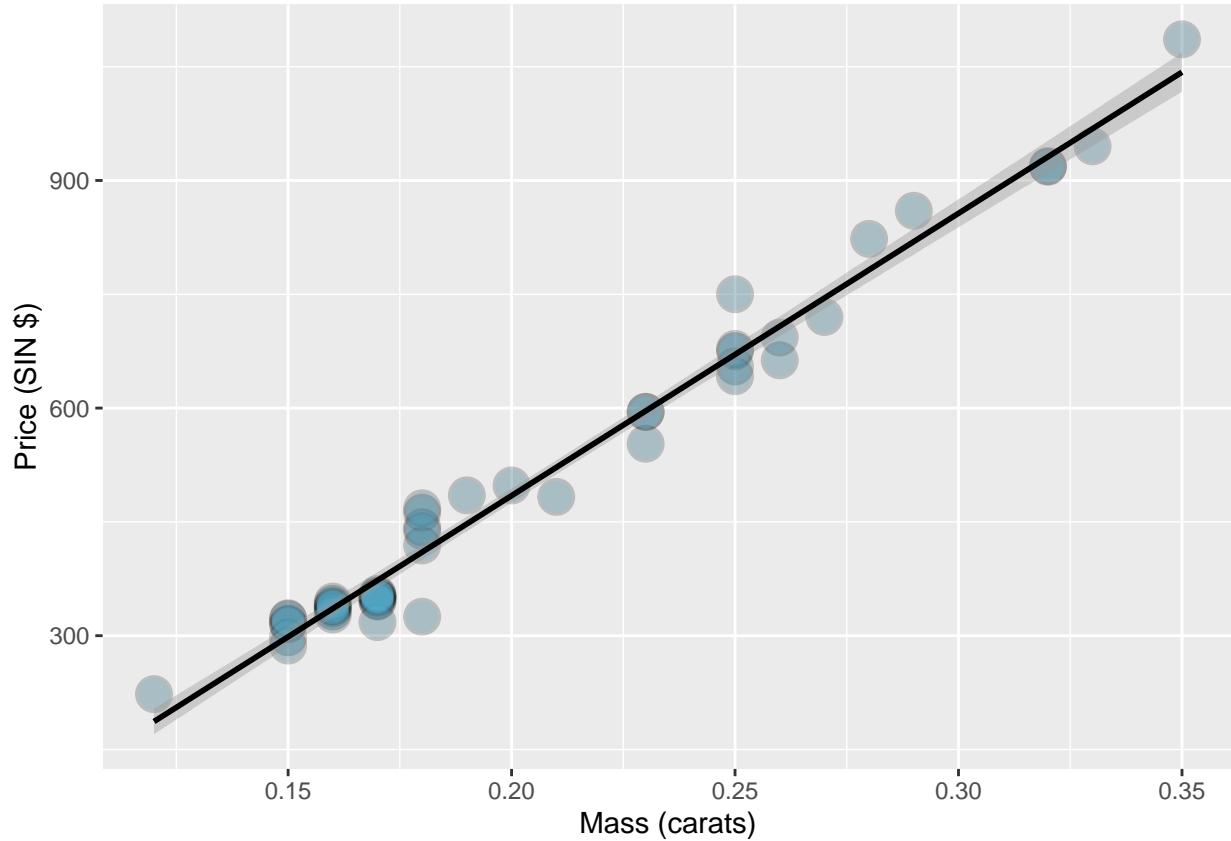
- We can get a prediction for  $Y$ ,  $\hat{y}$  by plugging in the  $X$  that we want into our model  

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$$

## Example using diamond Data

- The data in this example is diamond prices (in Singapore dollars) and diamond weight in carats (1 carat = 0.2 g).

```
library(UsingR); data(diamond); library(tidyverse)
plot <- ggplot(diamond, aes(carat, price)) +
 xlab("Mass (carats)") +
 ylab("Price (SIN $)") +
 geom_point(size = 6, colour = "#000000", alpha = 0.2) +
 geom_point(size = 5, colour = "#5BC2E7", alpha = 0.2)
plot + geom_smooth(method = "lm", colour = "#000000", formula = y ~ x)
```



#### Creating a Model

```
Fitting the linear regression model
fit <- lm(price ~ carat, data = diamond)
coef(fit)
```

```
(Intercept) carat
-259.6259 3721.0249
```

- We estimate an expected 3721.02 (SIN) dollar increase in price for every increase of 1 carat in mass of diamonds.
- The intercept, -259.63 is the expected price of a 0 carat diamond, which doesn't make sense to interpret.
  - + As such we'll mean center our reg. line

#### Centering Model on the Mean

```
cfit <- lm(price ~ I(carat - mean(carat)), data = diamond)
cfit$coef
```

```
(Intercept) I(carat - mean(carat))
500.0833 3721.0249
```

- To do arithmetic operations in the formula in `lm` you have to surround the operation with the `I` function
- The slope has not changed

- The intercept has changed to 500, the expected price for the average sized diamond of the data (0.204 carats).

## Changing Units in the Model

- Change unit to 1/10 of a carat

```
tenthfit <- lm(price ~ I(carat * 10), data = diamond)
coef(tenthfit)
```

```
(Intercept) I(carat * 10)
-259.6259 372.1025
```

- So now the slope is interpreted as a 372.1 dollar increase for every additional 0.1 carats of diamond.

## Estimating a Value

```
newDiamonds <- c(0.16, 0.27, 0.34)
#Computing manually
fit$coef[1] + fit$coef[2] * newDiamonds
```

```
[1] 335.7381 745.0508 1005.5225
```

#Using predict function

```
results <- predict(fit, newdata = data.frame(carat = newDiamonds))
names(results) <- as.character(newDiamonds) #renaming not required
results
```

```
0.16 0.27 0.34
335.7381 745.0508 1005.5225
```

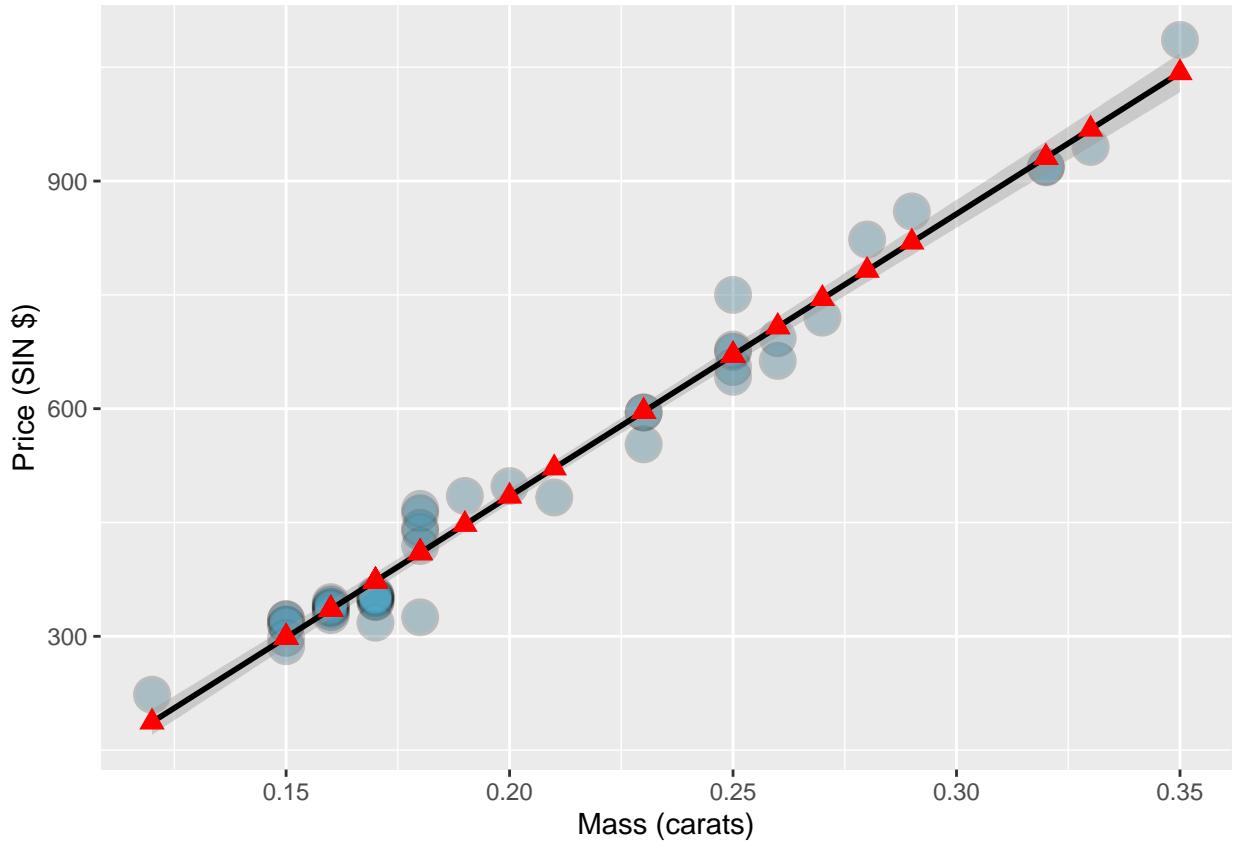
#Using predict without 'newdata' will return y-hat for given x values

```
predict(fit)
```

|    | 1        | 2         | 3        | 4        | 5        | 6        | 7        | 8        |
|----|----------|-----------|----------|----------|----------|----------|----------|----------|
| ## | 372.9483 | 335.7381  | 372.9483 | 410.1586 | 670.6303 | 335.7381 | 298.5278 | 447.3688 |
| ## | 9        | 10        | 11       | 12       | 13       | 14       | 15       | 16       |
| ## | 521.7893 | 298.5278  | 410.1586 | 782.2611 | 335.7381 | 484.5791 | 596.2098 | 819.4713 |
| ## | 17       | 18        | 19       | 20       | 21       | 22       | 23       | 24       |
| ## | 186.8971 | 707.8406  | 670.6303 | 745.0508 | 410.1586 | 335.7381 | 372.9483 | 335.7381 |
| ## | 25       | 26        | 27       | 28       | 29       | 30       | 31       | 32       |
| ## | 372.9483 | 410.1586  | 372.9483 | 410.1586 | 372.9483 | 298.5278 | 372.9483 | 931.1020 |
| ## | 33       | 34        | 35       | 36       | 37       | 38       | 39       | 40       |
| ## | 931.1020 | 298.5278  | 335.7381 | 335.7381 | 596.2098 | 596.2098 | 372.9483 | 968.3123 |
| ## | 41       | 42        | 43       | 44       | 45       | 46       | 47       | 48       |
| ## | 670.6303 | 1042.7328 | 410.1586 | 670.6303 | 670.6303 | 298.5278 | 707.8406 | 298.5278 |

```
plot + geom_smooth(method = "lm", colour = "#000000", formula = y ~ x) +
 geom_point(aes(y = as.numeric(predict(fit)))),
```

```
size = 3, color = "#FF0000", shape = 17)
```



## Residuals

### Residuals

- The residuals are the variation from the regression line, that is left unexplained by our model,  $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$  where  $\epsilon_i \sim N(0, \sigma^2)$ .
- Observed outcome  $i$  is  $Y_i$  at predictor value  $X_i$
- Predicted outcome  $i$  is  $\hat{Y}_i$  at predictor value  $X_i$  is  $\hat{Y}_i = \hat{\beta}_0 + \hat{\beta}_1 X_i$
- Residual,  $e_i$ , is the difference between the observed and predicted outcome:  $e_i = Y_i - \hat{Y}_i$ .  
+ This is the vertical distance between the observed data point and the regression line
- Least squares minimizes these residuals, the equation  $\sum_{i=1}^n e_i^2$
- The  $e_i$  can be thought of as estimates of the  $\epsilon_i$

### Properties of the Residuals

- $E[e_i] = 0$
- If an intercept is included,  $\sum_{i=1}^n e_i = 0$
- If a regressor variable,  $X_i$ , is included in the model  $\sum_{i=1}^n e_i X_i = 0$
- Residuals are useful for investigating poor model fit  
+ Residual plots can highlight these poor fits
- Residuals can be thought of as the outcome ( $Y$ ) with the linear association of the predictor ( $X$ ) removed.
- One differentiates residual variation (variation after removing the predictor) from systematic variation (variation explained by the regression model).

## Residuals, Coding Example

- Using diamond dataset again

```

data("diamond")
y <- diamond$price
x <- diamond$carat
fit <- lm(y ~ x)

e <- resid(fit) #Getting residuals

yhat <- predict(fit)

Showing residuals are the same as y - yhat (within a floating point error)
max(abs(e - (y - yhat)))

[1] 5.258016e-13

And again, but manually entering the equation for yhat
max(abs(e - (y - (coef(fit)[1] + coef(fit)[2] * x)))))

[1] 5.258016e-13

#Showing sum of resid and resid*x are both 0
sum(e)

[1] -3.93019e-14

sum(e * x)

[1] -1.249001e-15

#Plotting the residuals
plot <- ggplot(data.frame(x = x, y = y, resid = e), aes(x, resid)) +
 geom_segment(aes(xend = x, yend = 0), colour = "#FF0000") +

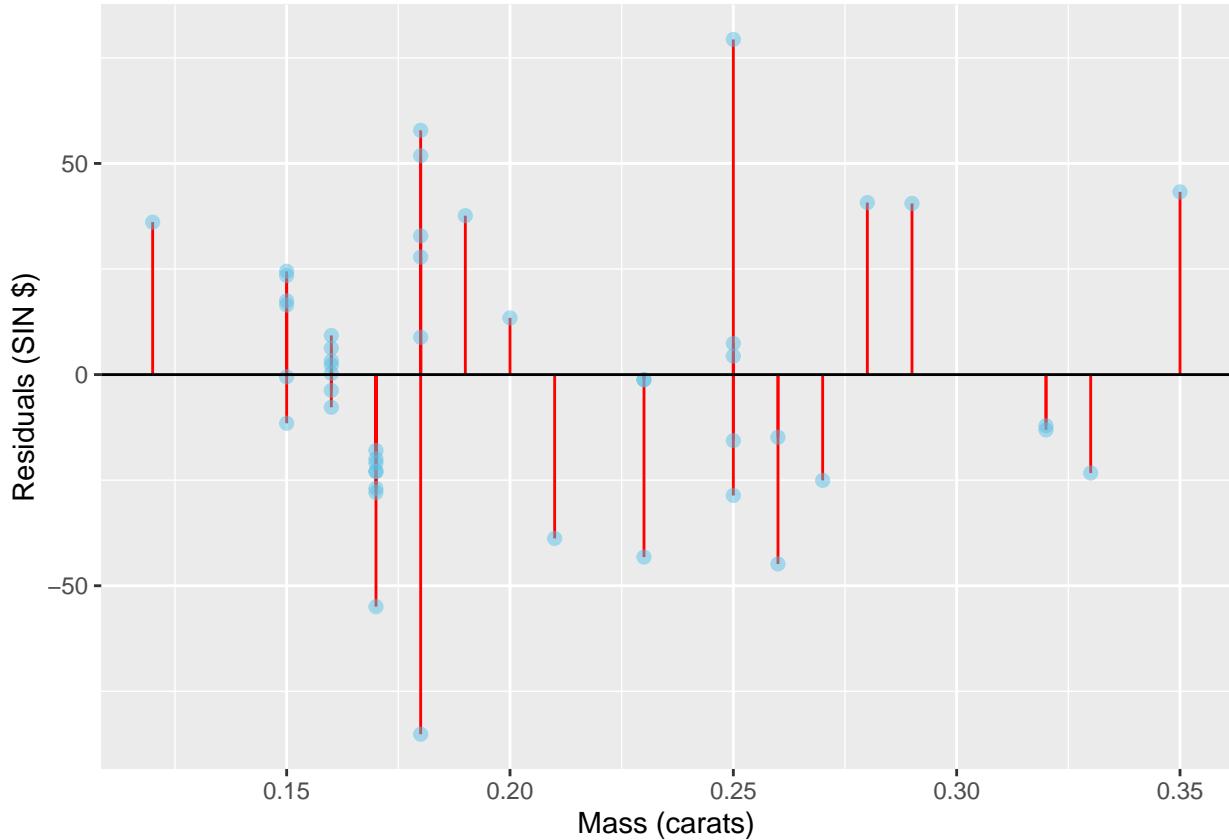
```

```

geom_point(size = 2, colour = "#5BC2E7", alpha = 0.5) +
xlab("Mass (carats)") +
ylab("Residuals (SIN $)") +
geom_hline(yintercept = 0, color = "#000000")

```

plot



### Using Residual Plot to Detect a Poorly Fit Model

- We're going to generate some data that looks linear but actually has an underlying relation to it that will become more apparent after plotting the residuals

```

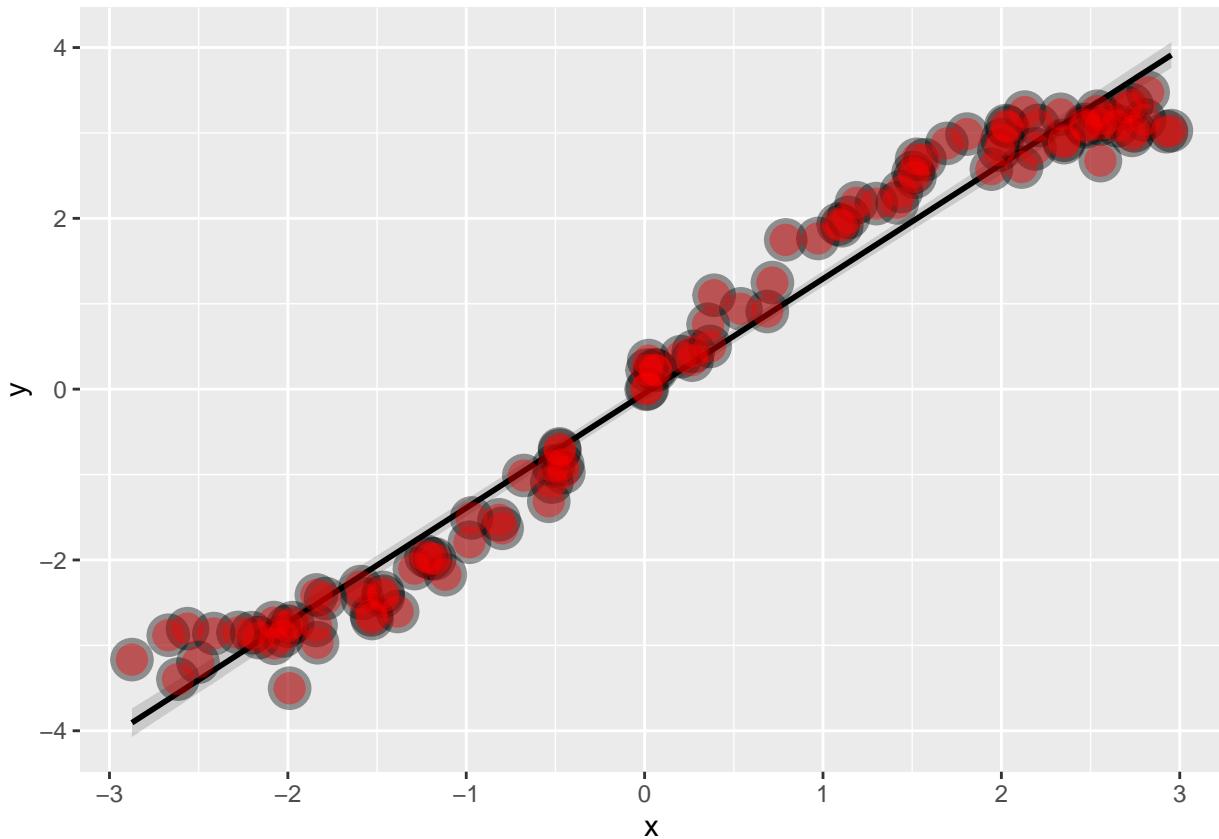
set.seed(1618033)
x <- runif(100, -3, 3)
y <- x + sin(x) + #Y is related with sin(x), lm will expose the sin(x) rel.
 rnorm(100, sd = .2) # For noise
plot <- ggplot(data.frame(x = x, y = y), aes(x,y)) +
 geom_smooth(method = "lm", colour = "#000000") +
 geom_point(size = 7, colour = "#000000", alpha = 0.4) +
 geom_point(size = 5, colour = "#FF0000", alpha = 0.4)

residplot <- ggplot(data.frame(x = x, resid = resid(lm(y ~ x))),
 aes(x, resid)) +

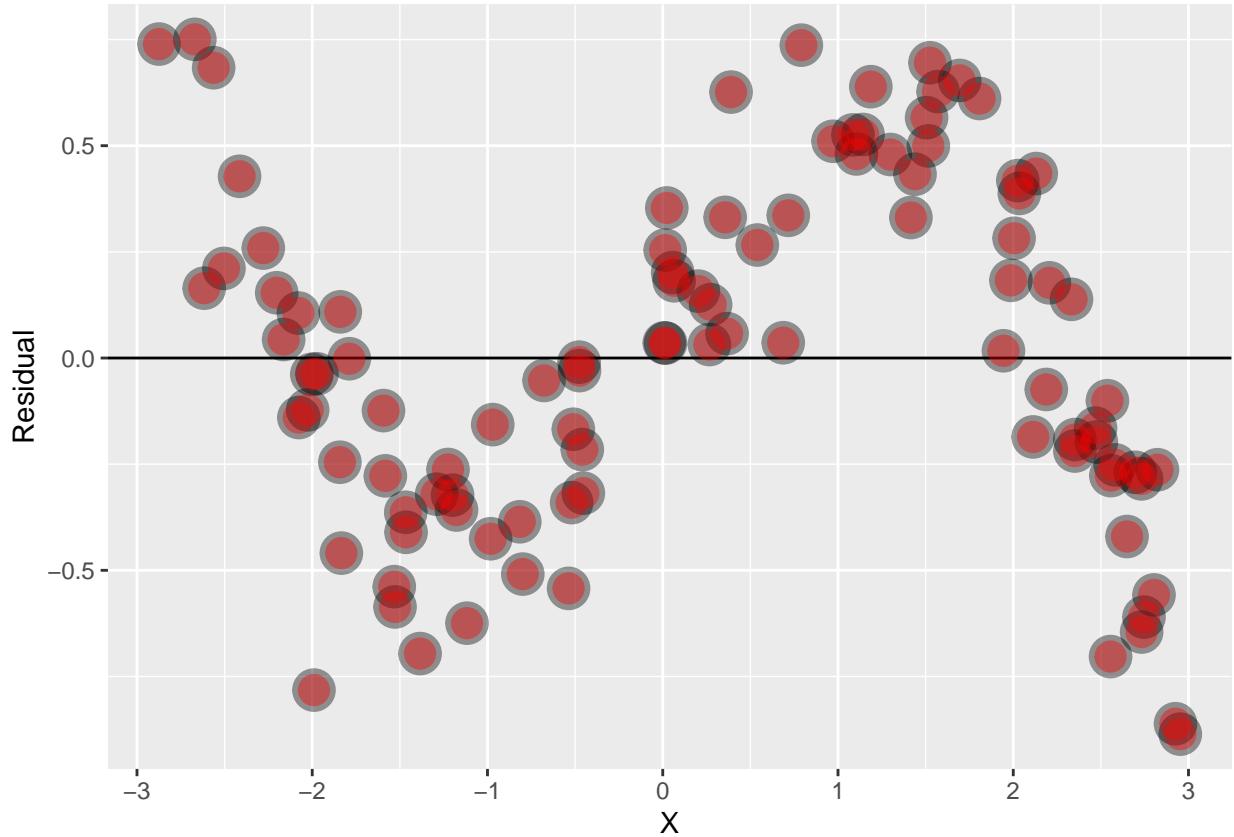
```

```
geom_hline(yintercept = 0) +
 geom_point(size = 7, colour = "#000000", alpha = 0.4) +
 geom_point(size = 5, colour = "#FF0000", alpha = 0.4) +
 labs(x = "X", y = "Residual")
plot

`geom_smooth()` using formula 'y ~ x'
```



```
residplot
```



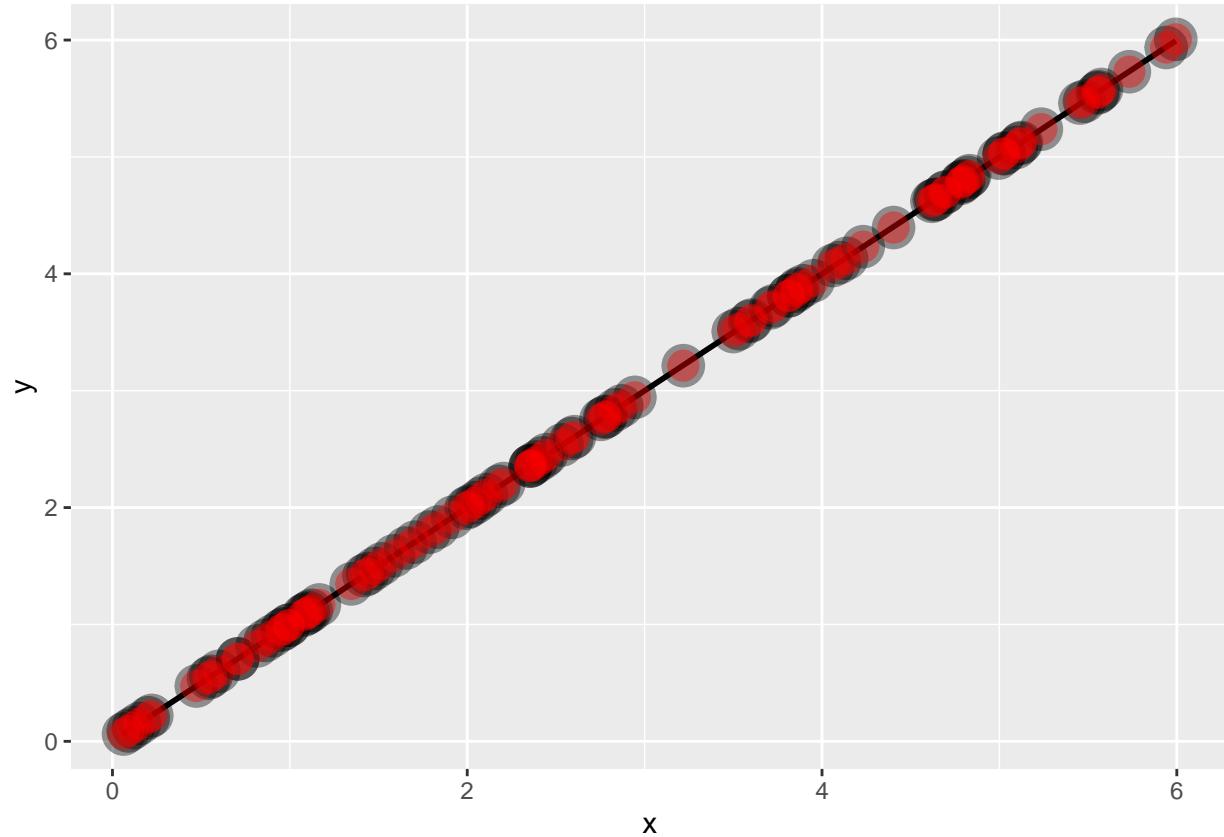
- A secondary pattern can be seen in the residual plot, indicating there might be a better model than a line.

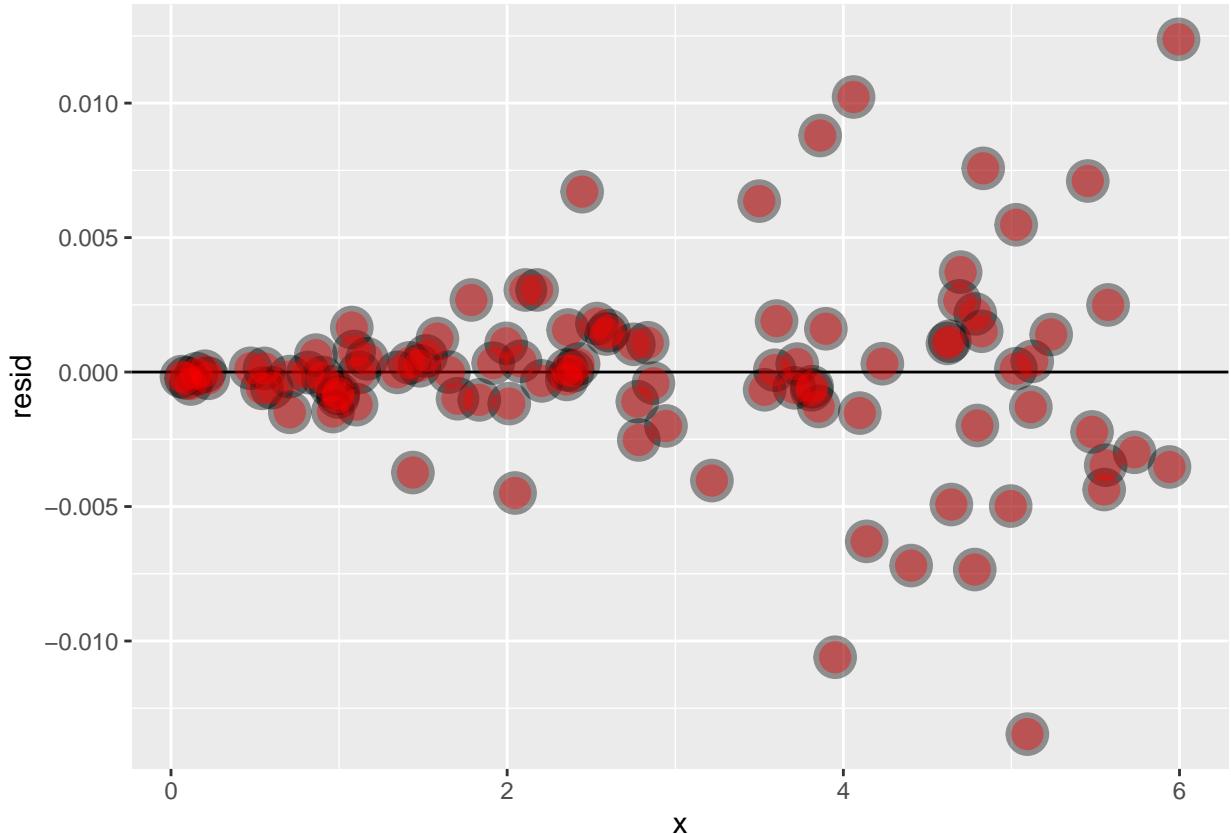
### Detecting Heteroskedasticity with a Residual Plot

```

x <- runif(100, 0, 6)
y <- x + rnorm(100, mean = 0, sd = 0.001 * x)#sd increases as x increases
plot <- ggplot(data.frame(x = x, y = y), aes(x,y)) +
 geom_smooth(method = "lm", colour = "black") +
 geom_point(size = 7, colour = "#000000", alpha = 0.4) +
 geom_point(size = 5, colour = "#FF0000", alpha = 0.4)
residplot <- ggplot(data.frame(x = x, resid = resid(lm(y ~ x))),
 aes(x,resid)) +
 geom_hline(yintercept = 0, colour = "#000000") +
 geom_point(size = 7, colour = "#000000", alpha = 0.4) +
 geom_point(size = 5, colour = "#FF0000", alpha = 0.4)
plot
'geom_smooth()' using formula 'y ~ x'

```





\* The plot looks linear, but plotting the residuals reveals an underlying pattern

## Residual Variance

### Estimating Residual Variance

- Model:  $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$  where  $\epsilon_i \sim N(0, \sigma^2)$
- The mean linear estimate of  $\sigma^2$  is  $\frac{1}{n} \sum_{i=1}^n e_i^2$ , the average squared residual
- Most people use:  

$$\hat{\sigma}^2 = \frac{1}{n-2} \sum_{i=1}^n e_i^2$$
+ with  $n - 2$  instead of  $n$  so that  $E[\hat{\sigma}^2] = \sigma^2$

### Diamond Example

```
y <- diamond$price
x <- diamond$carat
n <- length(y)

#Solving resid s.d. implicitly
sqrt(sum(resid(fit)^2) / (n - 2))
```

```

[1] 31.84052
#Getting resid deviation with functions
fit <- lm(y ~ x)
summary(fit)$sigma

[1] 31.84052
#You can see the value in the summary print out here:
summary(fit)

##
Call:
lm(formula = y ~ x)
##
Residuals:
Min 1Q Median 3Q Max
-85.159 -21.448 -0.869 18.972 79.370
##
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) -259.63 17.32 -14.99 <2e-16 ***
x 3721.02 81.79 45.50 <2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
Residual standard error: 31.84 on 46 degrees of freedom
Multiple R-squared: 0.9783, Adjusted R-squared: 0.9778
F-statistic: 2070 on 1 and 46 DF, p-value: < 2.2e-16

```

## Summarizing Variation

- **Total Variability** - the variability around an intercept (mean only regression) +  $\sum_{i=1}^n (Y_i - \bar{Y})^2$  + Sum of Regression & Error Variability
- **Regression Variability** - the variability that is explained by adding the predictor +  $\sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2$
- **Error Variability** - what's leftover around the regression line +  $\sum_{i=1}^n (Y_i - \hat{Y})^2$

## R Squared, the Coefficient of Determination

- R squared is the percentage of the total variability that is explained by the linear relationship with the predictor
- $$R^2 = \frac{\sum_{i=1}^n (\hat{Y}_i - \bar{Y})^2}{\sum_{i=1}^n (Y_i - \bar{Y})^2}$$

- $R^2$  is the percentage of variation explained by the regression model
- $0 \leq R^2 \leq 1$
- $R^2$  is the sample correlation squared
- $R^2$  can be a misleading summary of model fit
  - + Deleting data can inflate  $R^2$
  - + (For later,) Adding terms to a regression model always increases  $R^2$
- Execute `example(anscombe)` to see the following data:
  - + Basically same mean and variance of X and Y
  - + Identical correlations (hence the same  $R^2$  value)
  - + Same linear regression relationship

### Lesson with `swirl()`: Residual Variation

- `deviance` will calculate the sum of the squares of a lm

## Inference in Regression

### Inference in Regression

#### Recall Our Model and Fitted Values

- Model:
  - +  $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$
  - +  $\epsilon \sim N(0, \sigma^2)$ , an error term
  - +  $\hat{\beta}_1 = Cor(Y, X) \frac{Sd(Y)}{Sd(X)}$
  - +  $\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$
- We assume that the true model is known for most of this course

#### Review Some Statistical Inference Concepts

- Statistics like  $\frac{\hat{\theta} - \theta}{\hat{\sigma}_{\hat{\theta}}}$  often have the following properties:
  - + Is normally distributed and has a finite sample Student's T distribution if the estimated variance is replaced with a sample estimate (under normality assumptions).
  - + Can be used to test  $H_0 : \theta = \theta_0$  versus  $H_a : \theta >, <, \neq \theta_0$
  - + Can be used to create a confidence interval for  $\theta$  via  $\hat{\theta} \pm Q_{1-\alpha/2} \hat{\sigma}_{\hat{\theta}}$  where  $Q_{1-\alpha/2}$  is the relevant quantile from either a normal or T distribution
- In the case of regression with iid sampling assumptions and normal errors, our inferences will follow very similarly to what was discussed in the inference class.
- Under assumptions on the ways in which the  $X$  values are collected the iid sampling model, and mean model, the normal results hold to create intervals and confidence intervals

## Explanation

- Variance of our regression slope,  $\sigma_{\hat{\beta}_1}^2$ , tells both how variable points are around the regression line,  $\sigma^2$ , and how variable the points are from the mean  

$$\sigma_{\hat{\beta}_1}^2 = \text{Var}(\hat{\beta}_1) = \sigma^2 / \sum_{i=1}^n (X_i - \bar{X})^2$$
  - + This implies spreaded out points will give a lower variance for a slope
  - Thus large cluster of points very far apart would give the best variance, although this lm would assume the uncollected data between the clusters is linear
- Variance of the intercept,  $\sigma_{\hat{\beta}_0}^2$ , is less informative but still can provide some information.  

$$\sigma_{\hat{\beta}_0}^2 = \text{Var}(\hat{\beta}_0) = \left( \frac{1}{n} + \frac{\bar{X}^2}{\sum_{i=1}^n (X_i - \bar{X})^2} \right) \sigma^2$$
  - In both these cases, in practice,  $\sigma$  is replaced by its estimate
  - Under iid gaussian errors,  $\frac{\hat{\beta}_j - \beta_j}{\sigma_{\hat{\beta}_j}}$ , follows a  $t$  distribution with  $n - 2$  degrees of freedom and a normal distribution for large  $n$   
    - + This can be used to create confidence intervals and perform hypothesis tests.

## Coding Example

- Showing R is calculating all these values as we have given

```
library(UsingR); data(diamond)
y <- diamond$price
x <- diamond$carat
n <- length(y)
beta1 <- cor(y, x) * sd(y) / sd(x) #Slope
beta0 <- mean(y) - beta1 * mean(x) #y-intercept
e <- y - (beta0 + beta1 * x) #resids
sigma <- sqrt(sum(e^2) / (n-2)) #est. sd for resids
ssx <- sum((x - mean(x))^2) #Numerator of variance calculation
seBeta0 <- sqrt((1 / n + mean(x)^2 / ssx)) * sigma #s.e. of intercept
seBeta1 <- sigma / sqrt(ssx) #s.e. of slope
tBeta0 <- beta0 / seBeta0 #t statistic for intercept; H_0: beta0=0
tBeta1 <- beta1 / seBeta1 # t statistic for slope
#Relevant p values
pBeta0 <- 2 * pt(abs(tBeta0), df = n - 2, lower.tail = FALSE)
pBeta1 <- 2 * pt(abs(tBeta1), df = n - 2, lower.tail = FALSE)
coefTable <- rbind(c(beta0, seBeta0, tBeta0, pBeta0), c(beta1, seBeta1, tBeta1, pBeta1))
colnames(coefTable) <- c("Estimate", "Std. Error", "t value", "P(>|t|)")
rownames(coefTable) <- c("(Intercept)", "x")
coefTable

Estimate Std. Error t value P(>|t|)
(Intercept) -259.6259 17.31886 -14.99094 2.523271e-19
x 3721.0249 81.78588 45.49715 6.751260e-40
```

```

summary(lm(y ~ x))$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) -259.6259 17.31886 -14.99094 2.523271e-19
x 3721.0249 81.78588 45.49715 6.751260e-40

```

## Generating Confidence Intervals

```

fit <- lm(y ~ x)
sumCoef <- summary(fit)$coef

#Intercept
sumCoef[1, 1] + c(-1, 1) * qt(0.975, df = fit$df) * sumCoef[1, 2]

[1] -294.4870 -224.7649

#Slope; Change in x per 1 y unit
sumCoef[2, 1] + c(-1, 1) * qt(0.975, df = fit$df) * sumCoef[2, 2]

[1] 3556.398 3885.651

```

## Prediction

- Consider predicting  $Y$  at a value of  $X$ 
  - + Predicting the price of a diamond given the carat
  - + Predicting the height of a child given the height of the parents
- The obvious estimate for prediction at point  $x_0$  is  $\hat{\beta}_0 + \hat{\beta}_1 x_0$
- A standard error is needed to create a prediction interval
- There's a distinction between intervals for the regression line at points and the prediction of what a  $y$  would be at point  $x_0$
- Line at  $x_0$  std. error:  $\hat{\sigma} \sqrt{\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum_{i=1}^n (X_i - \bar{X})^2}}$ 
  - + Variance will be the least when predicting the average of  $x$
  - + The denominator is how variable the ' $x$ 's are, so the more variability the less this error
- Prediction interval std. error at  $x_0$ :  $\hat{\sigma} \sqrt{1 + \frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum_{i=1}^n (X_i - \bar{X})^2}}$

## Generating Prediction Intervals in Diamond Data Set

```

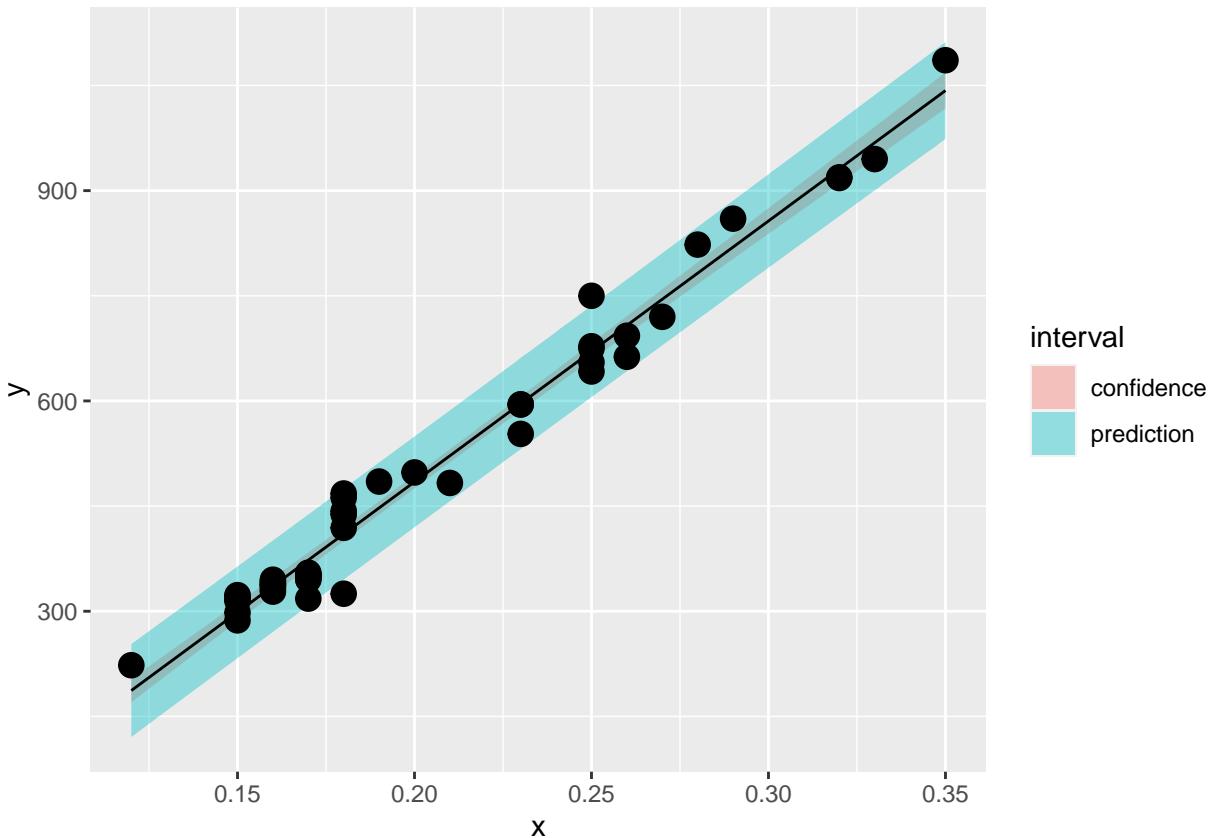
newx <- data.frame(x = seq(min(x), max(x), length = 100))
##Data Wranglin'
p1 <- data.frame(predict(fit, newdata = newx, interval = ("confidence")))
p2 <- data.frame(predict(fit, newdata = newx, interval = ("prediction")))
#p1 is giving confidence for each interval

```

```
#p2 is giving the actual prediction at that data value
```

```
p1$interval = "confidence"
p2$interval = "prediction"
p1$x <- newx$x
p2$x <- newx$x
dat <- rbind(p1, p2)
names(dat)[1] <- "y"

##Plotting
plot <- ggplot(dat, aes(x, y)) +
 geom_ribbon(aes(ymin = lwr, ymax = upr, fill = interval),
 alpha = 0.4) +
 geom_line() +
 geom_point(data = data.frame(x = x, y = y), size = 4)
plot
```



\* Blue is prediction area, salmon color is precipitating the line at each spot.

+ Both get narrower near middle since we're more confident as we are closer to the mean of x.

## Lesson with `swirl()`: Introduction to Multivariable Regression

- Once we identify one regression line we can eliminate it to reduce the dimensions of data
- By subtracting the mean from each variable, the regression line goes through the origin, hence its intercept is zero.
  - + thus we eliminate one of the two regressors, the constant, leaving just the predicting variable
  - + Subtracting the means is a special case of Gaussian Elimination
  - We pick one regressor and replace all other variables by the residuals of their regressions against that one
  - + Subtracting the mean is equivalent to replacing a variable by the residual of its regression against 1.
  - as such `lm(child ~ 1, galton)` will give an intercept of the mean, with a slope of 0.

## Eliminate Variable Function

- First we want a function to regress the given variable on the given predictor, suppressing the intercept, and return the residual.

```
regressOneOnOne <- function(predictor, other, dataframe){
 # Point A. Create a formula such as Girth ~ Height -1
 formula <- paste0(other, " ~ ", predictor, " - 1")
 # Use the formula in a regression and return the residual.
 resid(lm(formula, dataframe))
}
```

- Using that function we can write another function to eliminate the specified predictor from the dataframe by regressing all other variables on that predictor and returning a data frame containing the residuals of those regressions.

```
eliminate <- function(predictor, dataframe){
 # Find the names of all columns except the predictor.
 others <- setdiff(names(dataframe), predictor)
 # Calculate the residuals of each when regressed against
 # the given predictor with the previous function
 temp <- sapply(others, function(other)regressOneOnOne(predictor, other, dataframe))
 # convert matrix of resids to a data frame and return.
 as.data.frame(temp)
}
```

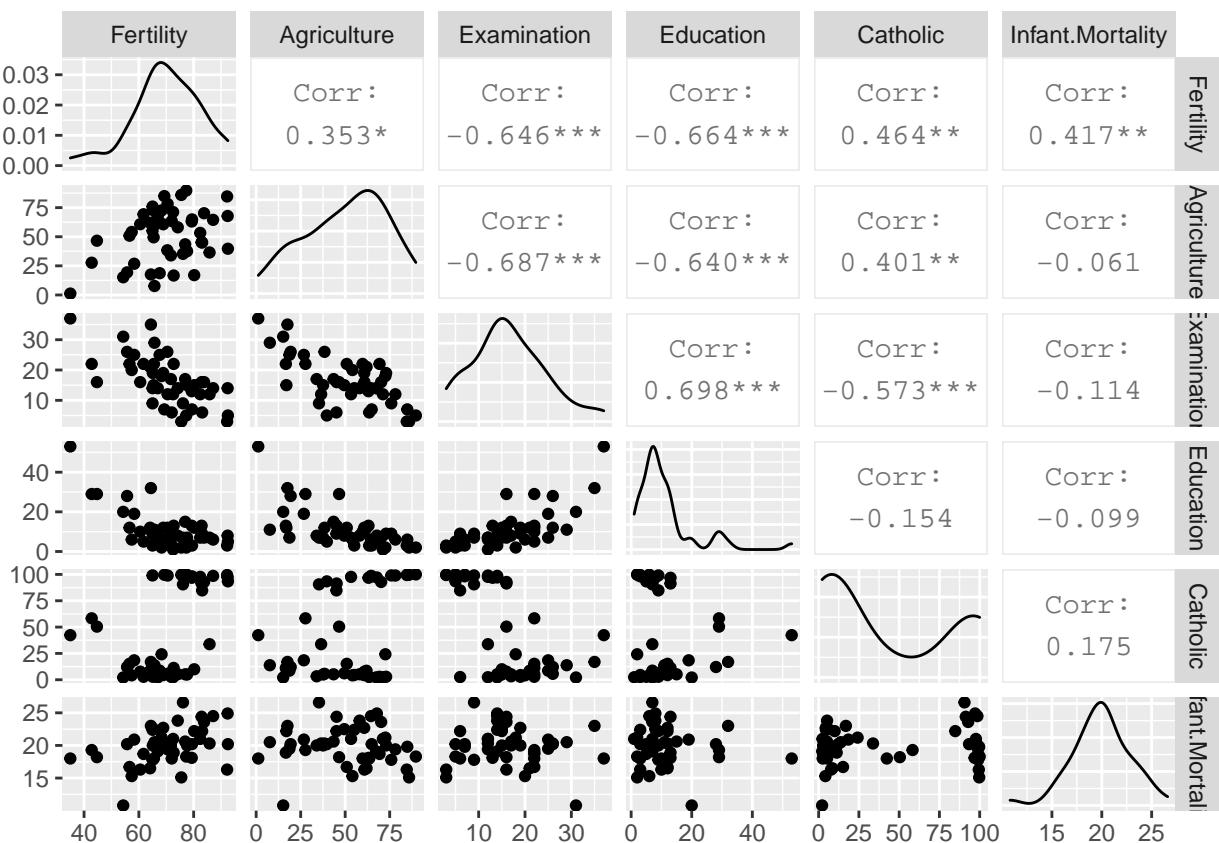
- We could use eliminate multiple times to get rid of more and more variables, each time essentially using Gaussian elimination to re-express all the terms such that they are plotted with the mean's intersection as the origin. This in turn replaces the outcome and all other regressors by their residuals against the chosen variable.

## Lesson with `swirl()`: MultiVar Examples

- This data was gathered in 1888 in Switzerland, below are explanations for the variables, all of which except fertility represent proportions of the population.

- + Fertility - a common standardized fertility measure
  - + Agriculture - % of males involved in agriculture as occupation
  - + Examination - % draftees receiving highest mark on army examination
  - + Education - % education beyond primary school for draftees
  - + Catholic - % catholic (as opposed to protestant)
  - + Infant.Mortality - live births who live less than 1 year
- Check out this 6 by 6 array of scatterplots showing pairwise relations between the variables.
  - + Lol, jk they just show the points plotted because I couldn't figure out `ggpairs`, I ought to use lattice for this task, but I'm just going to move on because I've spent enough time on it :(

```
data("swiss"); library(GGally)
ggpairs(swiss, lower = list(continuous = "smooth"))
```



## Reading Multiple Explanatory Variables

```
results <- summary(lm(Fertility ~ ., data = swiss))
results

##
Call:
lm(formula = Fertility ~ ., data = swiss)
##
Residuals:
```

```

Min 1Q Median 3Q Max
-15.2743 -5.2617 0.5032 4.1198 15.3213
##
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 66.91518 10.70604 6.250 1.91e-07 ***
Agriculture -0.17211 0.07030 -2.448 0.01873 *
Examination -0.25801 0.25388 -1.016 0.31546
Education -0.87094 0.18303 -4.758 2.43e-05 ***
Catholic 0.10412 0.03526 2.953 0.00519 **
Infant.Mortality 1.07705 0.38172 2.822 0.00734 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ',' 1
##
Residual standard error: 7.165 on 41 degrees of freedom
Multiple R-squared: 0.7067, Adjusted R-squared: 0.671
F-statistic: 19.76 on 5 and 41 DF, p-value: 5.594e-10

```

- The **Coefficients** table states the Estimate/Slope for each explanatory variable to the dependent variable. For example:
  - For every 1% increase in males involved in agriculture as an accupation we expect a .17 decrease in fertility, if all other variables are held constant.
  - For every 1% increase in Catholisism we expect a .10 increase in fertility, if all other variables are held constant.
  - For every 1% increase in education we expect a .87 decrease in fertility, if all other variables are held constant. + Etc., etc....
- The asterisks indicate what level of significance that explanatory variable has on the dependent variable, fertility. For example the alpha level of the t-test for Agriculture has one \* as such it is significant at an alpha level of 0.05
- However, if only Agriculture is listed as the independent variable we will see the coefficient change to positive, indicating that sometimes additional variables can affect the influence of an independent variable on a dependent one.

```
summary(lm(Fertility ~ Agriculture, swiss))$coef
```

```

Estimate Std. Error t value Pr(>|t|)
(Intercept) 60.3043752 4.25125562 14.185074 3.216304e-18
Agriculture 0.1942017 0.07671176 2.531577 1.491720e-02

```

- One last note: Adding additional, repeated info to a lm won't change the result, for example...

```

extra <- (swiss$Education + swiss$Agriculture)
extraLM <- lm(Fertility ~. + extra, swiss)$coef
extraLM

```

```

(Intercept) Agriculture Examination Education
66.9151817 -0.1721140 -0.2580082 -0.8709401
Catholic Infant.Mortality extra
0.1041153 1.0770481 NA

```

```

lm(Fertility ~ ., swiss)$coef - extraLM

Warning in lm(Fertility ~ ., swiss)$coef - extraLM: longer object length is not
a multiple of shorter object length

(Intercept) Agriculture Examination Education
0 0 0 0
Catholic Infant.Mortality extra
0 0 NA

```

- The above code returns NA for extra because it gave no additional info to the linear model, and when subtracting all the coefficients we can see there is no difference between the original and lm with extra

## Quiz 2

- Given...

```

x <- c(0.61, 0.93, 0.83, 0.35, 0.54, 0.16, 0.91, 0.62, 0.62)
y <- c(0.67, 0.84, 0.6, 0.18, 0.85, 0.47, 1.1, 0.65, 0.36)

```

Give a P-value for the two sided hypothesis test of whether  $\beta_1$  from a linear regression model is 0 or not

```

results <- summary(lm(y ~ x))
results$coefficients

Estimate Std. Error t value Pr(>|t|)
(Intercept) 0.1884572 0.2061290 0.9142681 0.39098029
x 0.7224211 0.3106531 2.3254912 0.05296439

```

- 0.05296

- Consider the previous problem, give the estimate of the residual standard deviation

```
results$sigma
```

```
[1] 0.2229981
```

- In the mtcars data set, fit a linear regression model of weight (predictor) on mpg (outcome). Get a 95% confidence interval for the expected mpg at the average weight. What is the lower endpoint?

```

fit <- lm(mpg ~ wt, mtcars)
results <- predict(fit, newdata = data.frame(wt = mean(mtcars$wt)),
 interval = "confidence")
results

```

```
fit lwr upr
1 20.09062 18.99098 21.19027
```

- 18.991

4. Refer to the help file for `mtcars`. What is the weight coefficient interpreted as?  
 \* Expected change in mpg/1,000 lb increase in weight.
5. Consider again the `mtcars` data set and a linear regression model with mpg as predicted by weight (1,000 lbs). A new car is coming weighing 3000 pounds. Construct a 95% **prediction** interval for its mpg. What is the upper endpoint?

```
results <- predict(fit, newdata = data.frame(wt = 3),
 interval = "prediction")
results

fit lwr upr
1 21.25171 14.92987 27.57355
* 27.57
```

6. Consider the `mtcars` data set again with mpg as predicted by weight. A “short” ton is defined as 2,000 lbs. Construct a 95% confidence interval for the expected change in mpg per 1 short ton increase in weight. Give the lower endpoint.

```
fit <- lm(mpg ~ I(wt * 1000/2000), mtcars)
coefs <- summary(fit)$coef
inter <- coefs[2,1]
slope <- coefs[2,2]
slopeInterval <- inter + c(-1, 1) * qt(0.975, df = fit$df) * slope
slopeInterval

[1] -12.97262 -8.40527
```

7. If my X from a linear regression is measured in centimeters and I convert it to meters what would happen to the slope coefficient?

- Slope is *rise/run*, or *change in y/change in x* since we’re changing the units of x we have to multiple by the conversion factor  $100 \text{ cm}/1 \text{ m}$ , which is to multiple the coefficient by 100.
8. I have an outcome,  $Y$ , and a predictor,  $X$  and fit a linear regression model with  $Y = \beta_0 + \beta_1 X + \epsilon$  to obtain  $\hat{\beta}_0$  and  $\hat{\beta}_1$ . What would be the consequence to the subsequent slope and intercept if I were to refit the model with a new regressor,  $X + c$  for some constant,  $c$ ?
- $$Y = \beta_0 + \beta_1 X + \epsilon = \beta_0 - c\beta_1 + c\beta_1 + \beta_1 X + \epsilon$$

$$= \beta_0 - c\beta_1 + \beta_1(X + c) + \epsilon$$
+ As such this new intercept is  $\beta_0 - c\beta_1$
9. Refer back to the `mtcars` data set with mpg as an outcome and weight (wt) as the predictor. About what is the ratio of the the sum of the squared errors,  $\sum_{i=1}^n (Y_i - \hat{Y}_i)^2$  when comparing a model with just an intercept (denominator) to the model with the intercept and slope (numerator)?

```
results <- summary(lm(mpg~wt, mtcars))
1 - results$r.squared

[1] 0.2471672
```

10. Do the residuals always have to sum to 0 in linear regression?

- Yes, if an intercept is included in the residuals.

## Multivariable Regression, Residuals, & Diagnostics

### Multivariable Regression

#### Multivariable Regression Part 1

##### Intro Scenario Example

- If one were to present evidence of a relationship between breath mint usage (mints per day,  $X$ ) and pulmonary function (lung health measurement in FEV (Forced Expiratory Volume))
  - + You may be skeptical, for smokers tend to use more breath mints than non-smokers, and smoking is related to a loss in pulmonary function.
  - + To counteract this skepticism one may want to investigate if non-smoking breath mint users had lower lung function than non-smoking non-breath mint users and if smoking breath mint users had lower lung function than smoking non-breath mint users.
  - + In other words, to consider these results, we would have to demonstrate that the results hold when smoking status is held fixed.

### The Linear Model

- How can one generalize Simple Linear Regression (SLR) to incorporate lots of regressors for the purpose of predictions?
- Are there consequences for adding lots of regressors?
  - + Omitting variables can cause incorrect predictions
  - + Adding too many variables into a model will eventually give 0 residuals, even if adding in random garbage predictors. This is called **over-fitting**
- The general linear model extends SLR by adding terms linearly into the model  
$$Y_i = \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_p X_{pi} + \epsilon_i = \sum_{k=1}^p (X_{ik} \beta_j) + \epsilon_i$$
  - + Here  $X_{1i} = 1$  typically, so that an intercept is included.
- Least squares (and hence Middle of Least squares estimates under iid Gaussianity of the errors) minimizes:  
$$\sum_{i=1}^n (Y_i - \sum_{k=1}^p X_{ki} \beta_j)^2$$
  - + Note that the important linearity is linearity in the coefficients, thus:
  - $$Y_i = \beta_1 X_{1i}^2 + \beta_2 X_{2i}^2 + \dots + \beta_p X_{pi}^2 + \epsilon_i$$
  - + The above equation is still a linear model, we've just squared the element of the predictors

#### Multivariable Regression Part 2

##### How to Get Estimates

- Recall that the Least Squares estimate for regression through the origin,  $E[Y_i] = X_{1i}\beta_1$ , was  $\sum(X_iY_i)/\sum X_i^2$
- For two regressors we would have a regression line of  $E[Y_i] = X_{1i}\beta_1 + X_{2i}\beta_2 = \mu_i$  and least squares would minimize  $\sum_{i=1}^n(Y_i - X_{1i}\beta_1 - X_{2i}\beta_2)^2$
- With two independent variables,  $X_1, X_2$ , being used in a lm  $X_1$  has to then be adjusted for  $X_2$  being involved
- So for the equation:  $E[Y_i] = X_{1i}\hat{\beta}_1 + X_{2i}\hat{\beta}_2$   

$$+ \hat{\beta}_1 = \frac{\sum_{i=1}^n(e_{i,Y|X_2}*e_{i,X_1|X_2})}{\sum_{i=1}^n e_{i,X_1|X_2}^2}$$
  - Remember that  $e_i$  is just a col. vector of 0's except at the  $i^{th}$  position, where there is a 1.
  - + Essentially the regression estimate for  $\beta_1$  is the regression through the origin estimate having regressed  $X_2$  out of both the response and the predictor.
  - Similarly, the regression estimate for  $\beta_2$  is the regression through the orgin estimate having regressed  $X_1$  out of both the response and the predictor.
  - + Multivariate regression estimates are exactly those having removed the linear relationship of the other variables from both the regressor and response.
- $Y_i = \beta_1 X_{1i} + \beta_2 X_{2i}$  where  $X_{2i} = 1$  is an intercept term.
- Notice the fitted coefficient of  $X_{2i}$  on  $Y_i$  is  $\bar{Y}$ 
  - + The residuals  $e_{i,Y|X_2} = Y_i - \bar{Y}$
- Notice the fitted coefficient of  $X_{2i}$  on  $X_{1i}$  is  $\bar{X}_1$ 
  - + The residuals  $e_{i,X_1|X_2} = X_{1i} - \bar{X}_1$
- Thus  

$$\begin{aligned} \hat{\beta}_1 &= \frac{\sum_{i=1}^n(e_{i,Y|X_2}*e_{i,X_1|X_2})}{\sum_{i=1}^n e_{i,X_1|X_2}^2} \\ &= \frac{\sum_{i=1}^n(X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n(X_i - \bar{X})^2} \\ &= Cor(X, Y) \frac{Sd(Y)}{Sd(X)} \end{aligned}$$
- For multiple variables this case just keeps adding terms, the least square solutions have to minimize:  

$$\sum_{i=1}^n(Y_i - X_{1i}\beta_1 - \dots - X_{pi}\beta_p)^2$$
  - + The least squares estimate for the coefficeint of a multivariate regression model is exactly regression through the orgin with the linear relationships with the other regressors removed from both the regressor and outcome by taking residuals.
  - + In this sense, multivariate regression “adjusts” a coefficient for the linear impact of the other variables.

## Multivariable Regression Continued

### Simulation

```

set.seed(1618033)
n <- 100
x1 <- rnorm(n)
x2 <- rnorm(n)
x3 <- rnorm(n)
y <- 1 + x1 + x2 + x3 +
 rnorm(n, sd = .1) #rand. noise as \epsilon term
ey <- resid(lm(y ~ x2 + x3))
ex <- resid(lm(x1 ~ x2 + x3))
sum(ey * ex) / sum(ex ^ 2) #Reg. through origin estimate

[1] 0.9889948

coef(lm(ey ~ ex - 1)) #Reg. est of ey to ex

ex
0.9889948

coef(lm(y ~ x1 + x2 + x3)) #Notice x1 is same as above

(Intercept) x1 x2 x3
1.0195807 0.9889948 1.0048448 1.0178638

```

### Interpretation of the Coeficients

- Our model is  $E[Y|X_1 = x_1, \dots, X_p = x_p] = \sum_{k=1}^p x_k \beta_k$   
+ Where  $\beta_k$  is the coefficient for each  $x_k$
- The interpretation of a multivarate regression coefficient is the expected change in the response per unit change in the regressor, holding all of the other regressors fixed. As such the diffrence between adding 1 to a regressor<sub>i</sub> and the orginal equation is just  $\beta_i$ , as seen below  
+ Adding 1:  $E[Y|X_1 = x_1 + 1, \dots, X_p = x_p] = (x_1 + 1)\beta_1 + \sum_{k=2}^p x_k \beta_k$   
+ Difference:  $(x_1 + 1)\beta_1 + \sum_{k=2}^p x_k \beta_k - \sum_{k=1}^p x_k \beta_k = \beta_1$

### Fitted Values, Residuals and Residual Variation

ALL of our Simple Linear Regression (SLR) quantites can be extended to linear models of multiple dimensions

- \* Model:  $Y_i = \sum_{k=1}^p X_{ik} \beta_k + \epsilon_i$  where  $\epsilon_i \sim N(0, \sigma^2)$
- \* Fitted Responses:  $\hat{Y}_i = \sum_{k=1}^p X_{ik} \hat{\beta}_k$
- \* Residuals:  $e_i = Y_i - \hat{Y}_i$
- \* Variance ewstiamte  $\hat{\sigma}^2 = \frac{1}{n-p} \sum_{i=1}^n e_i^2$
- \* To get predicted responses at new values,  $x_1, \dots, x_p$  simply plug them into the linear model  $\sum_{k=1}^p x_k \hat{\beta}_k$
- \* Each coefficient has their own standard error,  $\hat{\sigma}_{\hat{\beta}_k}$ , and as such  $\frac{\hat{\beta}_k - \beta_k}{\hat{\sigma}_{\hat{\beta}_k}}$  follows a  $T$  distribution with

$n - p$  degrees of freedom

\* Predicted responses have standard errors and we can calculate predicted and expected response intervals.

## Linear Models Summary

- Linear Models are the single most important applied statistical and machine learning technique, *by far*.
- They can... + decompose a signal into its harmonics  
+ flexibly fit complicated functions & curves.  
+ fit factor variables as predictors  
+ uncover complex multivariate relationships with the response  
+ build accurate prediction models

## Multivariable Regression Tips and Tricks

### Multivariable Regression Examples Part 1

```
library(datasets); data(swiss)
```

- `swiss` is a data frame with 47 observations on 6 variables
  - 1) *Fertility* - lg, a “common standardized fertility measure”
  - 2) *Agriculture* - % of males involved in agriculture as occupation
  - 3) *Examination* - % of draftees receiving highest mark on army examination
  - 4) *Education* - % education beyond primary school for draftees
  - 5) *Catholic* - % catholic (as opposed to protestant)
  - 6) *Infant.Mortality* - % of live births who live less than 1 year

- All variables but *Fertility* give proportions of the population

- These data are from Switzerland in 1888 from 47 French-speaking “provinces”

```
all <- summary(lm(Fertility ~ . , data = swiss))
all$coefficients
```

|                | Estimate   | Std. Error  | t value   | Pr(> t )     |
|----------------|------------|-------------|-----------|--------------|
| ## (Intercept) | 66.9151817 | 10.70603759 | 6.250229  | 1.906051e-07 |
| ## Agriculture | -0.1721140 | 0.07030392  | -2.448142 | 1.872715e-02 |
| ## Examination | -0.2580082 | 0.25387820  | -1.016268 | 3.154617e-01 |

```

Education -0.8709401 0.18302860 -4.758492 2.430605e-05
Catholic 0.1041153 0.03525785 2.952969 5.190079e-03
Infant.Mortality 1.0770481 0.38171965 2.821568 7.335715e-03

```

- Each estimate would be interpreted as: Our model estimates an expected Estimate Increase/Decrease in standardized fertility for every 1% increase in percentage of Explanatory Variable
- The Std. Error describes how precise the Estimate is
- The t-test for  $H_0 : \beta_{Agri} = 0$  versus  $H_a : \beta_{Agri} \neq 0$  is significant.
  - R gives the t value for this test and the  $\Pr(>|t|)$ , P-value
- Having only one predictor will change to coefficients

```

onlyAgri <- summary(lm(Fertility ~ Agriculture, data = swiss))
onlyAgri$coefficients

```

```

Estimate Std. Error t value Pr(>|t|)
(Intercept) 60.3043752 4.25125562 14.185074 3.216304e-18
Agriculture 0.1942017 0.07671176 2.531577 1.491720e-02
all$coefficients[1:2,]

```

```

Estimate Std. Error t value Pr(>|t|)
(Intercept) 66.915182 10.70603759 6.250229 1.906051e-07
Agriculture -0.172114 0.07030392 -2.448142 1.872715e-02

```

- This difference when including new factors is a version of Simpson's paradox

## Simulation

```

set.seed(1618033)
n <- 100
x2 <- 1:n
x1 <- 0.01 * x2 + runif(n, -0.1, 0.1)
y <- -x1 + x2 + rnorm(n, sd = 0.01)
summary(lm(y ~ x1))$coef

```

```

Estimate Std. Error t value Pr(>|t|)
(Intercept) 0.6769049 1.188377 0.5696045 5.702494e-01
x1 96.1246184 2.014714 47.7112934 1.208965e-69

```

- Only looking at x1 doesn't show the underneath pattern generated from x2 and gives a very large slope, however adding in x2 ...

```

summary(lm(y ~ x1 + x2))$coef

```

```

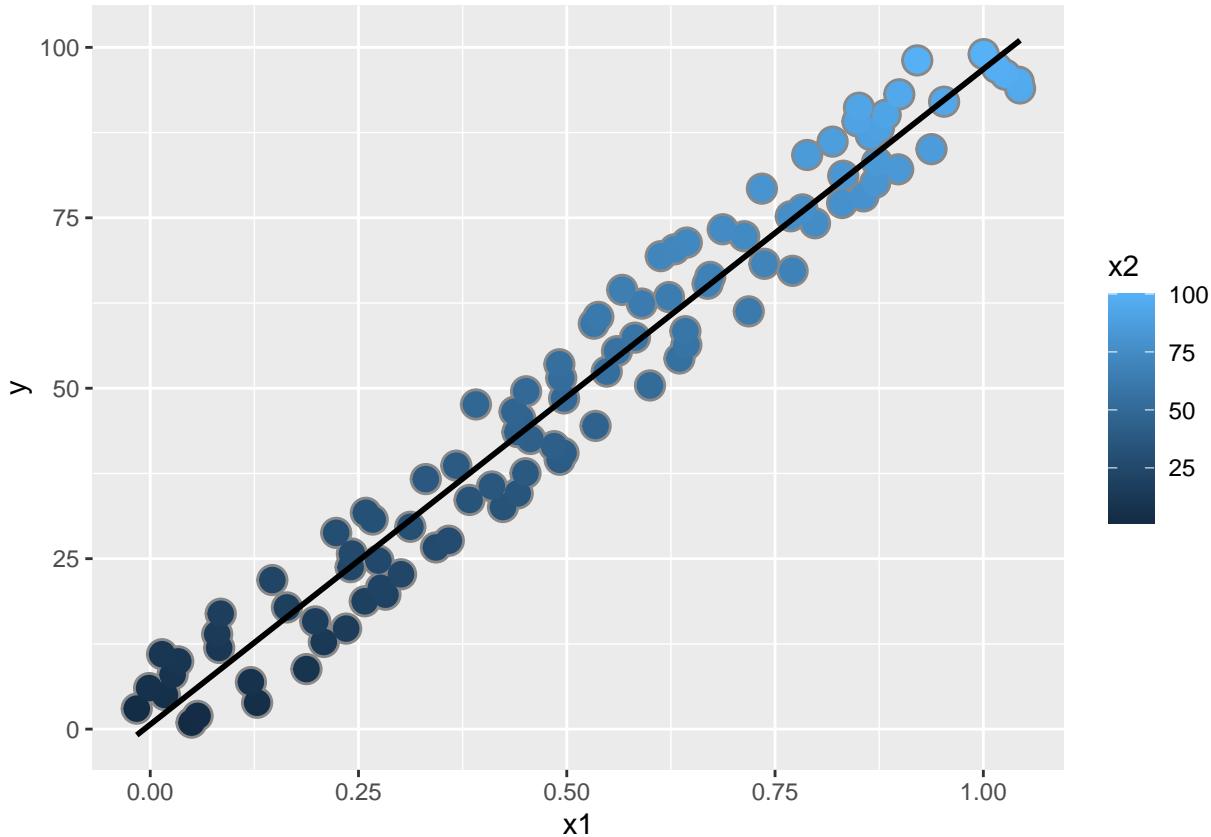
Estimate Std. Error t value Pr(>|t|)
(Intercept) 0.0003514301 0.0019009607 0.1848697 8.537172e-01
x1 -1.0257893318 0.0159992072 -64.1150103 3.195073e-81
x2 1.0002757089 0.0001613649 6198.8421136 2.580392e-273

```

- ... gives a more accurate representation of the true coefficients for  $y$ .

- Let's do some plotting:

```
library(tidyverse)
dat <- data.frame(y = y, x1 = x1, x2 = x2,
ey = resid(lm(y ~ x2)), ex1 = resid(lm(x1 ~ x2)))
plot <- ggplot(dat, aes(x1, y, colour = x2)) +
 geom_point(colour = "#888888", size = 5) +
 geom_point(size = 4) +
 geom_smooth(method = lm, se = FALSE, colour = "#000000")
plot
```

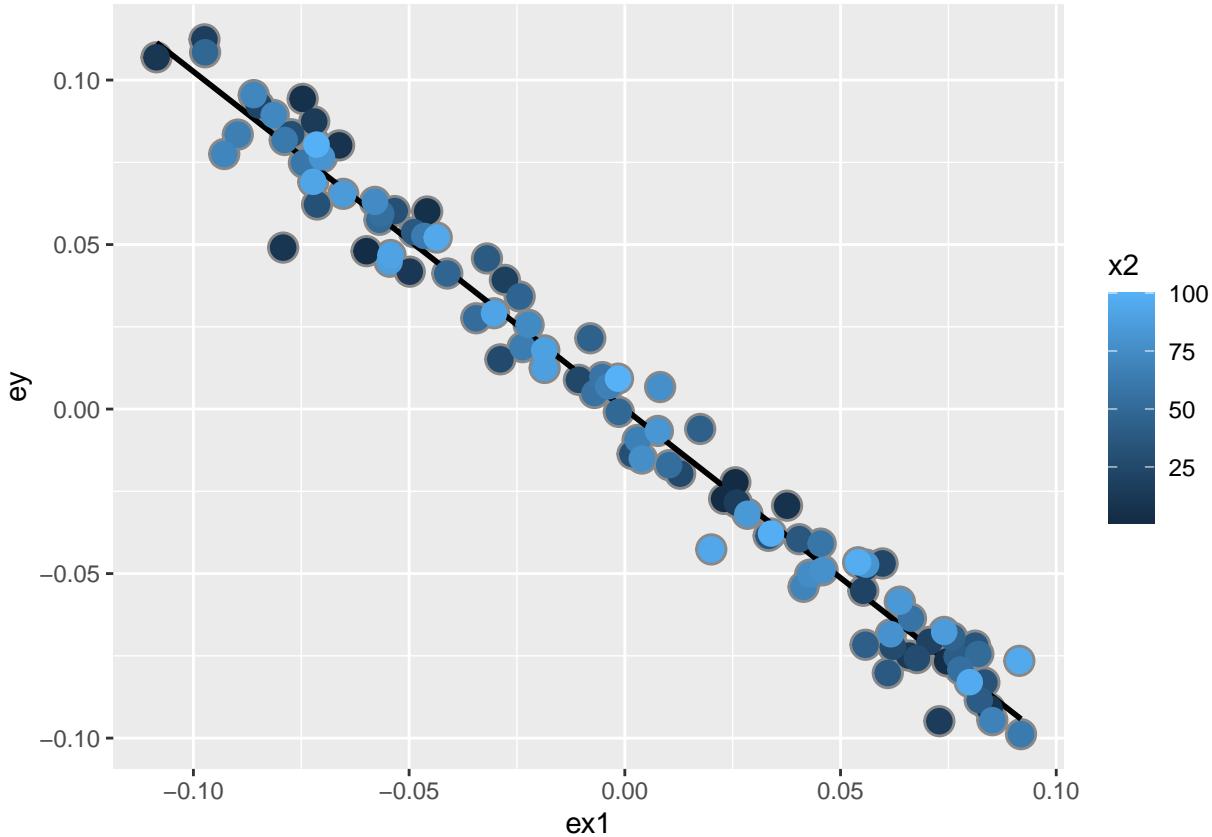


- Although  $y$  grows with  $x_1$  it can also be seen that  $x_2$  grows as  $x_1$  grows

- Now let's look at the residuals of  $y \sim x_2$  versus  $x_1 \sim x_2$

```
residplot <- ggplot(dat, aes(ex1, ey, colour = x2)) +
 geom_point(colour = "#888888", size = 5) +
 geom_smooth(method = lm, se = FALSE, colour = "#000000") +
 geom_point(size = 4)
residplot

'geom_smooth()' using formula 'y ~ x'
```



- This smoother has the slope of the estimate for  $x_1$  above

### Back to swiss Data Set

- The sign reverses itself with the inclusion of Examination and Education
- The percent of males in the province working in agriculture is negatively related to educational attainment (correlation of -0.6395) and Education and Examination (correlation of 0.6984) are obviously measuring similar things.
  - Is the positive marginal an artifact for not having accounted for the Education level? (Education does have a strong effect, btw)
- At the minimum, anyone claiming that provinces that focus more on agriculture have higher fertility rates would immediately be open to criticism due to these other cofactors present in the data.

### Multivariable Regression Examples Part 2

- Consider the linear model
 
$$Y_i = \beta_0 + X_{i1}\beta_1 + \epsilon_i$$
  - Where each  $X_{i1}$  is binary such that the value is 1 if measurement,  $i$ , is in a group and 0 otherwise.

- Then for people in the group,  $E[Y_i] = \beta_0 + \beta_1$
- And for people not in the group,  $E[Y_i] = \beta_0$ 
  - As such the mean of the treated group is  $\hat{\beta}_0 + \hat{\beta}_1$
  - Likewise, the mean for the control group is just  $\hat{\beta}_0$
- $\beta_1$  is interpreted as the increase or decrease in the mean of the treated group
  - This also gives the inference for the two groups, the same value a 2-sample t-test would give you

### Comparing More Than 2 Levels

- Since the value of  $\beta_1$  gives the value of a 2-sample t-test, you can extend this to perform t-tests across multiple variables
- Consider a multilevel factor level. For example a three level factor such as the US political party affiliation: Republican, Democrat, Independent.
- $Y_i = \beta_0 + X_{i1}\beta_1 + X_{i2}\beta_2 + \epsilon_i$ 
  - $X_{i1}$  is 1 for Republicans and 0 otherwise
  - $X_{i2}$  is 1 for Democrats and 0 otherwise
- If  $i$  is ... then  $E[Y_i] = \dots$ 
  - **Republican** =  $\beta_0 + \beta_1$
  - **Democrat** =  $\beta_0 + \beta_2$
  - **Independent** =  $\beta_0$
- $\beta_1$  compares Republicans to Independents
- $\beta_2$  compares Democrats to Independents
- $\beta_1 - \beta_2$  compares Republicans to Democrats
- (Choice of reference category changes the interpretation.)

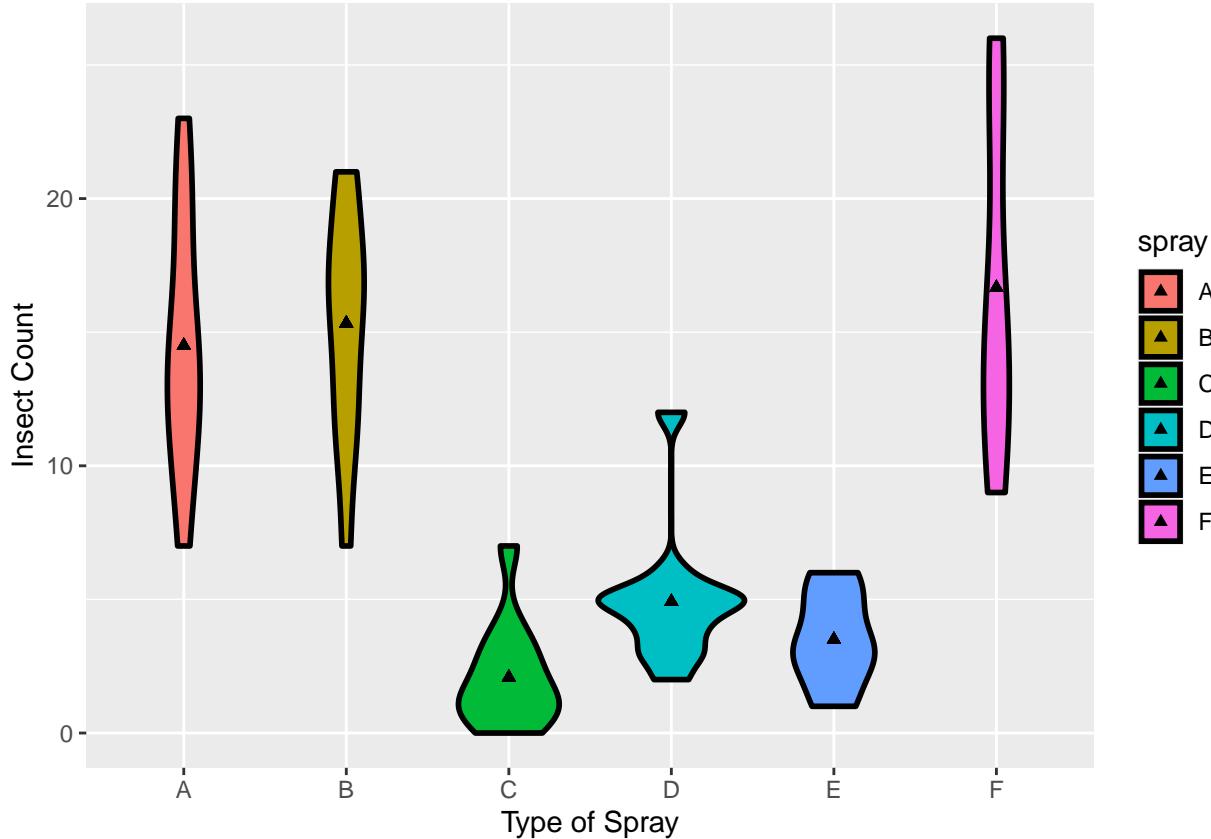
### Example in R with InsectSprays

```
data(InsectSprays); library(tidyverse)
mus <- InsectSprays %>% group_by(spray) %>% summarise(mean = mean(count))
dat <- merge(InsectSprays, mus, by = "spray")
plot <- ggplot(dat, aes(spray, count, fill = spray)) +
 geom_violin(colour = "#000000", size = 1) +
```

```

 labs(x = "Type of Spray", y = "Insect Count") +
 geom_point(aes(y=mean), shape = 17)
plot

```



- Fitting count as a response to `spray`, displaying mean with a triangle

```

res <- lm(count ~ spray, InsectSprays)
summary(res)$coef

```

|                | Estimate    | Std. Error | t value    | Pr(> t )     |
|----------------|-------------|------------|------------|--------------|
| ## (Intercept) | 14.500000   | 1.132156   | 12.8074279 | 1.470512e-19 |
| ## sprayB      | 0.8333333   | 1.601110   | 0.5204724  | 6.044761e-01 |
| ## sprayC      | -12.4166667 | 1.601110   | -7.7550382 | 7.266893e-11 |
| ## sprayD      | -9.5833333  | 1.601110   | -5.9854322 | 9.816910e-08 |
| ## sprayE      | -11.0000000 | 1.601110   | -6.8702352 | 2.753922e-09 |
| ## sprayF      | 2.1666667   | 1.601110   | 1.3532281  | 1.805998e-01 |

- It can be seen that `sprayA` is missing, this is because all of the other sprays are in comparison to `sprayA`, so the estimate values are indicating the change of each spray relative to spray A
- The average count of `sprayA` is just the `Estimate` of the `(Intercept)`
- The average count of `sprayC` would be the `Estimate` of the `(Intercept)` plus the `Estimate` of `sprayC`, as seen below

```

c(summary(res)$coef[1,1], mus$mean[1])

[1] 14.5 14.5

c(summary(res)$coef[1,1] + summary(res)$coef[3,1], mus$mean[3])

[1] 2.083333 2.083333

```

## Hard Coding the Dummy Variables

```

summary(lm(count ~
 I(1 * (spray == 'B')) + I(1 * (spray == 'C')) +
 I(1 * (spray == 'D')) + I(1 * (spray == 'E')) +
 I(1 * (spray == 'F')),
 data = InsectSprays))$coef

Estimate Std. Error t value Pr(>|t|)

(Intercept) 14.5000000 1.132156 12.8074279 1.470512e-19
I(1 * (spray == "B")) 0.8333333 1.601110 0.5204724 6.044761e-01
I(1 * (spray == "C")) -12.4166667 1.601110 -7.7550382 7.266893e-11
I(1 * (spray == "D")) -9.5833333 1.601110 -5.9854322 9.816910e-08
I(1 * (spray == "E")) -11.0000000 1.601110 -6.8702352 2.753922e-09
I(1 * (spray == "F")) 2.1666667 1.601110 1.3532281 1.805998e-01

summary(res)$coef

```

|                | Estimate    | Std. Error | t value    | Pr(> t )     |
|----------------|-------------|------------|------------|--------------|
| ## (Intercept) | 14.500000   | 1.132156   | 12.8074279 | 1.470512e-19 |
| ## sprayB      | 0.8333333   | 1.601110   | 0.5204724  | 6.044761e-01 |
| ## sprayC      | -12.4166667 | 1.601110   | -7.7550382 | 7.266893e-11 |
| ## sprayD      | -9.5833333  | 1.601110   | -5.9854322 | 9.816910e-08 |
| ## sprayE      | -11.0000000 | 1.601110   | -6.8702352 | 2.753922e-09 |
| ## sprayF      | 2.1666667   | 1.601110   | 1.3532281  | 1.805998e-01 |

- As such we can change what the  $\beta_0$  spray is, aka the reference level

```

summary(lm(count ~
 I(1 * (spray == 'A')) + I(1 * (spray == 'B')) +
 I(1 * (spray == 'D')) + I(1 * (spray == 'E')) +
 I(1 * (spray == 'F')),
 data = InsectSprays))$coef

Estimate Std. Error t value Pr(>|t|)

(Intercept) 2.083333 1.132156 1.840148 7.024334e-02
I(1 * (spray == "A")) 12.416667 1.601110 7.755038 7.266893e-11
I(1 * (spray == "B")) 13.250000 1.601110 8.275511 8.509776e-12
I(1 * (spray == "D")) 2.833333 1.601110 1.769606 8.141205e-02
I(1 * (spray == "E")) 1.416667 1.601110 0.884803 3.794750e-01
I(1 * (spray == "F")) 14.583333 1.601110 9.108266 2.794343e-13

```

- Instead of typing all the factors one can just `relevel` the factors

```

spray_ <- relevel(InsectSprays$spray, "C")
summary(lm(count ~ spray_, data = InsectSprays))$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 2.083333 1.132156 1.840148 7.024334e-02
spray_A 12.416667 1.601110 7.755038 7.266893e-11
spray_B 13.250000 1.601110 8.275511 8.509776e-12
spray_D 2.833333 1.601110 1.769606 8.141205e-02
spray_E 1.416667 1.601110 0.884803 3.794750e-01
spray_F 14.583333 1.601110 9.108266 2.794343e-13

```

- Including all the parameters will just return NA for one of the parameters (In this case it would be for sprayA)
- We can remove the intercept to get a set of all the levels

```
summary(lm(count ~ 0 + spray, data = InsectSprays))$coef
```

```

Estimate Std. Error t value Pr(>|t|)
sprayA 14.500000 1.132156 12.807428 1.470512e-19
sprayB 15.333333 1.132156 13.543487 1.001994e-20
sprayC 2.083333 1.132156 1.840148 7.024334e-02
sprayD 4.916667 1.132156 4.342749 4.953047e-05
sprayE 3.500000 1.132156 3.091448 2.916794e-03
sprayF 16.666667 1.132156 14.721181 1.573471e-22

```

## Summary

- If we treat spray as a factor, R includes an intercept and omits the alphabetically first level of the factor.
  - All t-tests are for comparisons of Sprays versus Spray A
  - Empirical mean for A is the (Intercept)
  - Other group means are the intercept plus their coefficient
- If we omit an intercept, then it includes terms for all levels of the factor.
  - Group means are the coefficients
  - Tests are tests of whether the groups are different than zero.
- If we want comparisons between, say between Spray B and C, we could refit the model with C (or B) as the reference level.

Additional Tid-bits on this data

\* Counts are bounded from below by 0, which violates the assumption of normality of the errors.  
 + Also there are counts near 0, so both the actual assumption and the intent of the assumption are violated.

\* Variance does not appear to be constant

\* Perhaps taking logs of the counts would help.

- + There are 0 counts, so perhaps  $\log(count + 1)$
- \* Poisson GLMs for fitting count data will be covered later in this course.
- \* Because of these issues our means are correct, but our inference would not be.

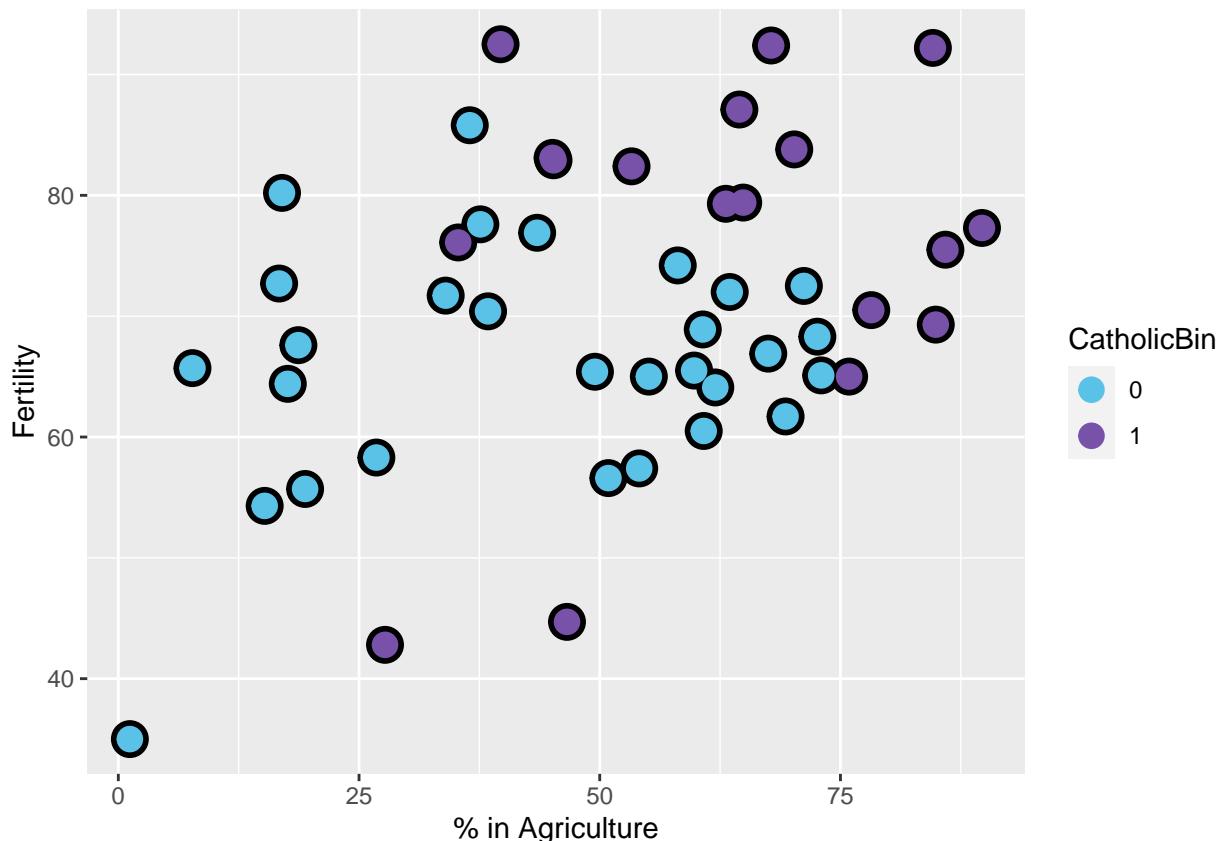
### Multivariable Regression Examples Part 3

We'll be looking at the `swiss` data set to see how we can translate a proportion to a binary variable

```
library(datasets); data(swiss)
##Create binary var from Catholic, 1 if > 50%, 0 o/w
dat <- swiss %>% mutate(CatholicBin = factor(1 * (Catholic > 50))) %>%
 select(Agriculture, Fertility, CatholicBin)
```

Now we can plot by subsets of variables

```
pal <- c("#5BC2E7", "#7851A9")
plot <- ggplot(dat, aes(Agriculture, Fertility, colour = CatholicBin)) +
 scale_colour_manual(values = pal) +
 geom_point(size = 6, colour = "#000000") + geom_point(size = 4) +
 labs(x = "% in Agriculture", y = "Fertility")
plot
```



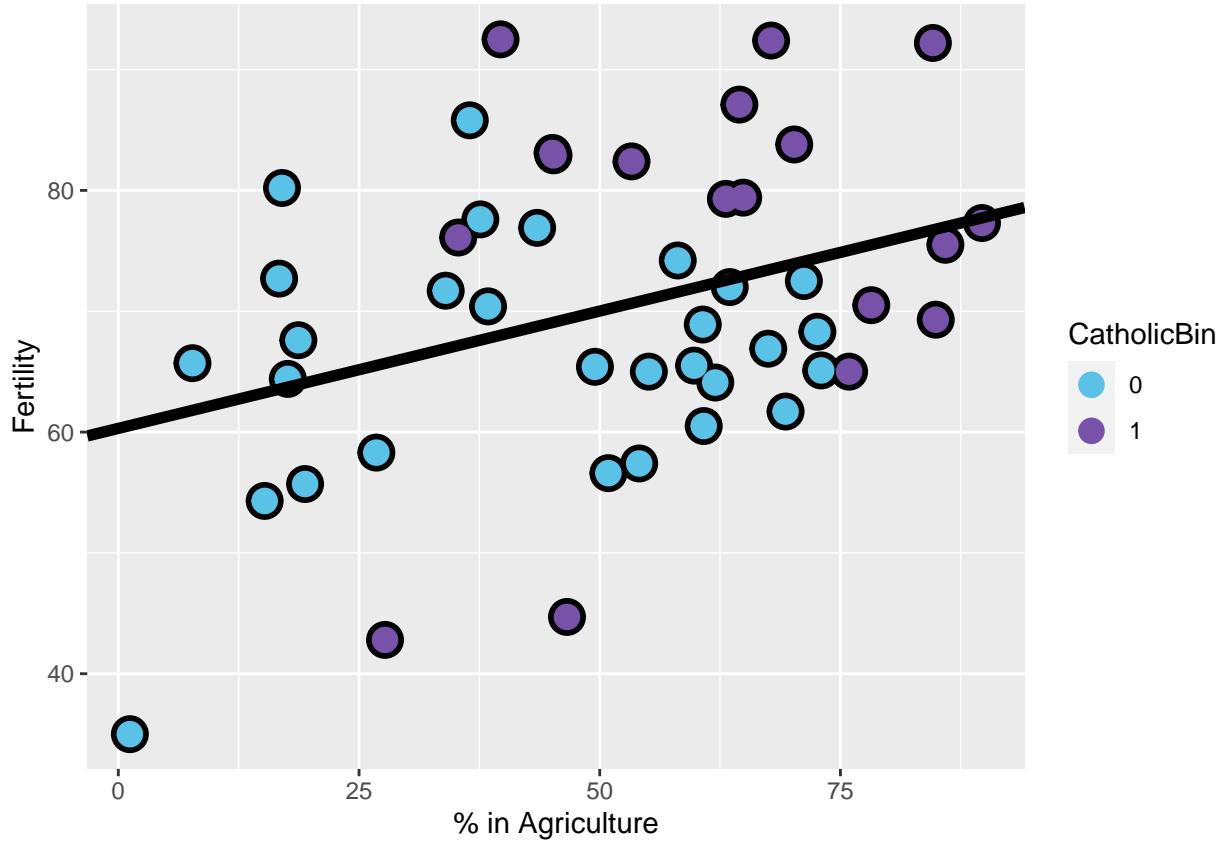
### Fitting the Model

- $y = Fertility$
- $x_1 = Agriculture$
- $x_2 = 1$  if over .50 Catholic, 0 otherwise  
Then our model could be:
- $E[y|x_1x_2] = \beta_0 + \beta_1x_1 + \beta_2x_2$   
Which would give us the expected values of:
- $E[y|x_2 = 0] = \beta_0 + \beta_1x_1$
- $E[y|x_2 = 1] = \beta_0 + \beta_2 + \beta_1x_1$   
These two models have the same slope,  $\beta_1$ , but different intercepts.  
To make a model that has different slopes we could fit it as such:
- $E[y|x_1x_2] = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_1x_2$   
We then have the following expected values:
- $E[y|x_2 = 0] = \beta_0 + \beta_1x_1$
- $E[y|x_2 = 1] = \beta_0 + \beta_1x_1 + \beta_2 + \beta_3x_1$   
 $= \beta_0 + \beta_2 + (\beta_1 + \beta_3)x_1$   
These models now have both different intercepts and slopes because we included an **interaction term**,  $\beta_3x_1x_2$

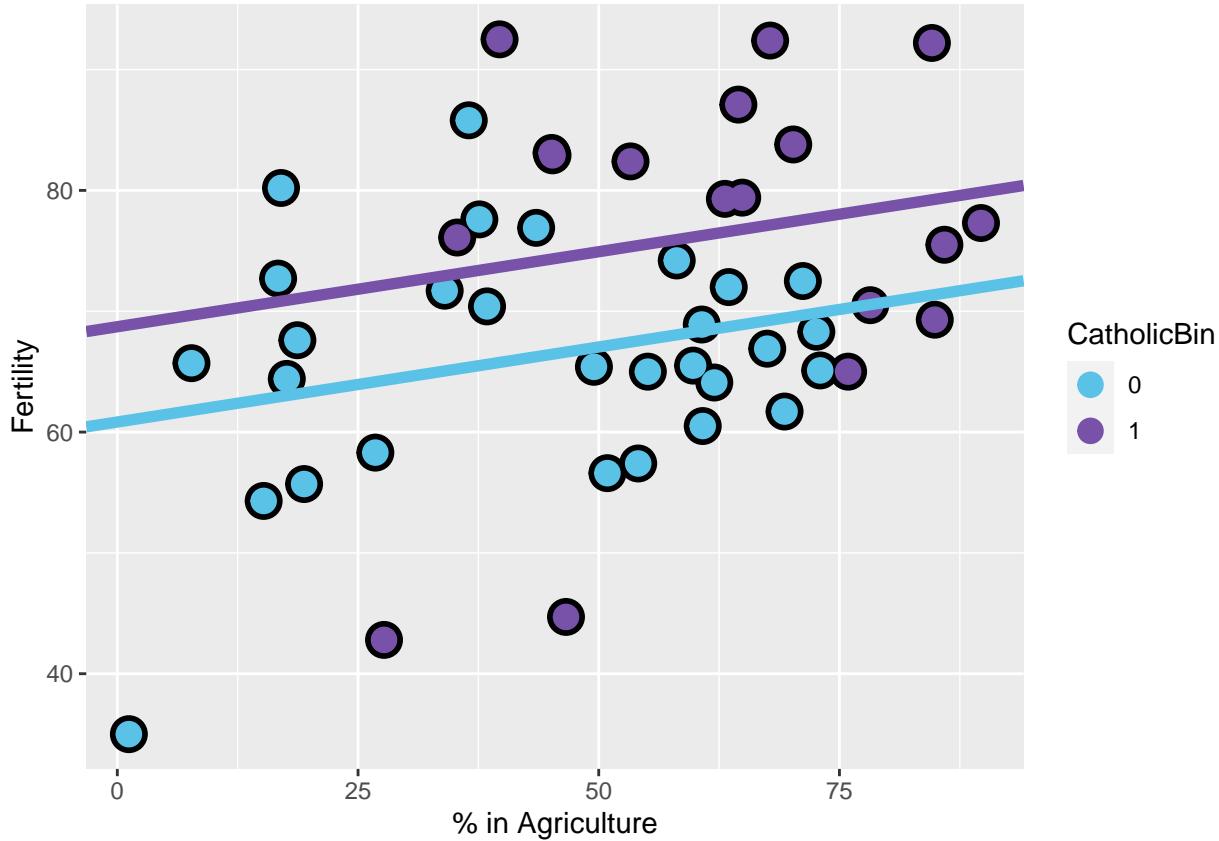
## Multivariable Regression Examples Part 4

We'll now look at an example of a model with an **interaction term**

```
##Fitting without CatholicBin
fafit <- lm(Fertility ~ Agriculture, data = dat)
plot + geom_abline(intercept = coef(fafit)[1],
 slope = coef(fafit)[2], size = 2)
```



```
##Fitting first proposed model
facfit <- lm(Fertility ~ Agriculture + CatholicBin, data = dat)
plot + geom_abline(intercept = coef(facfit)[1], slope = coef(facfit)[2],
 size = 2, colour = pal[1]) +
 geom_abline(intercept = coef(facfit)[1] + coef(facfit)[3],
 slope = coef(facfit)[2], size = 2, colour = pal[2])
```



```

##Fitting model with diffrent slopes
interactfit <- lm(Fertility ~ Agriculture * CatholicBin, dat)
##Let's look at the coef.s before tossing them all in
summary(interactfit)$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 62.04993019 4.78915566 12.9563402 1.919379e-16
Agriculture 0.09611572 0.09881204 0.9727127 3.361364e-01
CatholicBin1 2.85770359 10.62644275 0.2689238 7.892745e-01
Agriculture:CatholicBin1 0.08913512 0.17610660 0.5061430 6.153416e-01

coef(interactfit)

(Intercept) Agriculture CatholicBin1
62.04993019 0.09611572 2.85770359
Agriculture:CatholicBin1
0.08913512

```

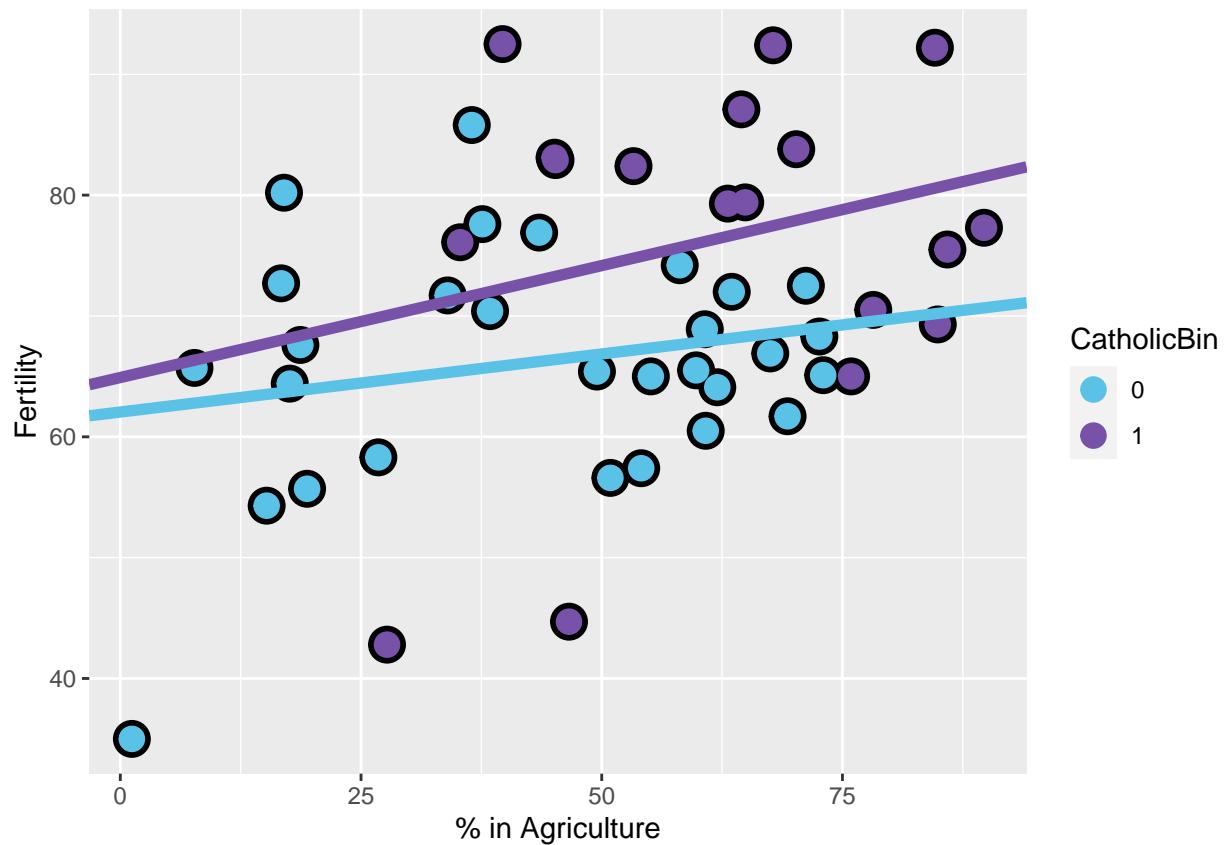
- We can see here that R has automatically made 3 variables:  $x_1$ , Agriculture;  $x_2$ , CatholicBin; and  $x_3$ , Agriculture\*CatholicBin.

```

plot + geom_abline(intercept = coef(interactfit)[1],
 slope = coef(interactfit)[2],
 size = 2, colour = pal[1]) +
 geom_abline(intercept = coef(interactfit)[1] + coef(interactfit)[3],#Int.Var

```

```
slope = coef(interactfit)[2] + coef(interactfit)[4],
size = 2, colour = pal[2])
```



### Lesson with swirl(): MultiVar Examples2

(No new content, review of InsectSprays example)

### Lesson with swirl(): MultiVar Examples3

This lesson looks at WHO data on hunger but essentially analyzes it the same as the swiss dataset.

## Adjustment

### Adjustment Examples

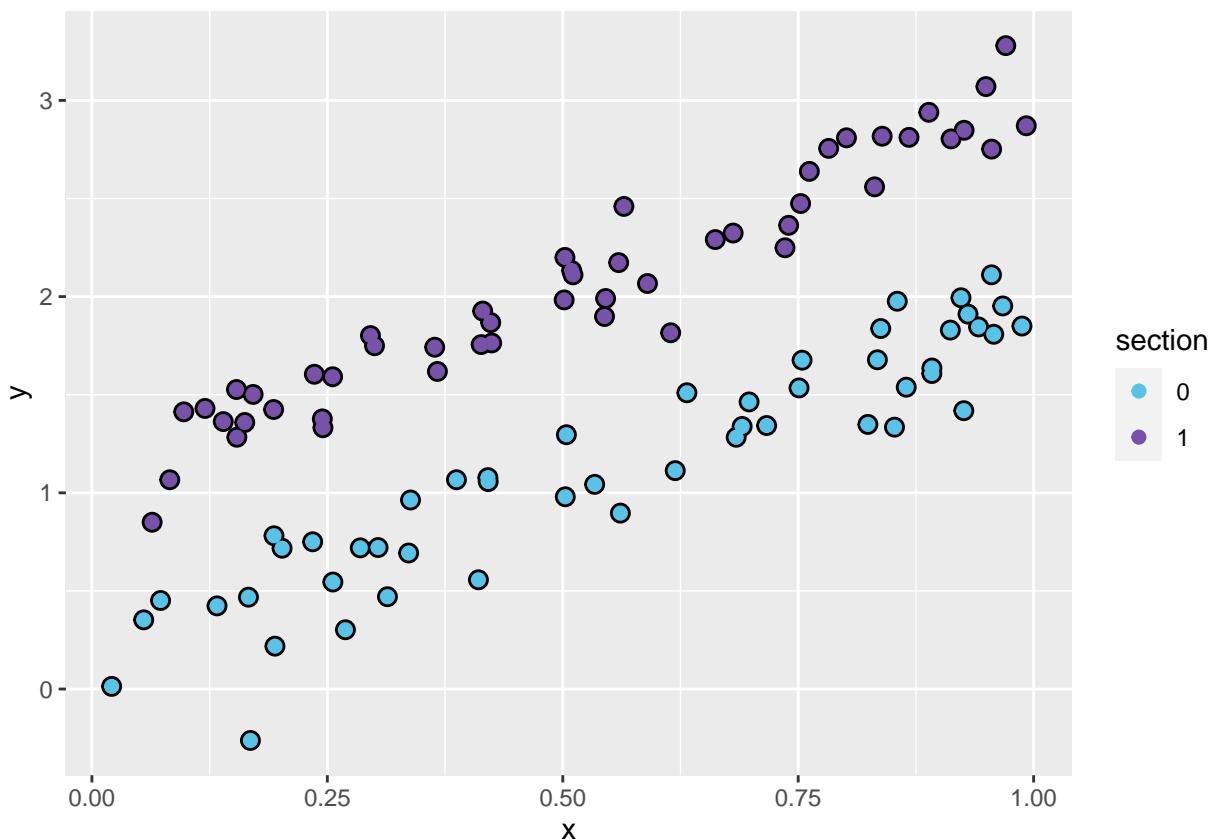
#### Simulation 1

```
set.seed(1618033)
n <- 100
t <- rep(c(0,1), each = n/2)
x <- c(runif(n/2), runif(n/2))
```

```

beta0 <- 0
beta1 <- 2
tau <- 1
sigma <- 0.2
y <- beta0 + x * beta1 + t * tau +
 rnorm(n, sd = sigma) #For noise
plot <- ggplot(data.frame(x = x, y = y, section = factor(t)),
 aes(x,y, colour = section)) +
 scale_colour_manual(values = pal) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2)
plot

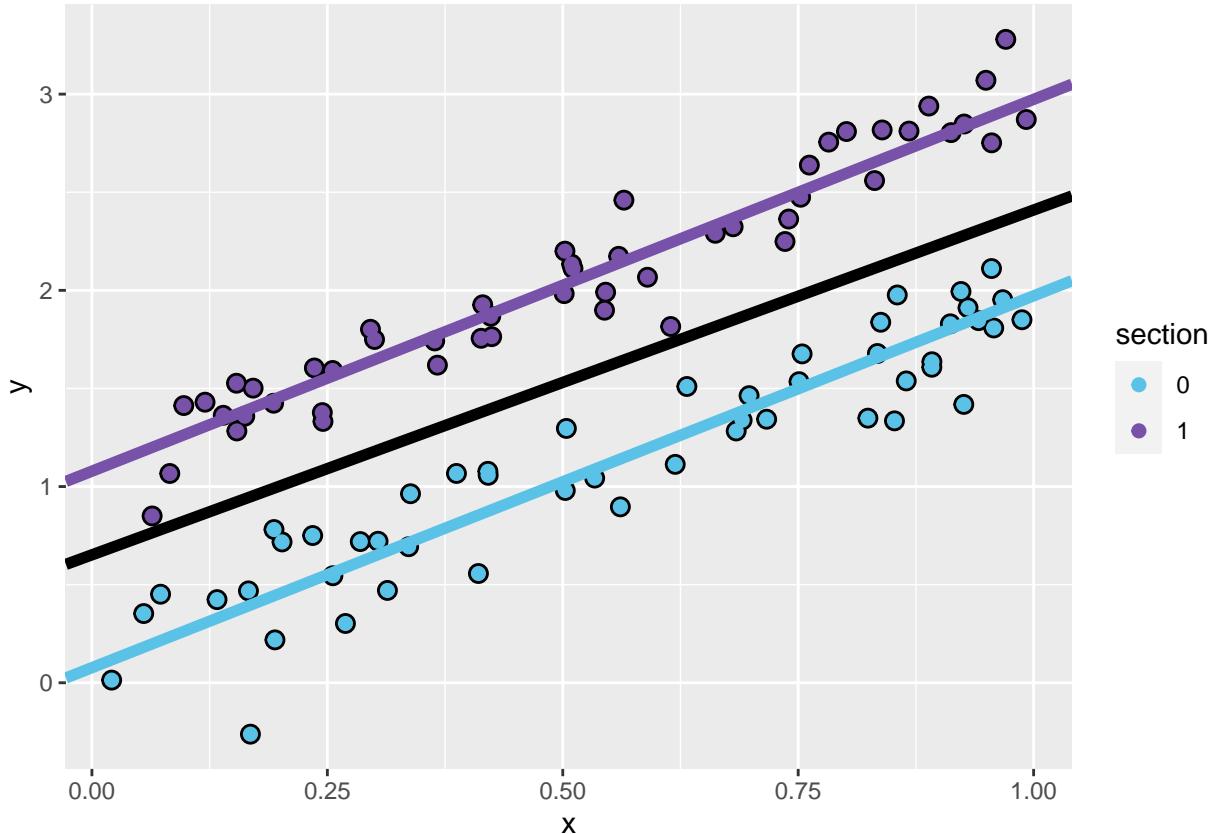
```



```

#Fiting a model to the points
xfit <- lm(y ~ x)
xtfit <- lm(y ~ x + t)
plot + geom_abline(intercept = xfit$coef[1], slope = xfit$coef[2],
 colour = "#000000", lwd = 2) + # All points, no t
 #With respect to t
 geom_abline(intercept = xtfit$coef[1], slope = xtfit$coef[2],
 colour = pal[1], lwd = 2) +
 geom_abline(intercept = xtfit$coef[1] + xtfit$coef[3], slope = xtfit$coef[2],
 colour = pal[2], lwd = 2)

```



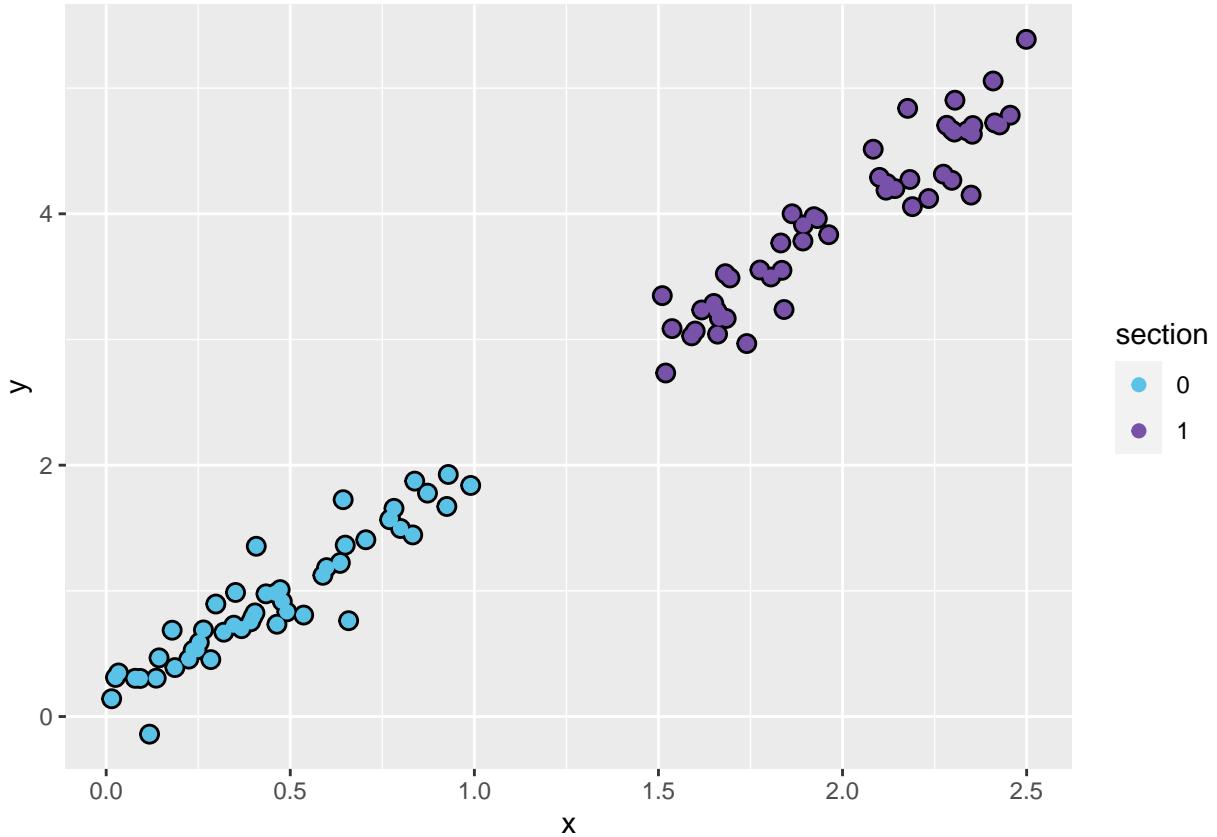
- The difference in the coloured and black lines is 0.1377 d
- Some things to note about this simulation
  - The X variable is unrelated to group status, t
  - The X variable is related to Y, but the intercept depends on group status
  - The group variable is related to Y
    - \* The relationship between group status and Y is constant depending on X
    - \* The relationship between group and Y disregarding X is about the same as holding X constant

## Simulation 2

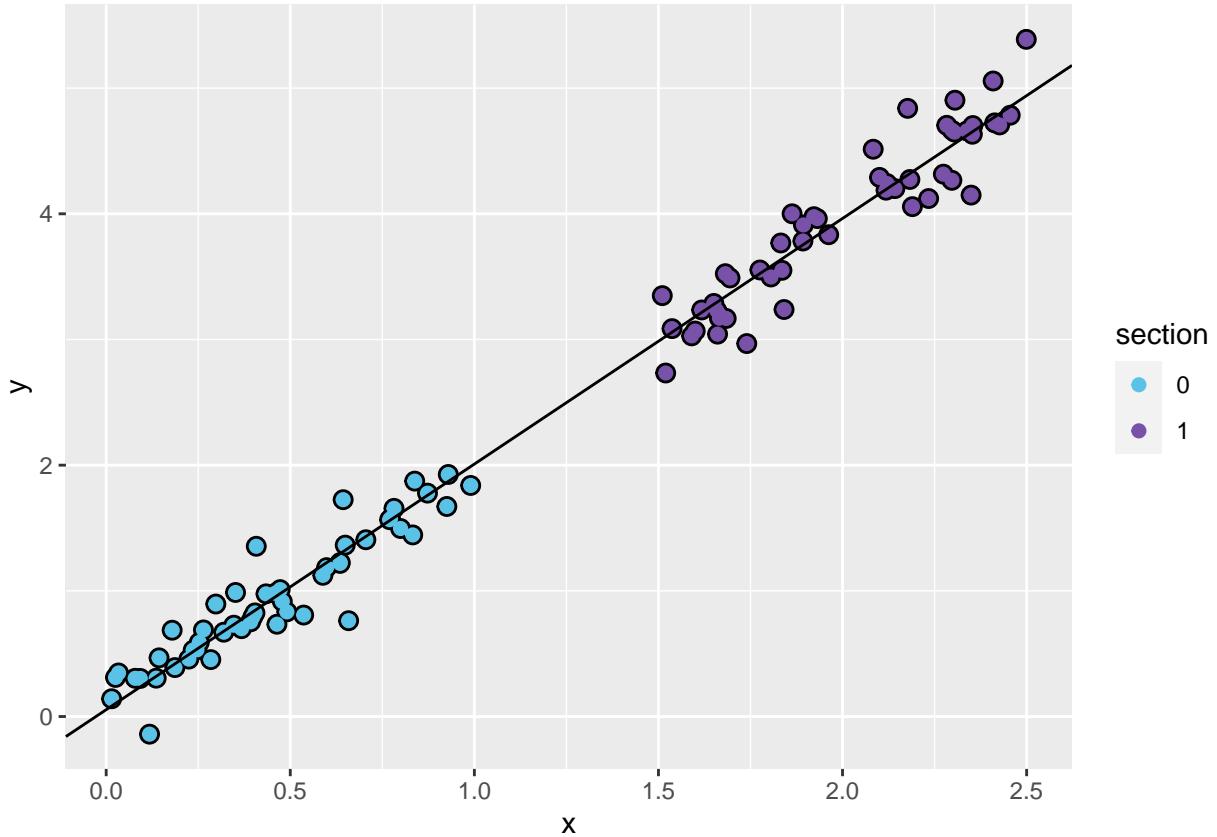
```

x <- c(runif(n/2), 1.5 + runif(n/2))
tau <- 0
y <- beta0 + x * beta1 + t * tau + rnorm(n, sd = sigma)
plot <- ggplot(data.frame(x = x, y = y, section = factor(t)),
 aes(x,y, colour = section)) +
 scale_colour_manual(values = pal) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2)
plot

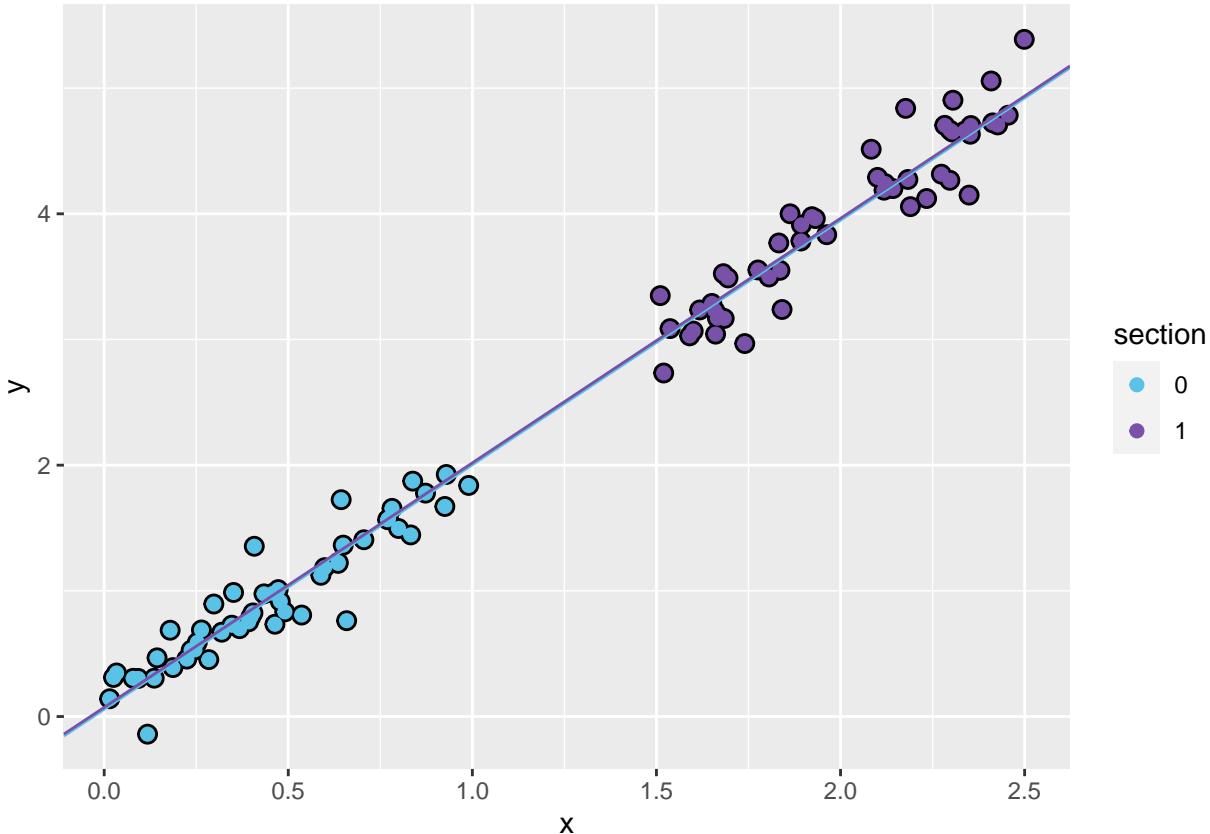
```



```
xfit <- lm(y ~ x)
xtfit <- lm(y ~ x + t)
plot + geom_abline(intercept = xfit$coef[1], slope = xfit$coef[2])
```



```
plot +
 geom_abline(intercept = xtfit$coef[1], slope = xtfit$coef[2],
 colour = pal[1]) +
 geom_abline(intercept = xtfit$coef[1] + xtfit$coef[3],
 slope = xtfit$coef[2], colour = pal[2])
```



```
xfit$coefficients
```

```
(Intercept) x
0.05432115 1.95436653
```

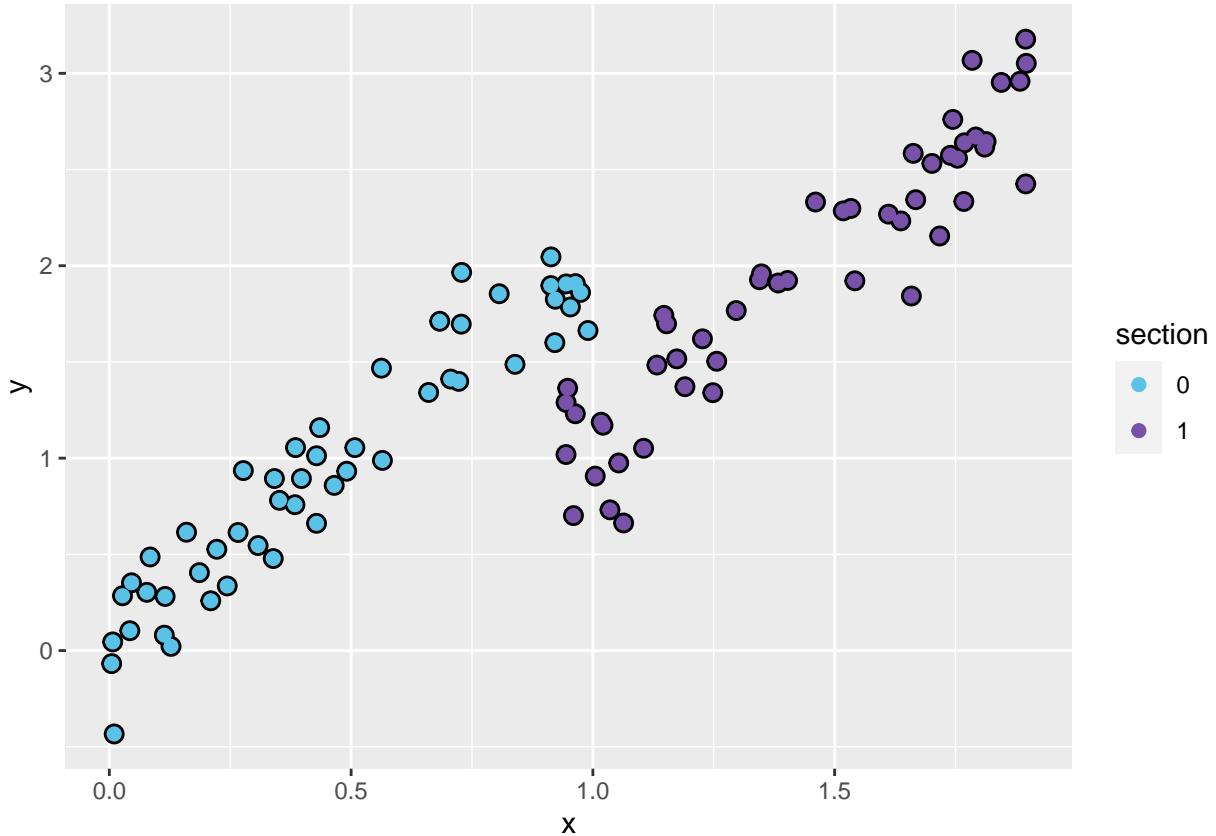
```
xtfit$coefficients
```

```
(Intercept) x t
0.05743077 1.94549763 0.01551007
```

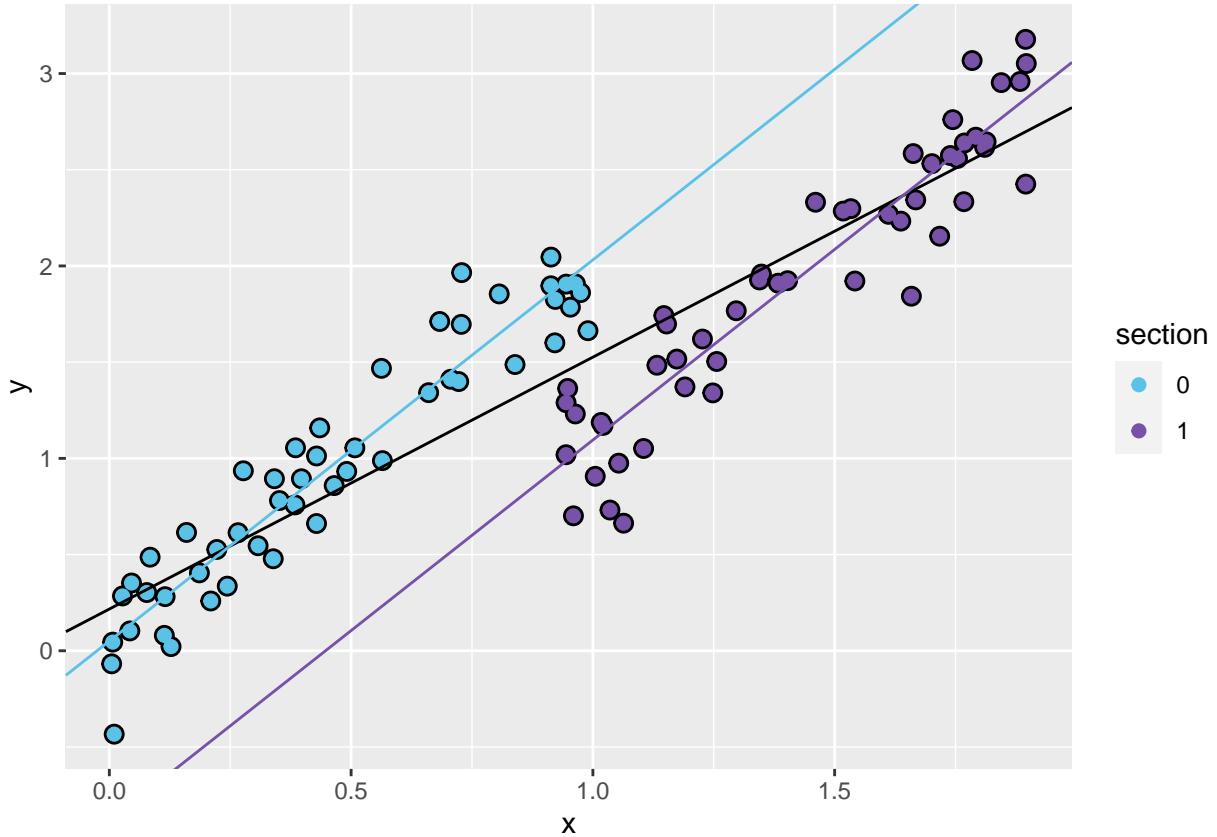
- These lines look nearly the same with only a slightly different intercept. However, the X variable is highly related to group status, below 1 is section 0 and above 1.5 is section 1.
- The X variable is realted to Y, and the intercept does not depend on the group varaible
  - The X variable remains related to Y holding group status constant
- The group variable is marginally related to Y disregarding X
- The model would estimate no adjusted effect due to group.
  - There isn't any data to inform the relationship between group and Y.
  - This conclusion is entirely based on the model
- This sort of model may occur if the x-variable is something related to being in the given section.

### Simulation 3

```
x <- c(runif(n/2), 0.9 + runif(n/2))
tau <- -1
y <- beta0 + x * betal + t * tau + rnorm(n, sd = sigma)
plot <- ggplot(data.frame(x = x, y = y, section = factor(t)),
 aes(x,y, colour = section)) +
 scale_colour_manual(values = pal) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2)
plot
```



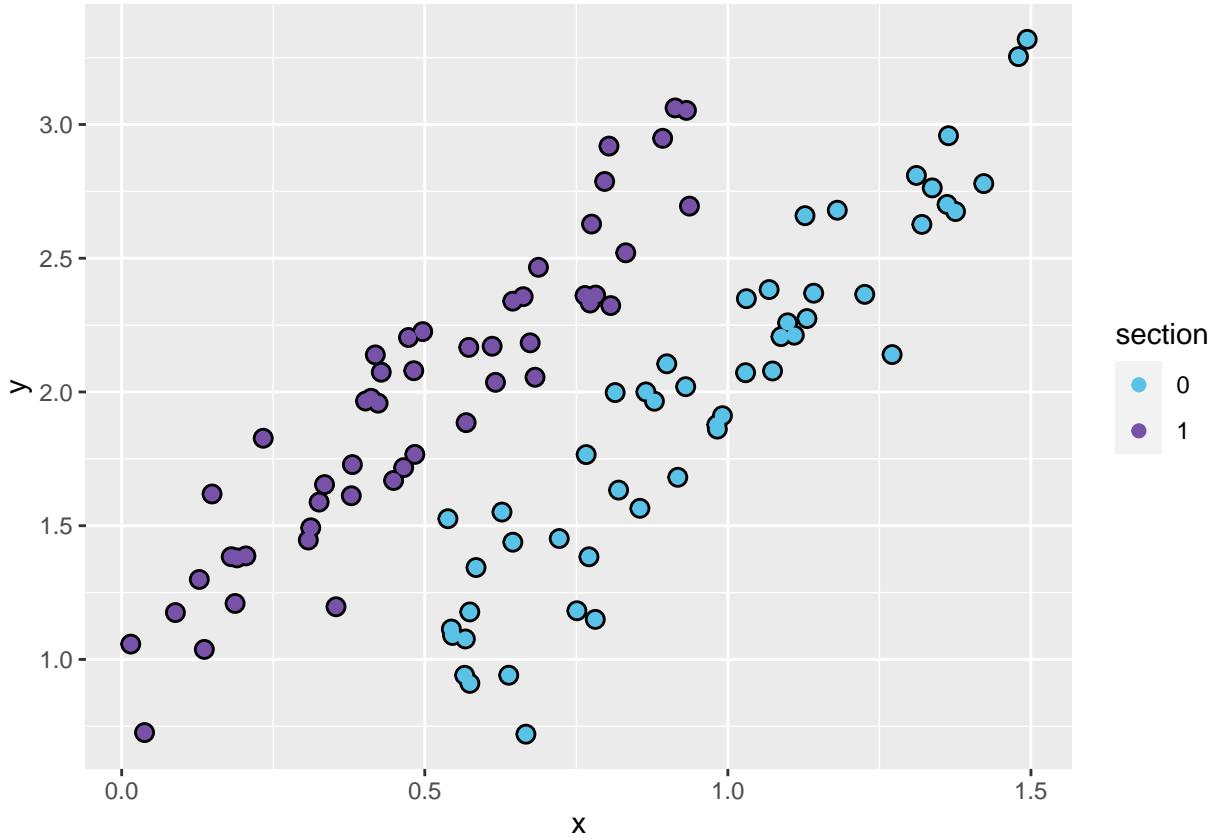
```
##Fitting a model
xfit <- lm(y ~ x)
xtfit <- lm(y ~ x + t)
plot + geom_abline(intercept = xfit$coef[1], slope = xfit$coef[2]) +
 geom_abline(intercept = xtfit$coef[1], slope = xtfit$coef[2],
 colour = pal[1]) +
 geom_abline(intercept = xtfit$coef[1] + xtfit$coef[3],
 slope = xtfit$coef[2], colour = pal[2])
```



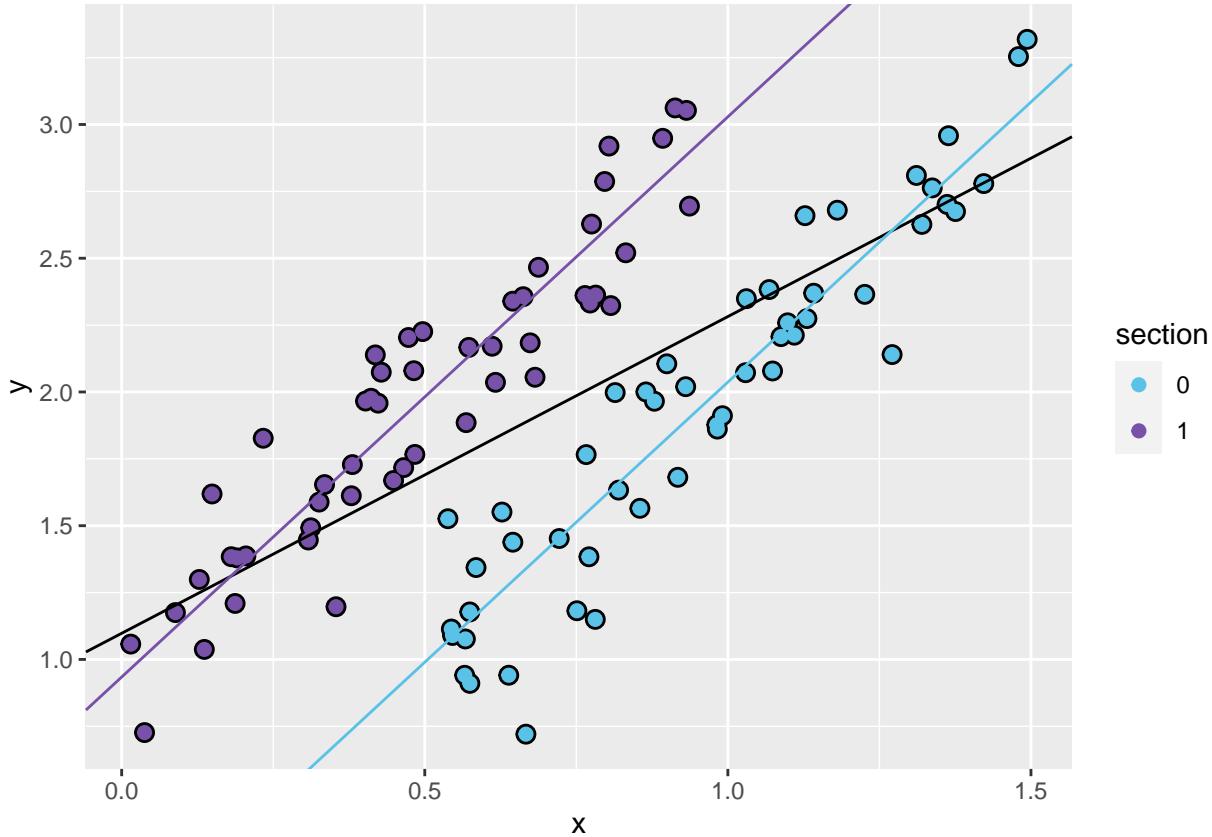
- Here we can see subsetting by group gives a clearly different slope and still has some shared x values between groups. They also have different intercepts
- This is an example of Simpson's "paradox"

#### Simulation 4

```
x <- c(0.5 + runif(n/2), runif(n/2))
tau <- 1
y <- beta0 + x * beta1 + t * tau + rnorm(n, sd = sigma)
plot <- ggplot(data.frame(x = x, y = y, section = factor(t)),
 aes(x,y, colour = section)) +
 scale_colour_manual(values = pal) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2)
plot
```



```
##Fitting a model
xfit <- lm(y ~ x)
xtfit <- lm(y ~ x + t)
plot + geom_abline(intercept = xfit$coef[1], slope = xfit$coef[2]) +
 geom_abline(intercept = xtfit$coef[1], slope = xtfit$coef[2],
 colour = pal[1]) +
 geom_abline(intercept = xtfit$coef[1] + xtfit$coef[3],
 slope = xtfit$coef[2], colour = pal[2])
```



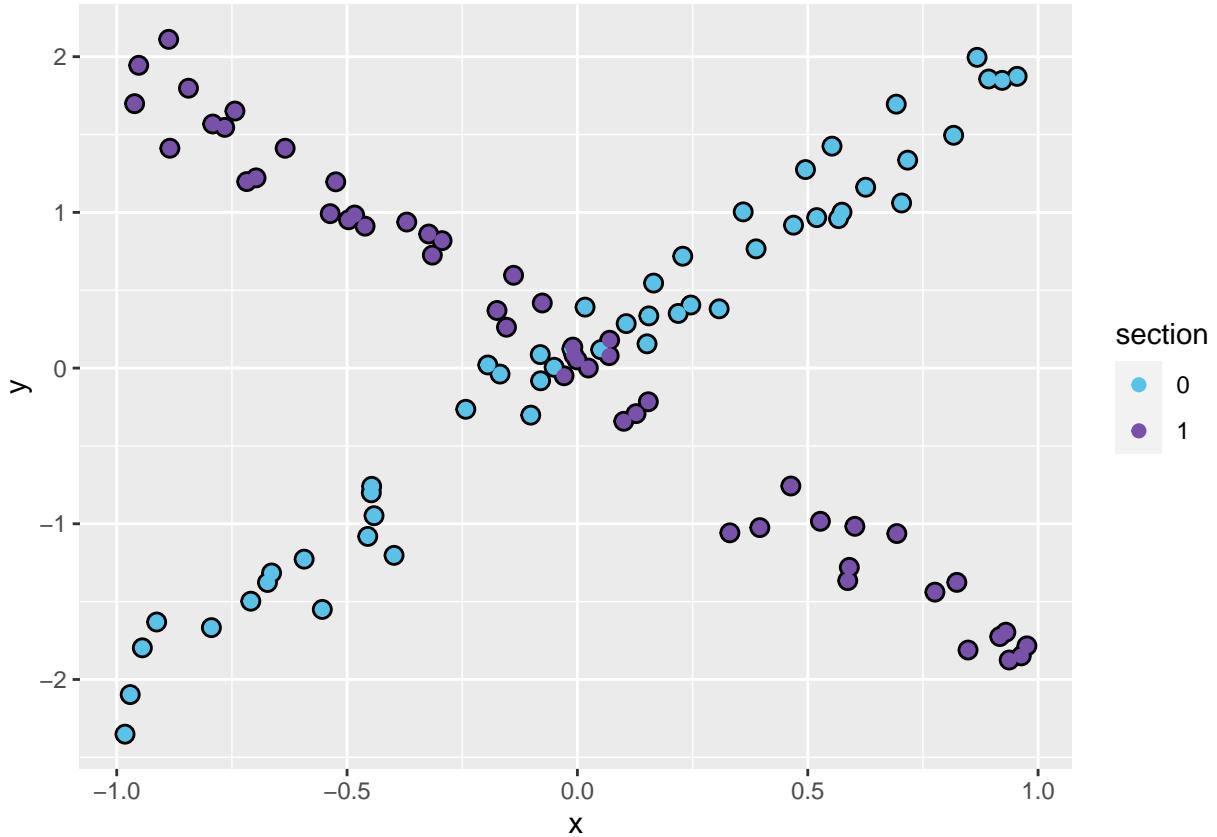
- In this example there is a large affect when adjusted for  $x$ , however the marginal affect is not great with  $\text{mean}(\text{section}==0) = 1.9586$  and  $\text{mean}(\text{section}==1) = 1.8103$ . Group status is not realted to  $X$  much with all the overlap.

### Simulation 5

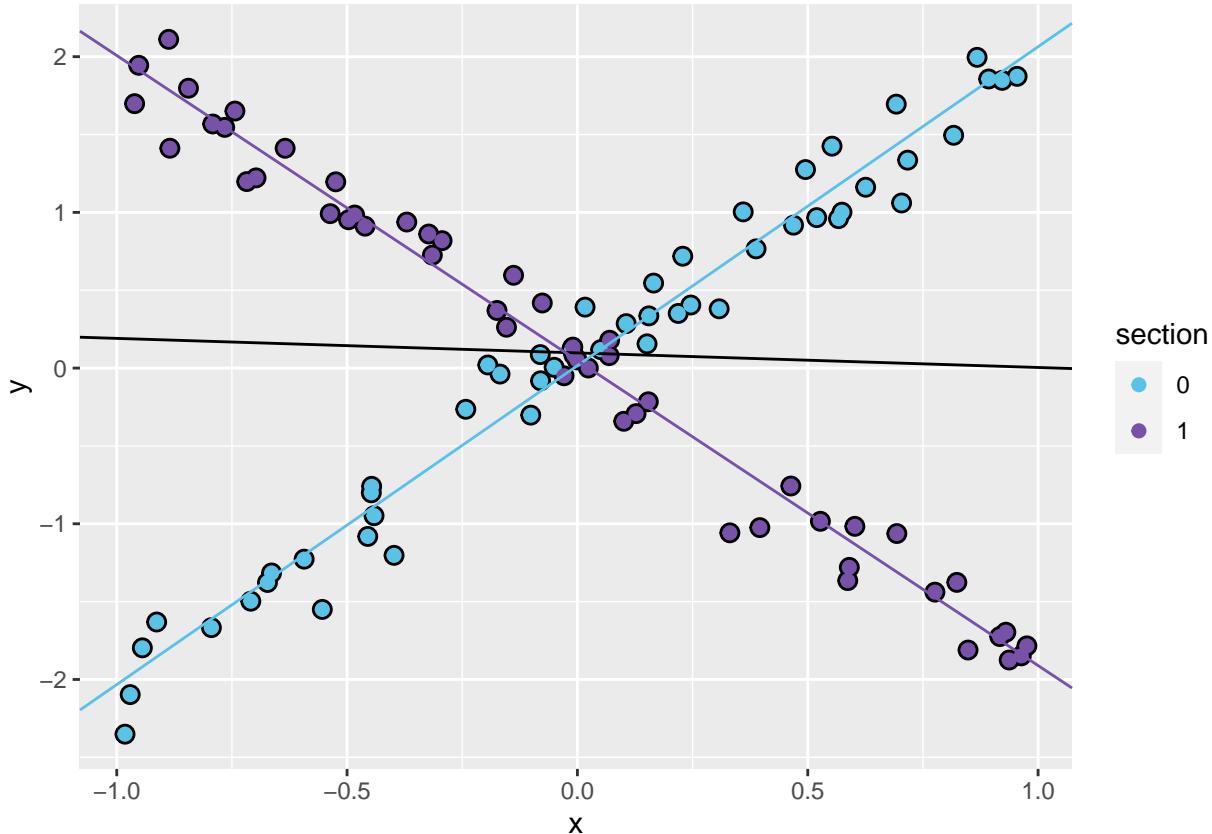
```

x <- c(runif(n/2, -1, 1), runif(n/2, -1, 1))
tau <- 0
tau1 <- -4
y <- beta0 + x * beta1 + t * tau + t * x * tau1 + rnorm(n, sd = sigma)
plot <- ggplot(data.frame(x = x, y = y, section = factor(t)),
 aes(x,y, colour = section)) +
 scale_colour_manual(values = pal) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2)
plot

```



```
##Fitting a model
xfit <- lm(y ~ x)
xtfit <- lm(y ~ x + t + I(x * t))
plot + geom_abline(intercept = xfit$coef[1], slope = xfit$coef[2]) +
 geom_abline(intercept = xtfit$coef[1], slope = xtfit$coef[2],
 colour = pal[1]) +
 geom_abline(intercept = xtfit$coef[1] + xtfit$coef[3],
 slope = xtfit$coef[2] + xtfit$coef[4],
 colour = pal[2])
```



```
c(mean(y[1:(n/2 - 1)]), mean(y[n/2:n]))
```

```
[1] 0.07857856 0.28650081
```

- Here the marginal difference is quite small, however looking at the trend by group shows that section 1 has a decrease in Y as X increases. However the intercept is quite high. As such one couldn't advise for a treatment if an individual had a low x.

### Some final thoughts

- Modeling multivariate relationships is difficult
- Play around with simulations to see how the inclusion or exclusion of another variable can change analyses
- The results of these analyses deal with the impact of variables on associations
  - Ascertaining mechanisms or cause are difficult subjects to be added on top of difficulty in understanding multivariate associations

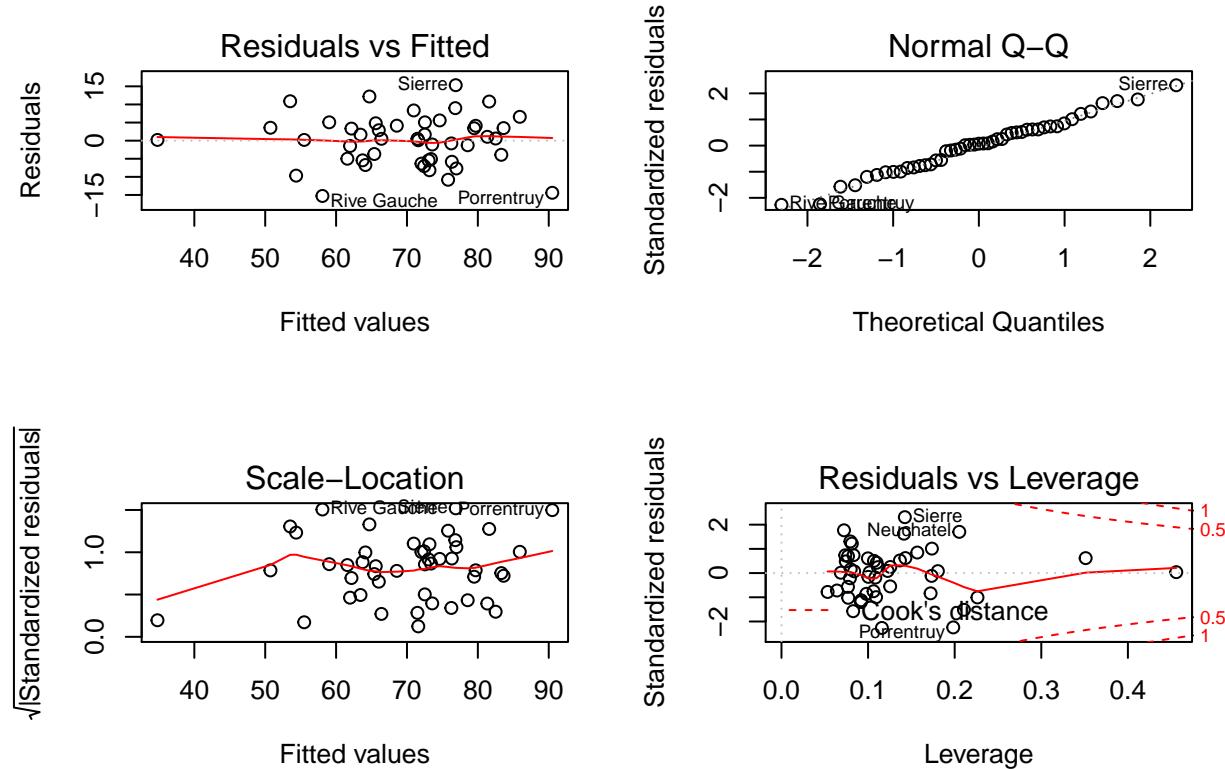
## Residuals Again

### Residuals and Diagnostics Part 1

#### The Linear Model

- Specified as  $Y_i = \sum_{k=1}^p X_{ik}\beta_j + \epsilon_i$
- We'll also assume here that  $\epsilon_i \sim^{iid} N(0, \sigma^2)$
- The residuals are defined as  $e_i = Y_i - \hat{Y}_i = Y_i - \sum_{k=1}^p X_{ik}\hat{\beta}_j$
- Our estimate of residual variation is  $\hat{\sigma}^2 = \frac{\sum_{i=1}^n e_i^2}{n-p}$ , the  $n-p$  is such that  $E[\hat{\sigma}^2] = \sigma^2$

```
data(swiss); par(mfrow = c(2,2))
fit <- lm(Fertility ~ ., data = swiss)
#R will splot the fit with a series of residual and diagnostic plots
plot(fit)
```



- Residuals vs. Fitted
  - Plots residuals against the predicted (fitted) values.
- Normal Q-Q
  - Tests for normality of error terms, plots true z-scores versus normalized z-scores

- Scale-Location
  - Plots standardized residuals against fitted values, scaled so they are like a t-statistic
- Residuals vs Leverage
  - Plots standardized residuals against their leverage

## Leverage

- Points that are far from the center of data have a lot of leverage, similar to a weight on the end of a lever having more leverage than if it's near the fulcrum

## Residuals and Diagnostics Part 2

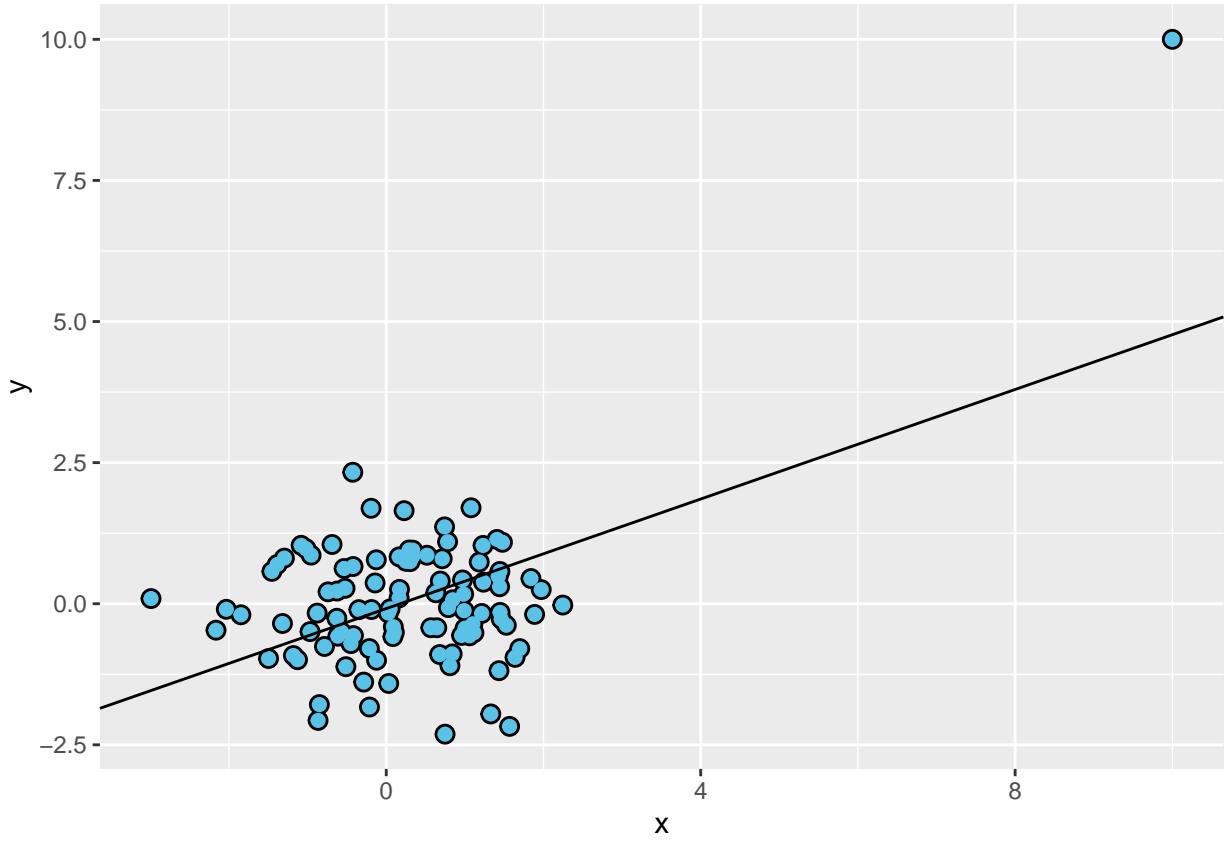
- Calling a point an outlier is vague.
  - They can be the result of real or spurious processes.
    - \* If they're real we don't want to just get rid of it
    - \* If they are spurious and there is evidence for this we can safely remove it from the model
- Outliers can conform to the regression relationship (being marginally outlying in X or Y, but not outlying given the regression relationship)
- There are various ways we can investigate outliers to ascertain how much influence they have.
- A list of default diagnostic measures can be obtained by executing `?influence.measures`
  - `rstandard` - standardized residuals, (residuals divided by their standard deviations)
  - `rstudent` - standardized residuals, residuals divided by their standard deviations, where the  $i^{th}$  data point was deleted in the calculation of the standard deviation for the residual to follow a t distribution
  - `hatvalues` - measures of leverage
    - \* Useful in finding data entry errors
  - `dffits` - change in the predicted response when the  $i^{th}$  point is deleted in fitting the model. Helps measure for influence
    - \* Returns a matrix of two values, slope & intercept
  - `dfbetas` - change in individual coefficients when the  $i^{th}$  point is deleted in fitting the model. Helps measure for influence
    - \* Returns a matrix of two values, slope & intercept
  - `cooks.distance` - overall change in the coefficients when the  $i^{th}$  point is deleted.
    - \* In a sense, summarizes `dfbetas`

- `resid` - returns the ordinary residuals
- `resid(fit) / (1 - hatvalues(fit))` where `fit` is the linear model fit returns the PRESS residuals, i.e. the leave one out cross validation residuals - the difference in the response and the predicted response at data point  $i$ , where it was not included in the model fitting.
- One can not just set a residual threshold and exclude points outside of it as that is akin to P-hacking
  - However one can declare points outliers from the residual, utilizing `rstandard` or `rstudent`
  - Although, it is better to look at them with respect to the overall cloud of data rather than just setting strict thresholds
- Linear algebra exploits are used such that the model doesn't have to be refit when using these models

## Residuals and Diagnostics Part 3

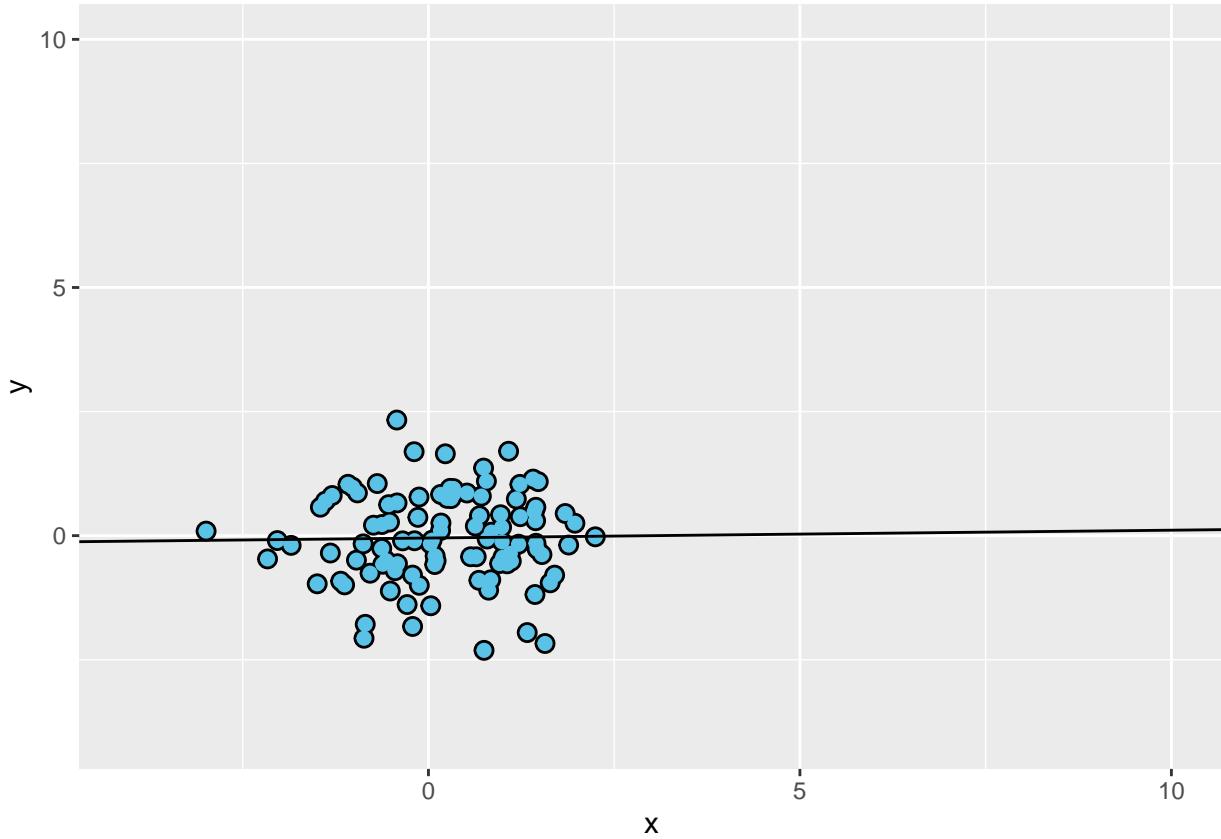
### Simulation 1

```
set.seed(1618033)
n <- 100
x <- c(10, rnorm(n))
y <- c(10, rnorm(n))
plot <- ggplot(data.frame(x = x, y = y), aes(x,y)) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2, colour = "#5BC2E7")
fit <- lm(y ~ x)
plot + geom_abline(intercept = fit$coef[1], slope = fit$coef[2])
```



- The point (10, 10) has created a strong regression relationship where there shouldn't be one

```
explot <- ggplot(data.frame(x = x[-1], y = y[-1]), aes(x,y)) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2, colour = "#5BC2E7") +
 xlim(c(-4,10)) + ylim(c(-4,10))
exfit <- lm(y[-1]~x[-1])
explot + geom_abline(intercept = exfit$coef[1], slope = exfit$coef[2])
```



Diagnostic values:

```
round(dfbetas(fit)[1:10, 2], 3) #Slope dif.
```

```
1 2 3 4 5 6 7 8 9 10
7.404 0.010 0.011 0.012 0.004 -0.125 -0.142 -0.008 -0.003 -0.148
```

```
round(hatvalues(fit)[1:10], 3) #
```

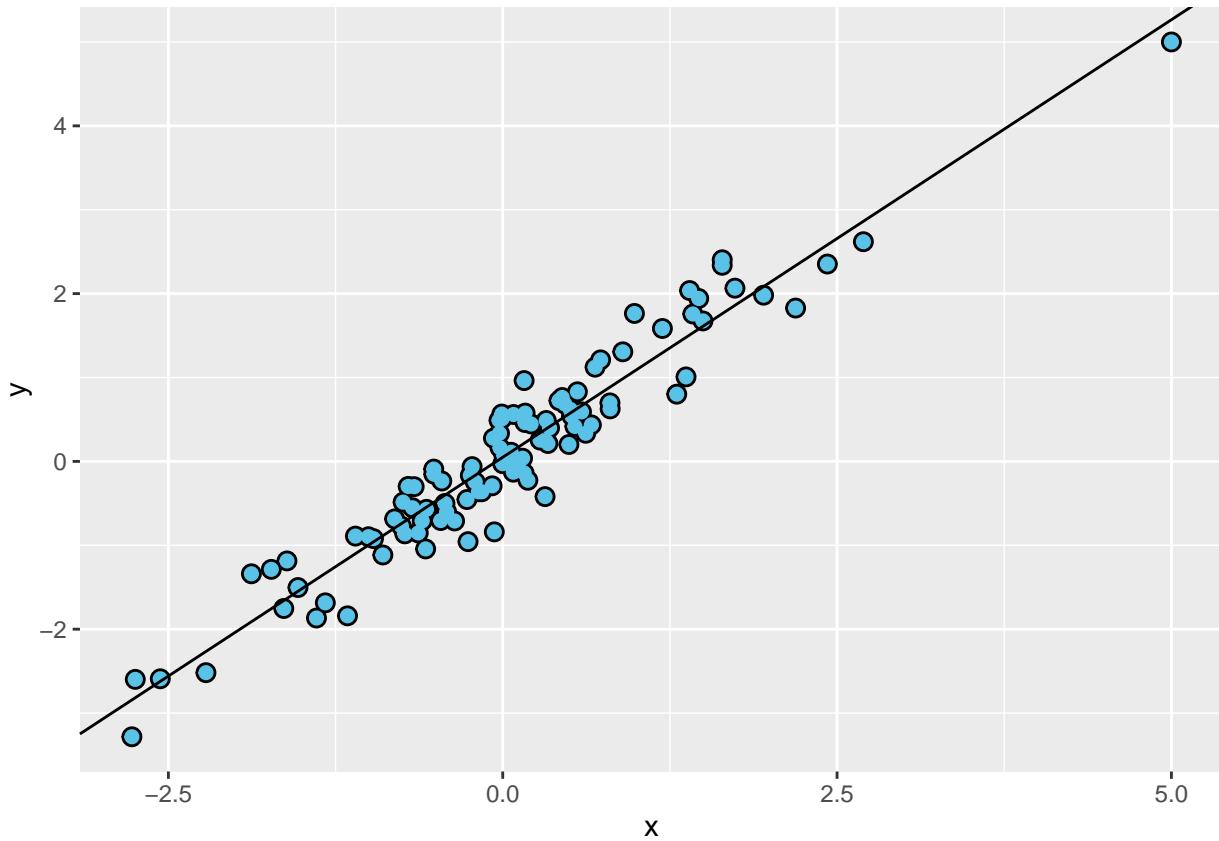
```
1 2 3 4 5 6 7 8 9 10
0.471 0.010 0.014 0.012 0.011 0.029 0.036 0.014 0.016 0.025
```

- Notice the first point in both instances are larger than the rest, indicating the first point's influence

## Simulation 2

```
x <- rnorm(n)
y <- x + rnorm(n, sd = 0.3)
x <- c(5, x)
y <- c(5, y)
plot <- ggplot(data.frame(x = x, y = y), aes(x,y)) +
 geom_point(size = 3, colour = "#000000") +
 geom_point(size = 2, colour = "#5BC2E7")
fit <- lm(y ~ x)
```

```
plot + geom_abline(intercept = fit$coef[1], slope = fit$coef[2])
```



- Here the outlier has a lot of leverage but adheres nicely to our model  
Diagnostic values:

```
round(dfbetas(fit)[1:10, 2], 3) #Slope dif.
```

```
1 2 3 4 5 6 7 8 9 10
-0.410 -0.057 -0.018 -0.075 0.011 0.242 -0.002 -0.057 -0.014 -0.290
```

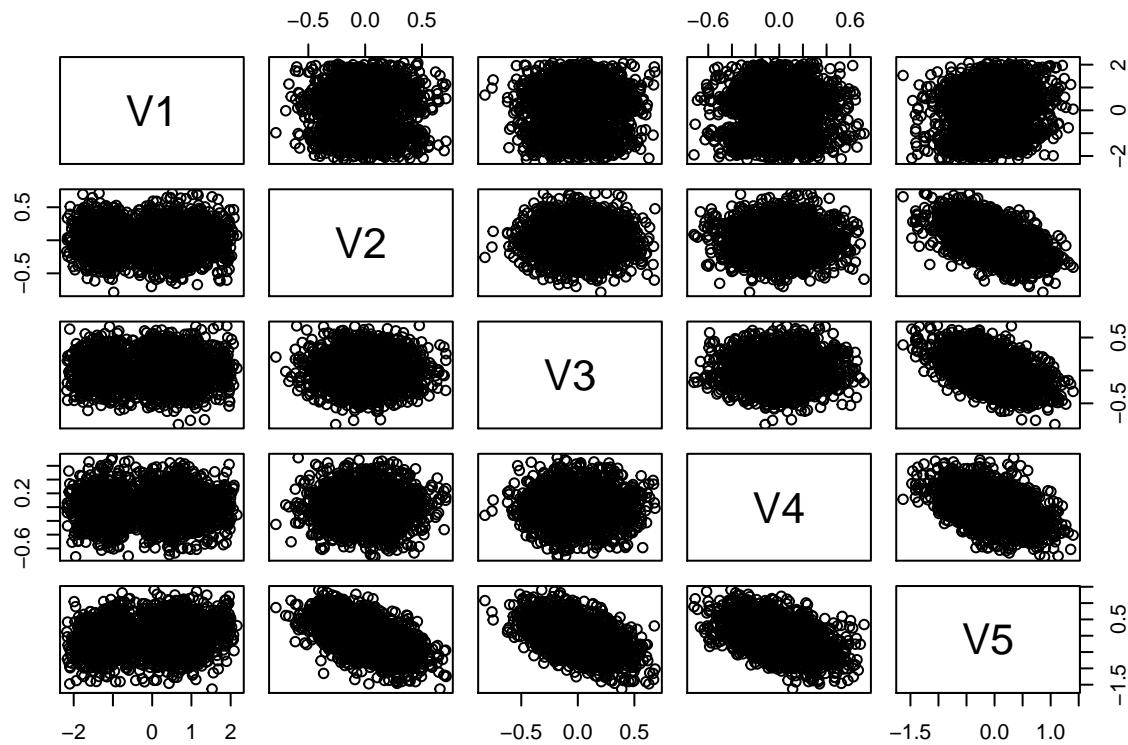
```
round(hatvalues(fit)[1:10], 3) #
```

```
1 2 3 4 5 6 7 8 9 10
0.194 0.012 0.010 0.014 0.010 0.029 0.010 0.010 0.011 0.038
```

- Here the first value is still large but not as much as previous
- The hatvalue is larger since its outside of the x values but still adheres to the relationship

## Using Residuals to See Underlying Patterns

```
dat <- read.table('http://www4.stat.ncsu.edu/~stefanski/NSF_Supported/Hidden_Images/orly_owl_f
pairs(dat) #Not much detail due to overplotting
```

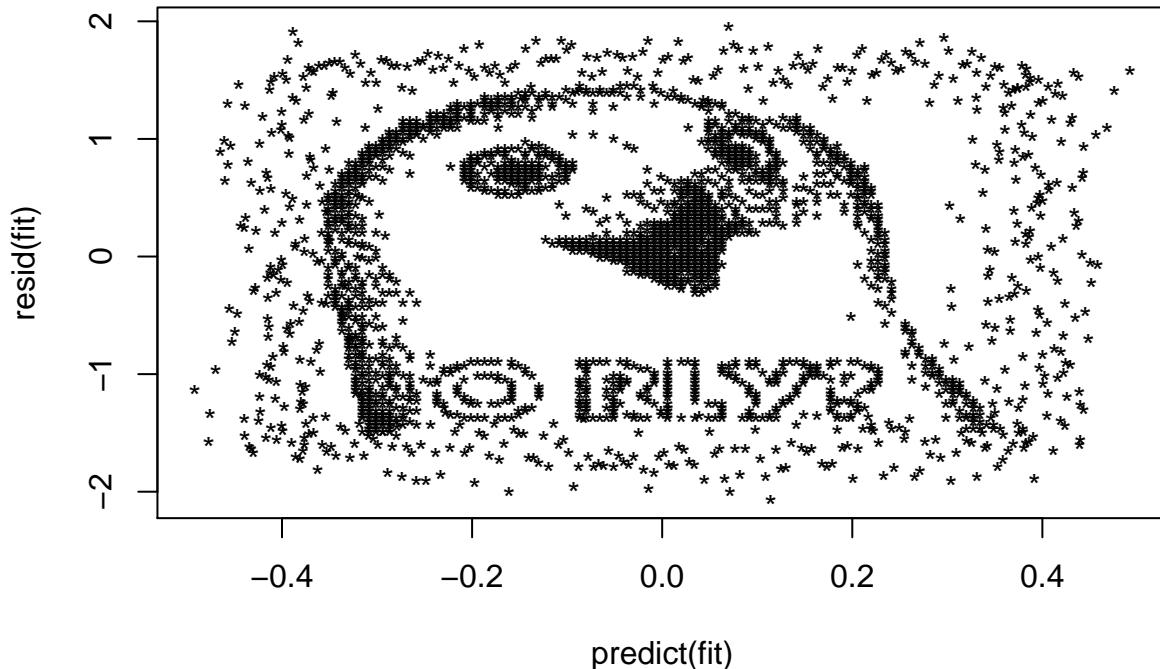


```
#Looking at P-values
fit <- lm(V1 ~ . - 1, dat)
summary(fit)$coef
```

|       | Estimate  | Std. Error | t value   | Pr(> t )     |
|-------|-----------|------------|-----------|--------------|
| ## V2 | 0.9856157 | 0.12798121 | 7.701253  | 1.989126e-14 |
| ## V3 | 0.9714707 | 0.12663829 | 7.671225  | 2.500259e-14 |
| ## V4 | 0.8606368 | 0.11958267 | 7.197003  | 8.301184e-13 |
| ## V5 | 0.9266981 | 0.08328434 | 11.126919 | 4.778110e-28 |

The fit shows that all the variables are significant, but we should still look at the residuals to see if there is an underlying pattern

```
plot(predict(fit), resid(fit), pch = '*')
```



And by plotting the residuals we can see that although the data is significant there is an underlying pattern

### **Lesson with `swirl()`: Residuals Diagnostics and Variation**

(No new content)

## **Model Selection**

How do we chose what variables to include in a regression model?

\* It depends, but that's what this section aims to look at

### **Model Selection Part 1**

#### **Multivariable Regression**

- The next course is on prediction and machine learning, here we're focusing on modeling
  - Prediction has a different set of criteria, needs for interpretability are less
  - In modeling, our interest lies in parsimonious, as simple as possible but no simpler. Interpretable representations enhance our understanding of the phenomena beneath the data

- “A model is a lense through which to look at your data” - Scott Zeger
  - \* Under this philosophy, what’s the right model? Whatever model connects the data to a true, parsimonious statement about what you’re studying
- There are nearly uncountable ways that a model can be wrong, this can be adjusted with variable inclusion and exclusion

## The Rumsfeldian Triplet

*“There are known knowns. These are things we know that we know. There are known unknowns. That is to say, there are things that we know we don’t know. But there are also unknowns. There are things we don’t know we don’t know.”* -Donald Rumsfeld

In terms of regression:

- \* (Known knowns) Regressors that we know we should check to include in the model and we have.
- \* (Known Unknowns) Regressors that we would like to include in the model, but simply don’t have.
- \* (Unknown Unknowns) Regressors that we don’t even know about that we should have included in the model.

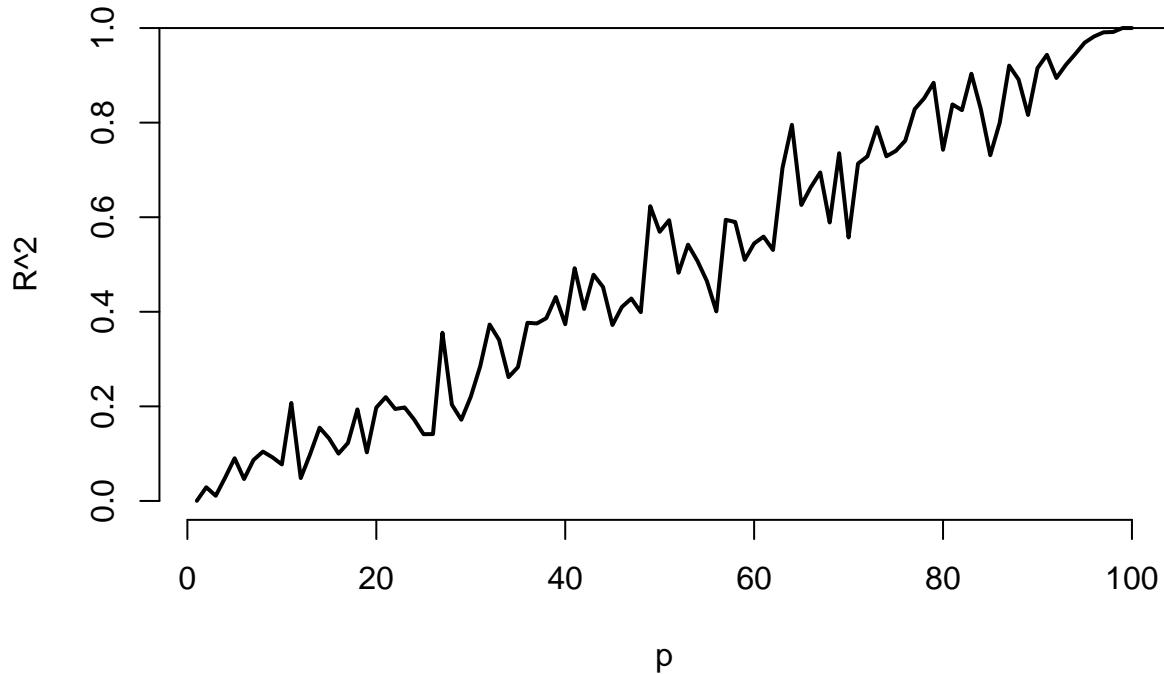
## General Rules

- Omitting variables results in bias in the coefficients of interest - unless their regressors are uncorrelated with the omitted ones.
    - This is why we randomize treatments, it attempts to uncorrelate our treatment indicator with variables that we don’t have access to to put in the model
    - (If there’s too many unobserved confounding variables, even randomization won’t help you.)
  - Including variables that we shouldn’t have increases standard errors of the regression variables
    - Including any new variables increases (actual, not estimated) standard errors of other regressors. So we don’t want to idly throw variables into the model.
  - The model must tend toward perfect fit as the number of non-redundant regressors approaches  $n$
  - $R^2$  increases monotonically as more regressors are included
  - The mean square error decreases monotonically as more regressors are included
- Below is a plot of  $R^2$  versus  $n$  where there is no true correlation underlying

```
set.seed(1618033)
n <- 100
plot(c(1, n), 0:1, type = "n", frame = FALSE, xlab = "p", ylab = "R^2")
r <- sapply(1:n, function(p)
{
 y <- rnorm(n); x <- matrix(rnorm(n * p), n, p)
 summary(lm(y ~ x))$r.squared
})

```

```
lines(1:n, r, lwd = 2)
abline(h = 1)
```



As we can see it trends towards 1, even though it's just noise being added

## Model Selection Part 2

### Variance Inflation

```
set.seed(1618033)
n <- 100
nosim <- 1000
x1 <- rnorm(n); x2 <- rnorm(n); x3 <- rnorm(n);
betas <- sapply(1:nosim, function(i){
 y <- x1 + rnorm(n, sd = .3)
 ##Looking at variable for each added variable
 c(coef(lm(y ~ x1))[2],
 coef(lm(y ~ x1 + x2))[2],
 coef(lm(y ~ x1 + x2 + x3))[2])
})
round(apply(betas, 1, sd), 5) #Look at sd of regression models
```

## x1 x1 x1

```
0.02898 0.02899 0.02899
```

- It can be seen that regardless of how many regressors we include the variance of the first coefficient does not increase

Let's look at a case where each variable relies on  $x_1$

```
x1 <- rnorm(n)
x2 <- x1/sqrt(2) + rnorm(n)/sqrt(2) #Slightly depends on x1
x3 <- x1 * 0.95 + rnorm(n) * sqrt(1 - 0.95^2); #Heavily depends on x1
betas <- sapply(1:nosim, function(i){
 y <- x1 + rnorm(n, sd = .3)
 c(coef(lm(y ~ x1))[2],
 coef(lm(y ~ x1 + x2))[2],
 coef(lm(y ~ x1 + x2 + x3))[2])
})
round(apply(betas, 1, sd), 5)
```

```
x1 x1 x1
0.03076 0.04443 0.10256
```

- Here we can see an increase in the variance of the coefficients as we include more variables that are all dependent on  $x_1$ 
  - If the variable you include is highly correlated to what you're investigating you'll get a higher standard error
    - \* If you can avoid it that's ideal but if you have to map both variables (such as height and weight) you'll get this drawback of a higher standard error

## Variance Inflation Factors

- Notice variance inflation was much worse when we included a variable that was highly related to  $x_1$ .
- We don't know  $\sigma$ , so we can only estimate the increase in the actual standard error of the coefficients for including a regressor
- However,  $\sigma$  drops out of the relative standard error. If one sequentially adds variables, one can check the variance inflation for including each one.
- When the other regressors are actually orthogonal to the regressor of interest, then there is no variance inflation
- The **variance inflation factor (VIF)** is the increase in the variance for the  $i^{th}$  regressor compared to the ideal setting where it is orthogonal to the other regressors
  - $\sqrt{VIF} =$  the increase in the sd
- Remember that variance inflation is only part of the picture. We want to include certain variables, even if they dramatically inflate our variance.

## Variance Inflation with Swiss Data

```

library(car) #Companion to Applied Regression
data(swiss)

fit <- lm(Fertility ~ ., swiss)
res <- data.frame(rbind(vif(fit), sqrt(vif(fit))))
rownames(res) <- c("Variance", "SD")
res

Agriculture Examination Education Catholic Infant.Mortality
Variance 2.284129 3.675420 2.774943 1.937160 1.107542
SD 1.511334 1.917138 1.665816 1.391819 1.052398

```

- [1,1] indicates that the Standard Error (SSE) of the Agriculture is double than what it would be if it were orthogonal to all the other regressors
- Infant.Mortality is low because it likely is already orthogonal to the other factors

### Model Selection Part 3

#### What About Residual Variance Estimation?

- Assuming that the model is linear with additive iid errors (with finite variance), we can mathematically describe the impact of omitting necessary variables or including unnecessary ones.
  - If we underfit the model, the variance estimate is biased due to that missing factor
  - If we correctly or overfit the model, including all necessary covariates and/or unnecessary covariates, the variance estimate is unbiased
    - \* However, the variance of the variance is larger if we include unnecessary variables
- Principal components or factor analytic models on covariates are often useful for reducing complex covariate spaces.
- Good design can often eliminate the need for complex model searches at analyses; though often control over the design is limited.
- If the models of interest are nested and without lots of parameters differentiating them, it's fairly uncontroversial to use nested likelihood ratio tests.

#### Nested Model Testing

```

fit1 <- lm(Fertility ~ Agriculture, swiss)
fit3 <- update(fit1, Fertility ~ Agriculture + Examination + Education)
fit5 <- update(fit3, Fertility ~ Agriculture + Examination + Education + Catholic + Infant.Mortality)
##Once can use ANOVA to perform an ANalysis Of the VAriance
anova(fit1, fit3, fit5)

Analysis of Variance Table

```

```


Model 1: Fertility ~ Agriculture
Model 2: Fertility ~ Agriculture + Examination + Education
Model 3: Fertility ~ Agriculture + Examination + Education + Catholic +
Infant.Mortality
Res.Df RSS Df Sum of Sq F Pr(>F)
1 45 6283.1
2 43 3180.9 2 3102.2 30.211 8.638e-09 ***
3 41 2105.0 2 1075.9 10.477 0.0002111 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ',' 1

```

- `Res.Df` - Number of Observations – Number of Parameters
- `RSS` - Residual Sum of Squares
- `Sum of Sq` - How many additional Df were added from previous model
- `F` - F-statistic
- `P....` - If significant it implies that the additional inclusion appears to be necessary

## Practice Exercise in Regression Modeling

My practice exercise can be found on GitHub

### Quiz 3

1. Consider the mtcars data set. Fit a model with mpg as the outcome that includes number of cylinders as a factor variable and weight as confounder. Give the adjusted estimate for the expected change in mpg comparing 8 cylinders to 4.

```

data(mtcars)
fit <- lm(mpg ~ factor(cyl) + wt, mtcars)
summary(fit)$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 33.990794 1.8877934 18.005569 6.257246e-17
factor(cyl)6 -4.255582 1.3860728 -3.070244 4.717834e-03
factor(cyl)8 -6.070860 1.6522878 -3.674214 9.991893e-04
wt -3.205613 0.7538957 -4.252065 2.130435e-04

```

- -6.071

2. Consider the mtcars data set. Fit a model with mpg as the outcome that includes number of cylinders as a factor variable and weight as a possible confounding variable. Compare the effect of 8 versus 4 cylinders on mpg for the adjusted and unadjusted by weight models. Here, adjusted means including the weight variable as a term in the regression model and unadjusted

means the model without weight included. What can be said about the effect comparing 8 and 4 cylinders after looking at models with and without weight included?

```
cylfit <- lm(mpg ~ factor(cyl), mtcars)
summary(fit)$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 33.990794 1.8877934 18.005569 6.257246e-17
factor(cyl)6 -4.255582 1.3860728 -3.070244 4.717834e-03
factor(cyl)8 -6.070860 1.6522878 -3.674214 9.991893e-04
wt -3.205613 0.7538957 -4.252065 2.130435e-04

summary(cylfit)$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 26.663636 0.9718008 27.437347 2.688358e-22
factor(cyl)6 -6.920779 1.5583482 -4.441099 1.194696e-04
factor(cyl)8 -11.563636 1.2986235 -8.904534 8.568209e-10
```

3. Consider the mtcars data set. Fit a model with mpg as the outcome that considers number of cylinders as a factor variable and weight as confounder. Now fit a second model with mpg as the outcome model that considers the interaction between number of cylinders (as a factor variable) and weight. Give the P-value for the likelihood ratio test comparing the two models and suggest a model using 0.05 as a type I error rate significance benchmark.

```
intfit <- lm(mpg ~ factor(cyl) * wt, mtcars)
summary(fit)$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 33.990794 1.8877934 18.005569 6.257246e-17
factor(cyl)6 -4.255582 1.3860728 -3.070244 4.717834e-03
factor(cyl)8 -6.070860 1.6522878 -3.674214 9.991893e-04
wt -3.205613 0.7538957 -4.252065 2.130435e-04

summary(intfit)$coef

Estimate Std. Error t value Pr(>|t|)
(Intercept) 39.571196 3.193940 12.3894599 2.058359e-12
factor(cyl)6 -11.162351 9.355346 -1.1931522 2.435843e-01
factor(cyl)8 -15.703167 4.839464 -3.2448150 3.223216e-03
wt -5.647025 1.359498 -4.1537586 3.127578e-04
factor(cyl)6:wt 2.866919 3.117330 0.9196716 3.661987e-01
factor(cyl)8:wt 3.454587 1.627261 2.1229458 4.344037e-02

anova(fit, intfit)

Analysis of Variance Table
##
Model 1: mpg ~ factor(cyl) + wt
Model 2: mpg ~ factor(cyl) * wt
Res.Df RSS Df Sum of Sq F Pr(>F)
1 28 183.06
```

```
2 26 155.89 2 27.17 2.2658 0.1239
```

4. Consider the mtcarsmtcarsmtcars data set. Fit a model with mpg as the outcome that includes number of cylinders as a factor variable and weight included in the model as:

```
weightedfit <- lm(mpg ~ I(wt * 0.5) + factor(cyl), mtcars)
```

How is the wt coefficient interpreted?

```
summary(weightedfit)$coef
```

```
Estimate Std. Error t value Pr(>|t|)
(Intercept) 33.990794 1.887793 18.005569 6.257246e-17
I(wt * 0.5) -6.411227 1.507791 -4.252065 2.130435e-04
factor(cyl)6 -4.255582 1.386073 -3.070244 4.717834e-03
factor(cyl)8 -6.070860 1.652288 -3.674214 9.991893e-04
```

- Change in mpg per 1 ton increase in weight while holding cyl constant (Since wt is originally in 1000 lbs)

5. Given:

```
x <- c(0.586, 0.166, -0.042, -0.614, 11.72)
y <- c(0.549, -0.026, -0.127, -0.751, 1.344)
```

Give the hat diagonal for the most influential point

```
fit <- lm(y ~ x)
hatvalues(fit)
```

```
1 2 3 4 5
0.2286650 0.2438146 0.2525027 0.2804443 0.9945734
```

6. With the same dataset as question 5, give the slope dfbeta for the point with the highest hat value.

```
dfbetas(fit)
```

```
(Intercept) x
1 1.06212391 -0.37811633
2 0.06748037 -0.02861769
3 -0.01735756 0.00791512
4 -1.24958248 0.67253246
5 0.20432010 -133.82261293
```

7. Consider a regression relationship between Y and X with and without adjustment for a third variable Z. Which of the following is true about comparing the regression coefficient between Y and X with and without adjustment for Z.

- It is possible to have a scenario of Simpson's paradox where the coefficient can be positive and strongly significant before adjusting for Z, then negative and strongly significant after the adjustment.

# Logistic Regression and Poisson Regression

## GLMs

### Intro

- Generalized linear models (GLMs) framework came from **Nelder and Wedderburn in 1972** in the Journal of the Royal Statistical Society.
- GLM is a family of models that includes linear models.
  - Extending the family handles many of the issues with linear models, but at the expense of some complexity and loss of some of the mathematical “tidiness”
- A GLM involves three components
  - An *exponential family* model for the response
  - A systematic component via a linear predictor
  - A link function that connects the means of the response to the linear predictor
- The three most famous cases of GLMs are:
  - Linear models
  - Binomial and binary regression
  - Poisson regression

### Limitations of Linear Models

- Additive response models don't make much sense if the response is discrete, or strictly positive.
- Additive error models often don't make sense, for example if the outcome has to be positive.
- Transformations are often hard to interpret
  - There's value in modeling the data on the scale that it was collected.
  - Particularly interpretable transformations, natural logarithms in specific, aren't applicable for negative or zero values.

### Example, Linear Models

- Assume that  $Y_i \sim N(\mu_i, \sigma^2)$  (the Gaussian distribution is an exponential family distribution.)
- Define the linear predictor to be  $\eta_i = \sum_{k=1}^p X_{ik}\beta_k$ .

- The link function,  $g$ , is defined such that  $g(\mu) = \eta$ .
  - For linear models  $g(\mu) = \mu$  as such  $\mu_i = \eta_i$
- The difference here is that we state  $Y_i$  is normally distributed as opposed to saying the error is normally distributed. We then connect the mean to the linear predictor.
  - The reasoning for this will be more clear when we look at different settings

### Example, Logistic Regression

- Assume that  $Y_i \sim Bernoulli(\mu_i)$  (Binomial) such that  $E[Y_i] = \mu_i$  where  $0 \leq \mu_i \leq 1$ .
- Linear predictor  $\eta_i = \sum_{k=1}^p X_{ik}\beta_k$
- Link function  $g(\mu) = \eta = \log(\frac{\mu}{1-\mu})$ 
  - $g$  is the (natural) log odds, referred to as the **logit**
  - \*  $\log$  is assumed to be base  $e$
- The parameter estimates are obtained by maximizing the following likelihood function:  

$$\prod_{i=1}^n \mu_i^{y_i} (1 - \mu_i)^{1-y_i} = \exp(\sum_{i=1}^n y_i \eta_i) \prod_{i=1}^n (1 + \eta_i)^{-1}$$

### Example, Poisson Regression

- Assume that  $Y_i \sim Poisson(\mu_i)$  such that  $E[Y_i] = \mu_i$  where  $0 \leq \mu_i$
- Linear predictor (is still)  $\eta_i = \sum_{k=1}^p X_{ik}\beta_k$
- Link function  $g(\mu) = \eta = \log(\mu)$

### Some Things to Note

- In each case, the only way in which the likelihood depends on the data is through:  

$$\sum_{i=1}^n y_i \eta_i = \sum_{i=1}^n y_i \times \sum_{k=1}^p X_{ik}\beta_k = \sum_{k=1}^p \beta_k \times \sum_{i=1}^n X_{ik}y_i$$
- Thus if we don't need the full data, only  $\sum_{i=1}^n X_{ik}y_i$ .
  - This simplification is a consequence of choosing so-called “canonical” link functions
- All models achieve their maximum likelihood at the root of the so called normal equations  

$$0 = \sum_{i=1}^n \frac{(Y_i - \mu_i)}{Var(Y_i)} W_i$$
  - Where  $W_i$  are the derivative of the inverse of the link function.

### About Variances

- For the linear model  $Var(Y_i) = \sigma^2$  is constant.
- For the Bernoulli model  $Var(Y_i) = \mu_i(1 - \mu_i)$

- As such the variance depends on which observation you’re looking at
- For the Poisson model  $Var(Y_i) = \mu_i$ 
  - This variance also depends on the observation you’re looking at
- There are **quasi-** versions of the non-constant functions that can be used when the data doesn’t follow the variance assumptions of the GLMs.

## Odds & Ends

- The normal equations ahve to be solved iteratively. Resulting in  $\hat{\beta}_k$  and, if included,  $\hat{\phi}$ .
- Predicted linear predictor responses can be obtained as  $\hat{\eta} = \sum_{k=1}^p X_k \hat{\beta}_k$
- Predicted mean responses as  $\hat{\mu} = g^{-1}(\hat{\eta})$
- Coefficients are interpreted as  

$$g(E[Y|X_k = x_k + 1, X_{\sim k} = x_{\sim k}]) - g(E[Y|X_k = x_k, X_{\sim k} = x_{\sim k}]) = \beta_k$$
- This is interpreted as the change in the link function of the expected response per unit change in  $X_k$  holding other regressors constant.
  - Variations on Newon/Raphson’s algorithm are used to do it.
  - Asyptotics are used for inference usually, which means they require larger sample sizes for the P-values to be informative.
  - Many of the ideas from linear models can be brought over to GLMs

## Logistic Regression

- This section coverers logistic regression for binary outcomes

### Logistic Regression Part 1

#### Key Ideas

- Frequently we care about outcomes that ahve two values
  - Alive/Dead
  - Win/loss
  - Success/Failure
  - etc.

- Called binary, Bernoulli or 0/1 outcomes
- Collection of exchangeable binary outcomes for the same covariate data are called binomial outcomes.

### Example Baltimore Ravens win/loss

```
url <- "https://github.com/bcaffo/courses/raw/master/07_RegressionModels/03_02_binaryOutcomes.rda"
loc <- "./data/ravensData.rda"
download.file(url, destfile = loc)
load(loc)
head(ravensData)

ravenWinNum ravenWin ravenScore opponentScore
1 1 W 24 9
2 1 W 38 35
3 1 W 28 13
4 1 W 34 31
5 1 W 44 13
6 0 L 23 24
```

- This data contains a 1/0 for W/L which are the first two columns.
- The other two columns are the Raven's score and their opponent's score for the game

### Linear Regression

$RW_i = b_0 + b_1 RS_i + e_i$

- \*  $RW_i$  - 1 if a Ravens win, 0 if not
- \*  $RS_i$  - Number of points Ravens scored
- \*  $b_0$  - probability of a Ravens win if they score 0 points
- \*  $b_1$  - increase in probability of a Ravens win for each additional point
- \*  $e_i$  - residual variation

```
lmRavens <- lm(ravensData$ravenWinNum ~ ravensData$ravenScore)
res <- summary(lmRavens)$coef
res
```

|                           | Estimate   | Std. Error  | t value  | Pr(> t )   |
|---------------------------|------------|-------------|----------|------------|
| ## (Intercept)            | 0.28503172 | 0.256643165 | 1.110615 | 0.28135043 |
| ## ravensData\$ravenScore | 0.01589917 | 0.009058997 | 1.755069 | 0.09625261 |

```
res[1,]

Estimate Std. Error t value Pr(>|t|)
0.2850317 0.2566432 1.1106149 0.2813504
```

- From this model we can ascertain:
  - $b_0 = 0.285$

–  $b_1 = 0.0159$

## Odds

- Binary Outcome 0/1  
 $RW_i$

- Probability (0, 1)  
 $Pr(RW_i|RS_i, b_0, b_1) =$

$p$

- Odds  $(0, \infty)$

$$\frac{Pr(RW_i|RS_i, b_0, b_1)}{1 - Pr(RW_i|RS_i, b_0, b_1)} =$$

$$\frac{p}{1-p}$$

– Probability can be derived from the odds:  $p = \frac{O}{1+O}$

- Log odds  $(-\infty, \infty)$  (**logit**)

$$\log\left(\frac{Pr(RW_i|RS_i, b_0, b_1)}{1 - Pr(RW_i|RS_i, b_0, b_1)}\right) =$$

$$\log\left(\frac{p}{1-p}\right)$$

## Linear vs. Logistic Regression

- Linear definitions

–  $RW_i = b_0 + b_1 RS_i + e_i$

–  $E[RW_i|RS_i, b_0, b_1] = b_0 + b_1 RS_i$

- Logistic definitions

–  $Pr(RW_i|RS_i, b_0, b_1) = \frac{\exp(b_0 + b_1 RS_i)}{1 + \exp(b_0 + b_1 RS_i)}$

–  $\log\left(\frac{Pr(RW_i|RS_i, b_0, b_1)}{1 - Pr(RW_i|RS_i, b_0, b_1)}\right) = b_0 + b_1 RS_i$

## Interpreting Logistic Regression

$$\log\left(\frac{Pr(RW_i|RS_i, b_0, b_1)}{1 - Pr(RW_i|RS_i, b_0, b_1)}\right) = b_0 + b_1 RS_i$$

\*  $b_0$  - Log odds of a Ravens win if they score 0 points

\*  $b_1$  - Log odds ratio of win probability for each point scored (compared to zero points)

\*  $\exp(b_1)$  - Odds ratio of win probability for each point scored (compared to zero points)

## Odds

- Imagine that you are playing a game where you have a probability of success of  $p$  (Perhaps drawing cards or something).

- If you win you gain  $X$  dollars, if you lose you lose  $Y$  dollars.

- What should we set  $X$  and  $Y$  for the game to be fair?

$$E[\text{earnings}] = Xp - Y(1-p) = 0$$

- Which implies...

$$Xp = Y(1-p)$$

$$\frac{Xp}{1-p} = Y$$

$$\frac{p}{1-p} = \frac{Y}{X}$$

- The odds can be said as “How much should you be willing to pay for a  $p$  probability of winning a dollar?”

– If  $p > 0.5$  you have to pay more if you lose than you get if you win.

– If  $p < 0.5$  you have to pay less if you lose than you get if you win.

## Logistic Regression Part 2

### Visualizing Fitting Logistic Regression Curves

```
set.seed(1618033)
library(manipulate)
x <- seq(-10, 10, length = 1000)
manipulate can only be executed in an R Studio session
manipulate(
 plot(x, exp(beta0 + beta1 * x) / (1 + exp(beta0 + beta1 * x)),
 type = "l", lwd = 3, frame = FALSE),
 beta1 = slider(-2, 2, step = 0.1, initial = 2),
 beta0 = slider(-2, 2, step = 0.1, initial = 0)
)
```

- This is a unit step function with a large enough `beta1`
  - `beta1` adjusts the steepness and direction of the slope
  - `beta0` determines at what `x` value the function changes from 0 to 1 (or 1 to 0 if `beta1` is negative)
- This model is:  $\frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$

## Logistic Regression Part 3

### Ravens Logistic Regression

```
#If we had bounded count data we'd have to give a sample size
Assumes you want the logistic link function, which we do in this case
logRegRavens <- glm(ravensData$ravenWinNum ~ ravensData$ravenScore,
 family = "binomial")
summary(logRegRavens)
```

```


Call:

glm(formula = ravensData$ravenWinNum ~ ravensData$ravenScore,

family = "binomial")

Deviance Residuals:

Min 1Q Median 3Q Max

-1.7575 -1.0999 0.5305 0.8060 1.4947

Coefficients:

Estimate Std. Error z value Pr(>|z|)

(Intercept) -1.68001 1.55412 -1.081 0.28

ravensData$ravenScore 0.10658 0.06674 1.597 0.11

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 24.435 on 19 degrees of freedom

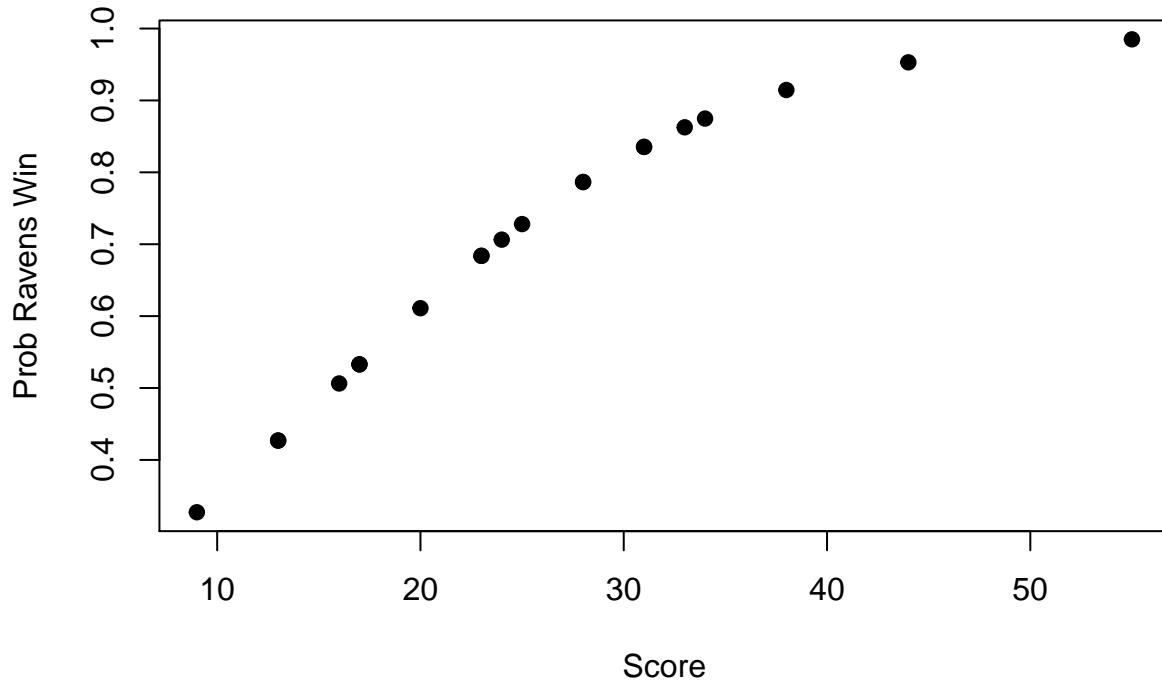
Residual deviance: 20.895 on 18 degrees of freedom

AIC: 24.895

Number of Fisher Scoring iterations: 5
• Fitted Values

plot(ravensData$ravenScore, logRegRavens$fitted,
 pch = 19, xlab = "Score", ylab = "Prob Ravens Win")

```



- These fitted values are only showing part of the curve since no data existed with a low enough score (i.e. 0)

### Odds Ratios & Confidence Intervals

```
#Suggests a [1,2] increase in p(winning) for each point scored
exp(logRegRavens$coeff)
```

```
(Intercept) ravensData$ravenScore
0.1863724 1.1124694
```

```
#Looking on exponential scale
exp(confint(logRegRavens))
```

```
Waiting for profiling to be done...
```

```
2.5 % 97.5 %
(Intercept) 0.005674966 3.106384
ravensData$ravenScore 0.996229662 1.303304
```

### ANOVA for Logistic Regression

```
#Additional models could be added, or factor variables could be considered
anova(logRegRavens, test = "Chisq")
```

```

Analysis of Deviance Table
##
Model: binomial, link: logit
##
Response: ravensData$ravenWinNum
##
Terms added sequentially (first to last)
##
Df Deviance Resid. Df Resid. Dev Pr(>Chi)
NULL 19 24.435
ravensData$ravenScore 1 3.5398 18 20.895 0.05991 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

## Interpreting Odds Ratios

- Not probabilities
- Odds ratio of 1 = no difference in odds
- Log odds ratio of 0 = no difference in odds
- Odds ratio < 0.5 or > 2 is commonly referred to as a “moderate effect”
  - Depends on the context
- Relative risk,  $\frac{Pr(RW_i|RS_i=10)}{Pr(RW_i|RS_i=0)}$ , is often easier to interpret, but harder to estimate. It is just the ratio of two probabilities.
- For small probabilities Relative Risk  $\approx$  Odds Ratio but **they are not the same!**

## Further Resources

- **Wikipedia on Logistic Regression**
- **Logistic regression and glms in R PDFs:**
  - Brian Caffo’s lecture notes on: **Simpson’s paradox, Case-control studies**
- **Open Intro Chapter on Logistic Regression**

## Lesson with `swirl()`: Variance Inflation Factors

- Variance Inflation is when you include any new variables to a model, it will increase the standard errors of the other regressors

- This lesson looks at `vifSims.R`, which has some simulations of this in effect
  - `rgp1()` has constant variance since each variable is uncorrelated
  - `rgp2()` has different variances since `x2` & `x3` are correlated with `x1` \* This lesson also looks at the `swiss` data set again, which was previously covered.

### **Lesson with `swirl()`: Overfitting and Underfitting**

- This lesson has largely already been covered, however it does use `fitting.R` to simulate fitting different models
- The `shapiro.test()` function will test a parameter for normality, with a null-hypothesis that it is normally distributed
  - This can be used to test residuals, if they aren't normally distributed it suggests there may be an underlying pattern

### **Lesson with `swirl()`: Binary Outcomes**

(No additional content)

## **Poisson Regression**

### **Poisson Regression Part 1**

#### **Overview**

- Many data take the form of unbounded count data, for example:
  - Number of calls to a call center
  - Number of flu cases in an area
  - Number of hits to a web site
- Data may also be in the form of rates
  - Percent of children passing a test
  - Percent of hits to a website from a country
  - Pump failure per a unit of time
- In some cases the counts are clearly bounded. However, modeling the counts as unbounded is often done when the upper limit is not known or very large relative to the number of events.
- The starting point for most count analysis is the Poisson distribution

## Poisson Distribution

- The Poisson distribution is a useful model for counts and rates
- Here a rate is count per some monitoring time
- Some examples:
  - Modeling web traffic hits
  - Incidence rates (Number of new cases per person time at risk)
  - Approximating binomial probabilities with small  $p$  and large  $n$
  - Analyzing contingency table data

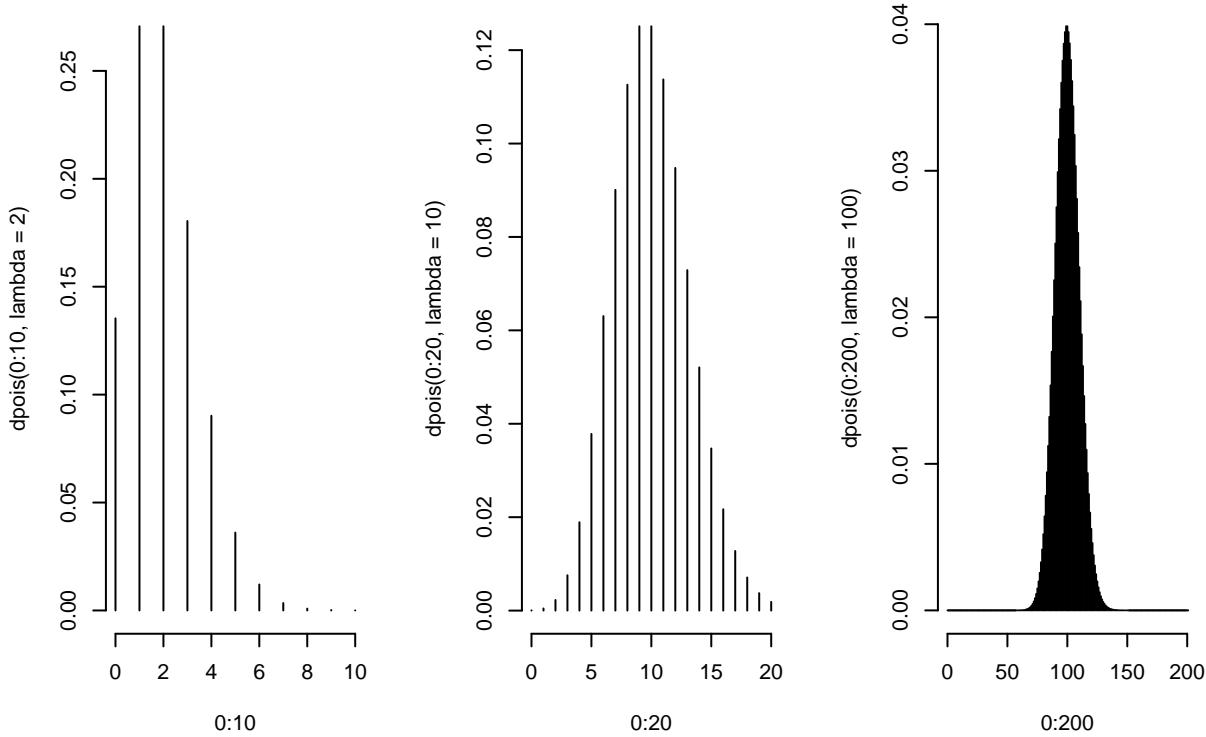
## The Poisson Mass Function

- $X \sim Poisson(t\lambda)$  if  
$$P(X = x) = \frac{(t\lambda)^x e^{-t\lambda}}{x!}$$
  - For  $x = 0, 1, \dots$
- The mean of the Poisson is  $E[X] = t\lambda$ , thus  $E[X/t] = \lambda$
- The variance of the Poisson is  $Var(X) = t\lambda$ .
  - This is an assumption that we can check
- The Poisson tends to a normal as  $t\lambda$  gets large.

## Simulation

- Simulate three sizes of poisson data and see that as the sample size grows it starts to appear normal

```
par(mfrow = c(1, 3))
set.seed(1618033)
plot(0:10, dpois(0:10, lambda = 2), type = "h", frame = FALSE)
plot(0:20, dpois(0:20, lambda = 10), type = "h", frame = FALSE)
plot(0:200, dpois(0:200, lambda = 100), type = "h", frame = FALSE)
```



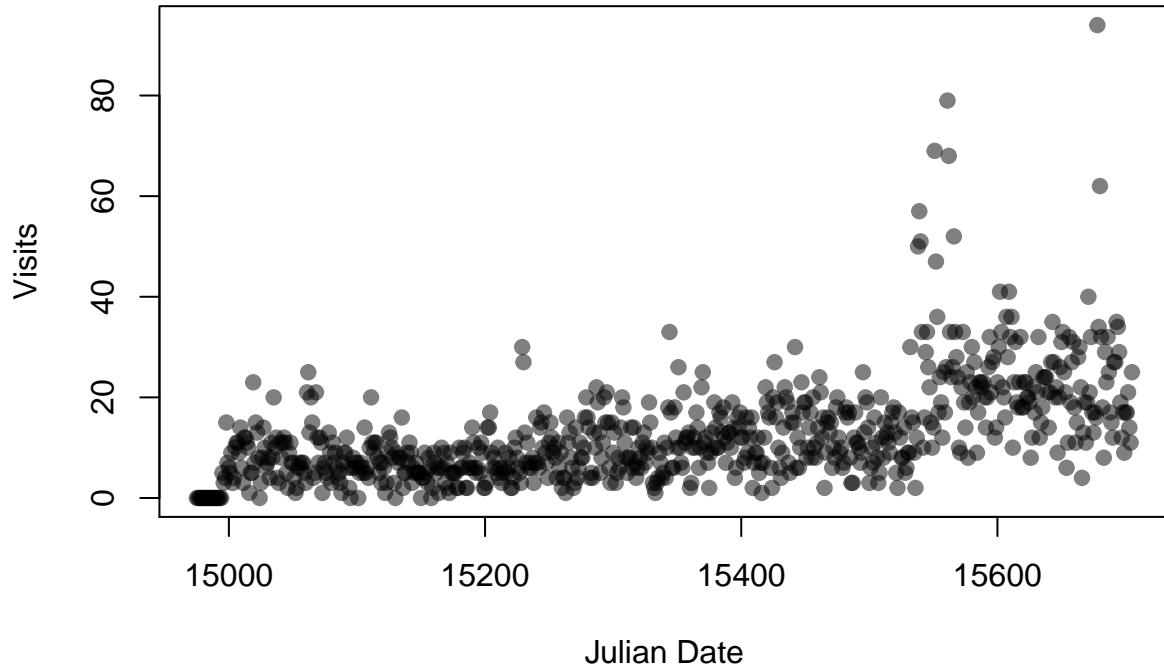
### Example: Leek Group Website Traffic

- Consider the daily counts to **Jeff Leek's web site**
- Since the unit of time is always one day, we set  $t = 1$  and then the Poisson mean is interpreted as web hits per day.
  - If we set  $t = 24$ , it would be web hits per hour.

```
download.file("https://github.com/bcaffo/courses/raw/master/07_RegressionModels/03_03_countOut.rda")
 destfile= "./data/gaData.rda")
load("./data/gaData.rda")
#Change char date to numeric julian date format
gaData$julian <- julian(gaData$date)
head(gaData)

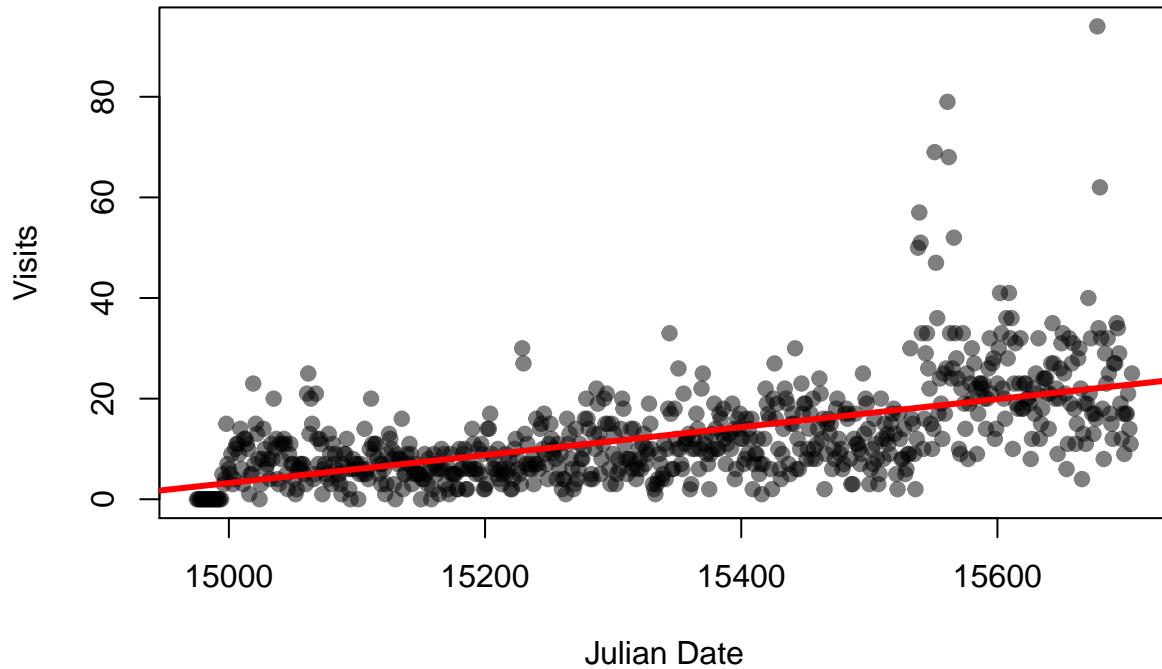
date visits simplystats julian
1 2011-01-01 0 0 14975
2 2011-01-02 0 0 14976
3 2011-01-03 0 0 14977
4 2011-01-04 0 0 14978
5 2011-01-05 0 0 14979
6 2011-01-06 0 0 14980
```

```
plot(gaData$julian, gaData$visits, pch = 19, col = rgb(0,0,0,0.5),
 xlab = "Julian Date", ylab = "Visits")
```



- Since this is count data a linear model won't be perfect at modeling it, even though it will trend to normality.

```
plot(gaData$julian, gaData$visits, pch = 19, col = rgb(0,0,0,0.5),
 xlab = "Julian Date", ylab = "Visits")
lm1 <- lm(gaData$visits ~ gaData$julian)
abline(lm1, col = "#FF0000", lwd = 3)
```



### Natural Log

- Taking the natural log of the outcome has a specific interpretation, consider the model:  
 $\log(NH_i) = b_0 + b_1 JD_i + e_i$ 
  - $NH_i$  - number of hits to the website
  - $JD_i$  - day fo the year (Julian day)
  - $b_0$  - log number of hits on Julian day 0 (1970-01-01)
  - $b_1$  - increase in log number of hits per unit day
  - $e_i$  - variation due to everything we didn't measure

### Exponentiating Coefficients

- $e^{E[\log(Y)]}$  geometric mean of  $Y$ .
  - With no covariates, this is estimated by  $e^{\frac{1}{n} \sum_{i=1}^n \log(y_i)} = (\prod_{i=1}^n y_i)^{1/n}$
- When you take the natural log fo outcomes and fit a regression model, you exponentiated coefficients estimate things about geometric means.

- $e^{\beta_0}$  estimated geometric mean hits on day 0
- $e^{\beta_1}$  estimated relative increase or decrease in geometric “successes” per day
- There’s an error from *log* if you have values of zero, adding a constant is a work-around for this

```
res <- round(exp(coef(lm(
 I(log(gaData$visits + 1)) ~ gaData$julian))), 5)
res

(Intercept) gaData$julian
0.00000 1.00231
```

- Since this is on the log scale you would subtract 1 from the suggested slope to get the percent increase per day, in this case it would be 0.23% increase in hits per day.

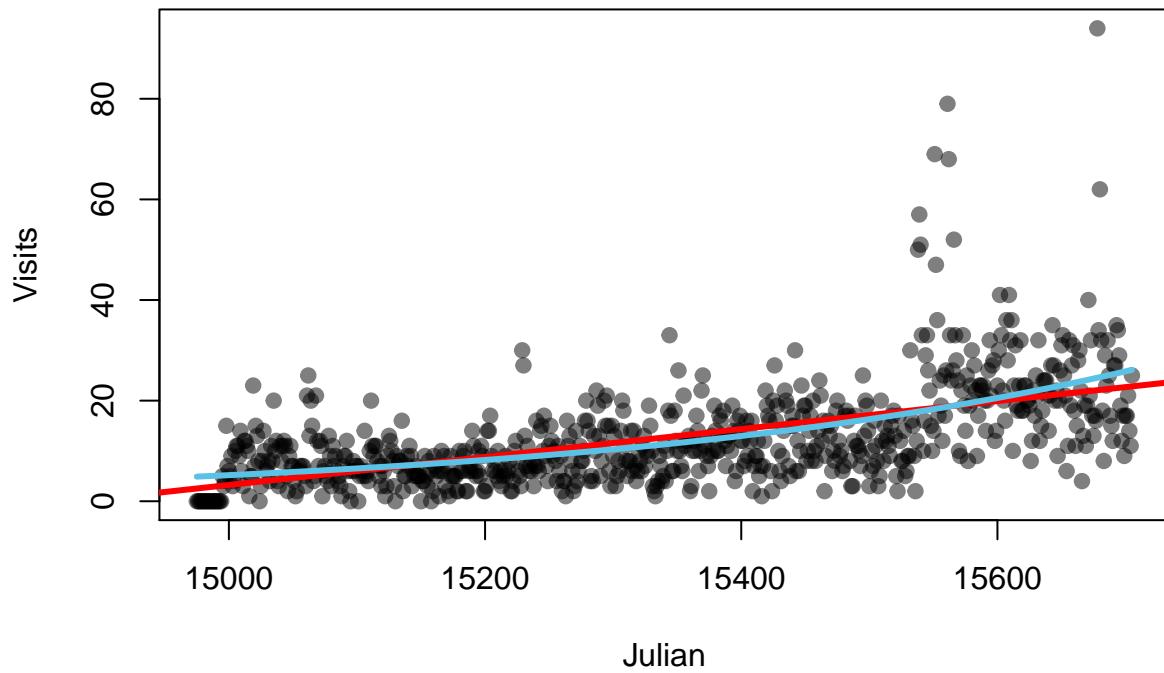
## Poisson Regression Part 2

### Linear vs. Poisson Regression

- Linear
  - $y_i = b_0 + b_1 x_i + e_i$
  - Also stated as:  $E[y_i|x_i, b_0, b_1] = b_0 + b_1 JD_i$
- Poisson/log-linear
  - $\log(E[y_i|x_i, b_0, b_1]) = b_0 + b_1 x_i$ 
    - \*  $\log$  of expected is linear
  - reevaluating this we get:  $E[y_i|x_i, b_0, b_1] = \exp(b_0 + b_1 JD_i)$
  - If  $x_i$  is increased by one unit,  $E[y_i|x_i, b_0, b_1]$  is *multiplied* by  $\exp(b_1)$

### Poisson Regression in R

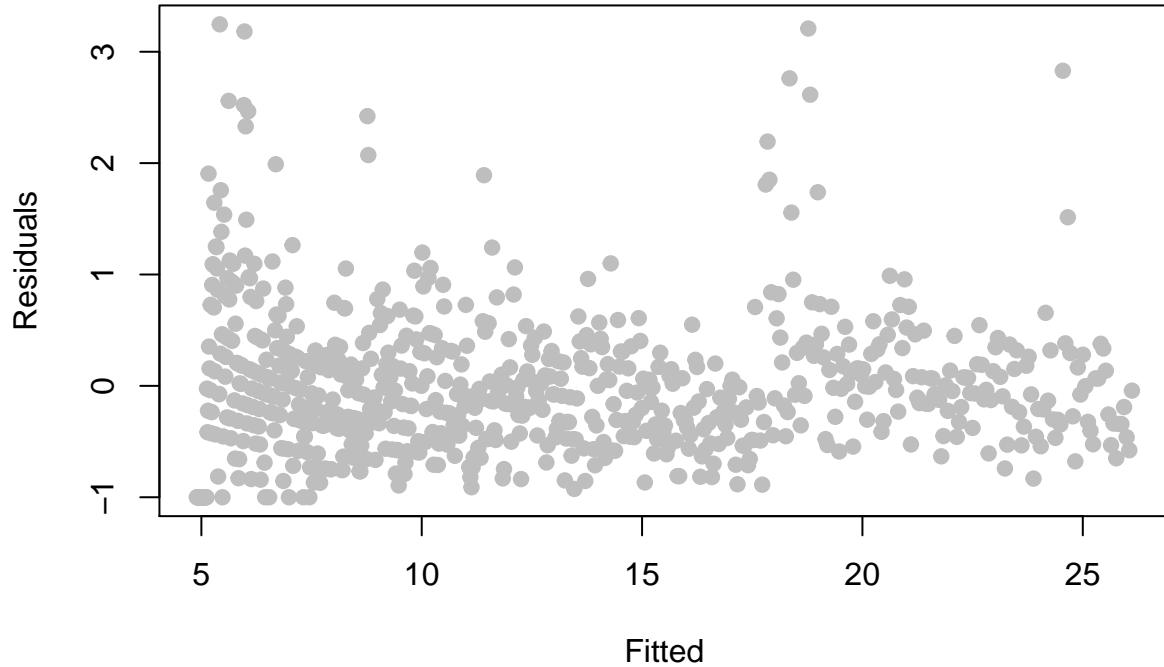
```
plot(gaData$julian, gaData$visits,
 pch = 19, col = rgb(0,0,0,0.5), xlab = "Julian", ylab = "Visits")
glm1 <- glm(gaData$visits ~ gaData$julian, family = "poisson")
abline(lm1, col = "#FF0000", lwd = 3)
lines(gaData$julian, glm1$fitted, col = "#5BC2E7", lwd = 3)
```



- As you can see here the poisson model does a better job at modeling the earlier slow rise of hits ( $\sim 15200$  -  $\sim 15600$ ) and the later growth ( $> \sim 15600$ )

#### Checking the Mean-Variance Relationship

```
plot(glm1$fitted, glm1$residuals,
 pch = 19, col = "grey", ylab = "Residuals", xlab = "Fitted")
```



- In this model the variance has to equal the mean, as such the variance should go up as the mean goes up. However, we can see in this plot that the variance is quite large for lower values. There are ways to account for this difference though.
- Estimating Confidence Intervals

```
confint(glm1)
```

```
Waiting for profiling to be done...
2.5 % 97.5 %
(Intercept) -34.346577587 -31.159715656
gaData$julian 0.002190043 0.002396461
```

## Rates

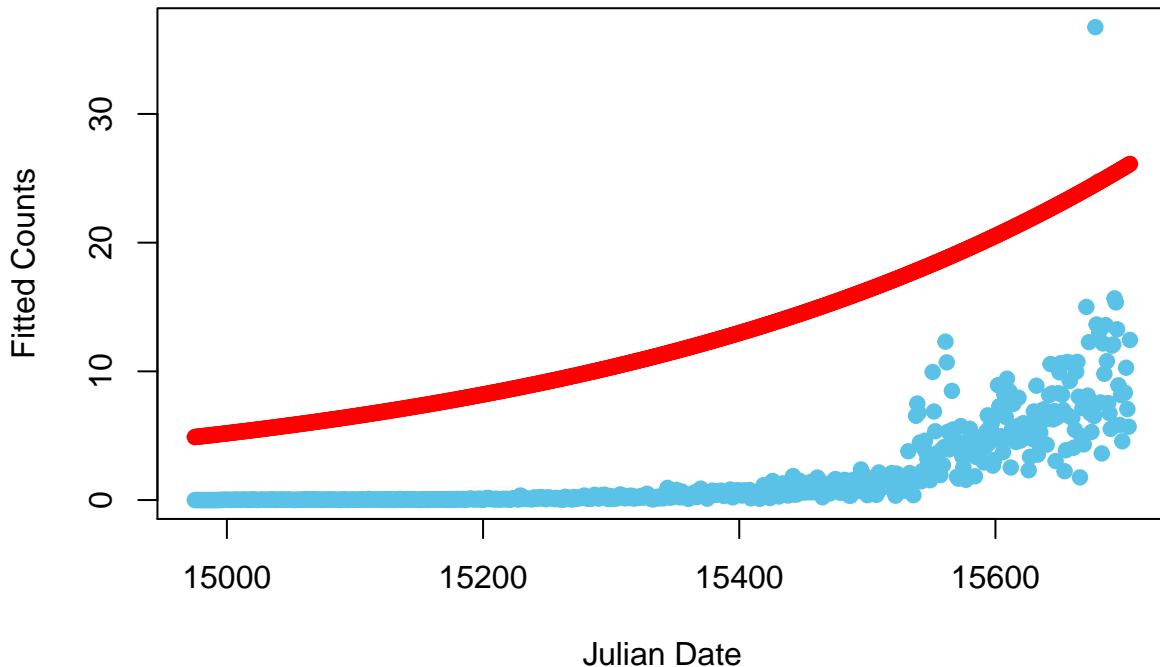
- In some cases one may want to compare a variable,  $NHSS_i$  (Number Hits for Simply Statistics) in our example, to a total variable,  $NH_i$  in our example, like a country's population if you're looking at a disease outbreak.
- We want to model:  $E[NHSS_i|JD_i, b_0, b_1]/NH_i = \exp(b_0 + b_1 JD_i)$ 
  - $JD_i$  is the Julian Date
  - This can be evaluated as:

$$* \log(E[NHSS_i|JD_i, b_0, b_1]) - \log(NH_i) = b_0 + b_1 JD_i$$

$$* \log(E[NHSS_i|JD_i, b_0, b_1]) = \log(NH_i) + b_0 + b_1 JD_i$$

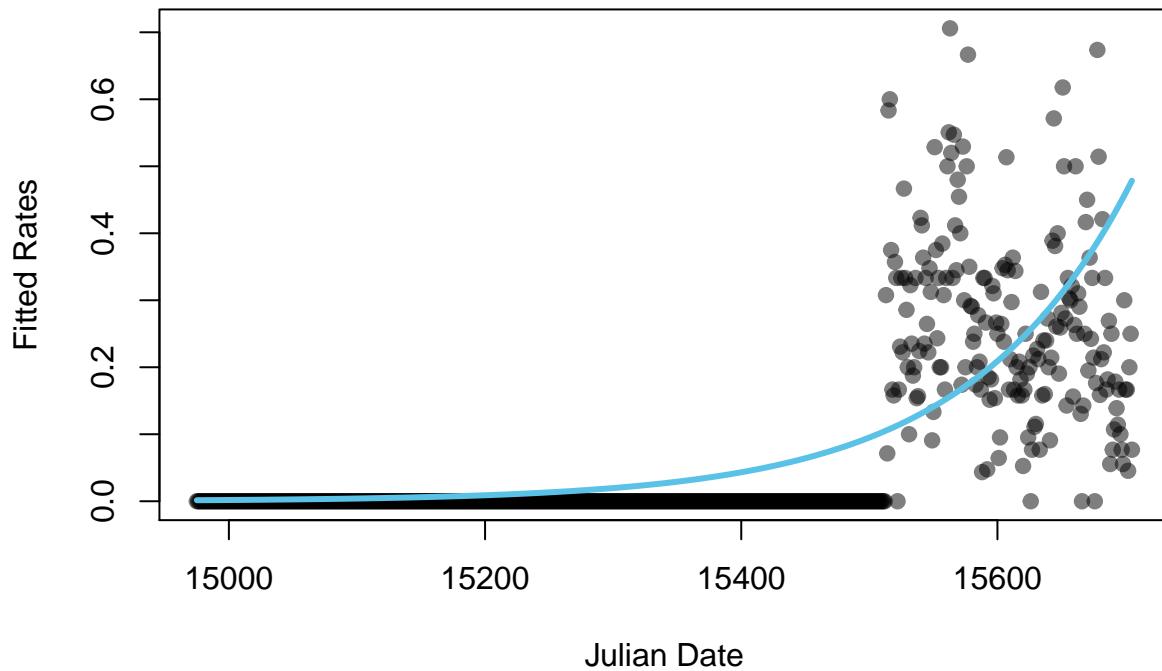
- Fitting in R

```
glm2 <- glm(gaData$simplystats ~ julian(gaData$date), offset = log(visits + 1),
 family = "poisson", data = gaData)
plot(julian(gaData$date), glm2$fitted,
 col = "#5BC2E7", pch = 19, xlab = "Julian Date", ylab = "Fitted Counts")
points(julian(gaData$date), glm1$fitted, col = "red", pch = 19)
```



- The blue points are adjusted for the red points, this will let the fit follow a more appropriate poisson model.

```
plot(julian(gaData$date), gaData$simplystats/(gaData$visits + 1),
 col = rgb(0,0,0,0.5), pch = 19,
 xlab = "Julian Date", ylab = "Fitted Rates")
lines(julian(gaData$date), glm2$fitted/(gaData$visits + 1),
 col = "#5BC2E7", lwd = 3)
```



### More Information

- Log-linear models and multiway tables
- Wikipedia on Poisson regression, Wikipedia on overdispersion
- PDF: Regression models for count data in R
- pscl package - the function `zeroinfl` fits zero inflated models.

### Lesson with `swirl()`: Count Outcomes

(No new content)

### Hodgepodge

#### Non-Linear Models

#### Fitting Functions Using Linear Models

- Consider a model where a function is applied to X:  $Y_i = f(X_i) + \epsilon$ .

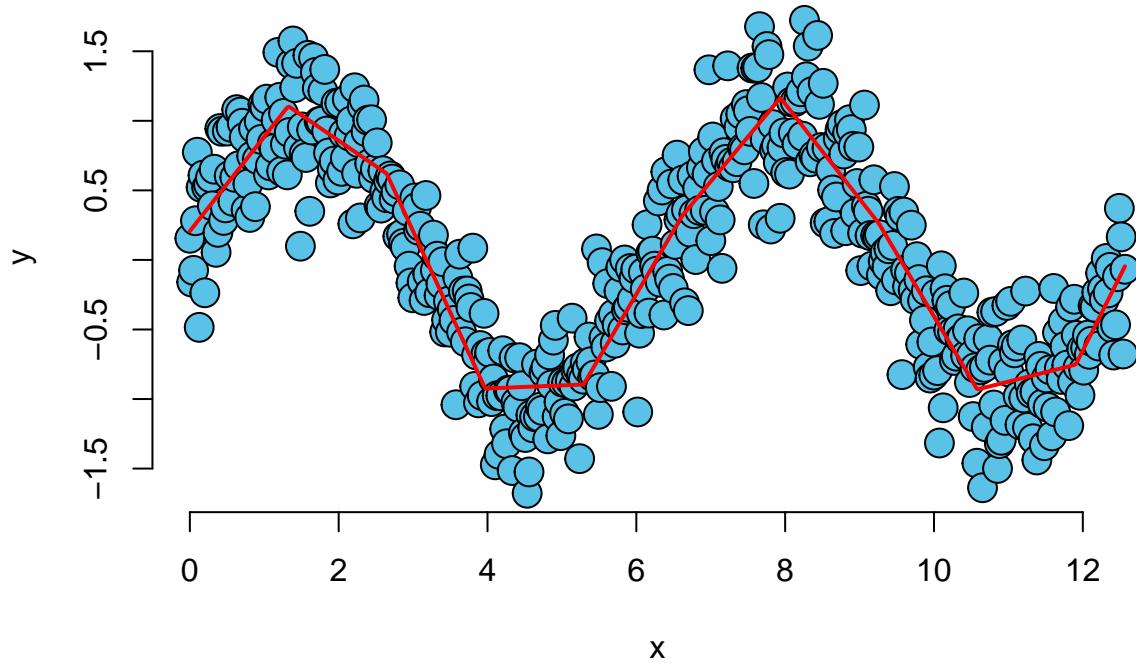
- How can we fit such a model using linear models (called scatterplot smoothing)
- Consider the model:  

$$Y_i = \beta_0 + \beta_1 X_i + \sum_{k=1}^d (x_i - \xi_k)_+ \gamma_k + \epsilon_i$$
  - where  $(a)_+ = a$  if  $a > 0$  and 0 otherwise and  $\xi_1 \leq \dots \leq \xi_d$  are known knot points ( $\xi$  is pronounced ' $k$ -si).
- The mean function,  

$$\beta_0 + \beta_1 X_i + \sum_{k=1}^d (x_i - \xi_k)_+ \gamma_k,$$
is continuous at the knot points.

### Example

```
set.seed(1618033)
n <- 500
x <- seq(0, 4 * pi, length = n)
y <- sin(x) + rnorm(n, sd = 0.3)
knots <- seq(0, 8 * pi, length = 20)
splineTerms <- sapply(knots,
 function(knot){(x > knot) * (x - knot)})
xMat <- cbind(1, x, splineTerms)
yhat <- predict(lm(y ~ xMat - 1)) #int. is included in x
plot(x, y, frame = FALSE, pch = 21, bg = "#5BC2E7", cex = 2)
lines(x, yhat, col = "#FF0000", lwd = 2)
```



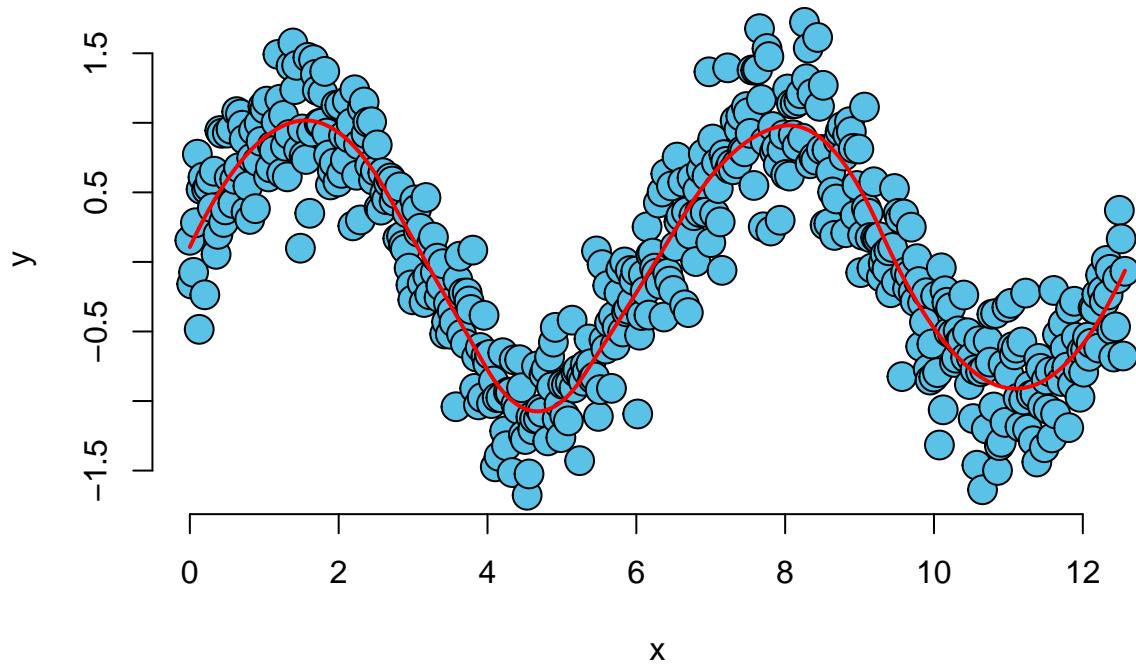
- The breakpoints are where the line turns at a sharp angle, the derivative of this line at these points would not be continuous
- Adding squared terms will fix this continuity issue

### Adding Squared Terms

- Adding squared terms makes it continuously differentiable at the knot points.  

$$Y_i = \beta_0 + \beta_1 X_i + \beta_2 X_i^2 + \sum_{k=1}^d (x_i - \xi_k)_+^2 \gamma_k + \epsilon_i$$
- Adding cubic terms makes it continuously differentiable on the second derivative, and so on for higher powers

```
splineTerms <- sapply(knots, function(knot) (x > knot) * (x - knot)^2)
xMat <- cbind(1, x, x^2, splineTerms)
yhat <- predict(lm(y ~ xMat - 1))
plot(x, y, frame = FALSE, pch = 21, bg = "#5BC2E7", cex = 2)
lines(x, yhat, col = "#FF0000", lwd = 2)
```



## Notes

- The collection of regressors is called a basis.
  - This is one of the main focuses of linear algebra
- Single knot point terms can fit “hockey stick like” processes (Models with one knot point).
- These bases can be used in GLMs as well.
- An issue with these approaches is the large number of parameters introduced.
  - Requires some method of so called regularization.

## Harmonics Using Linear Models

- A chord finder from playing the white keys on a piano from c4 (c in 4th octave) to c5.

```
notes4 <- c(261.63, 293.66, 329.63, 349.23, 392.00, 440.00, 493.88, 523.25)
t <- seq(0, 2, by = 0.001)
n <- length(t)
c4 <- sin(2 * pi * notes4[1] * t)
e4 <- sin(2 * pi * notes4[3] * t)
g4 <- sin(2 * pi * notes4[5] * t)
```

```

chord <- c4 + e4 + g4 + rnorm(n, 0, 0.3)
x <- sapply(notes4,
 function(freq) sin(2 * pi * freq * t))
fit <- lm(chord ~ x - 1)

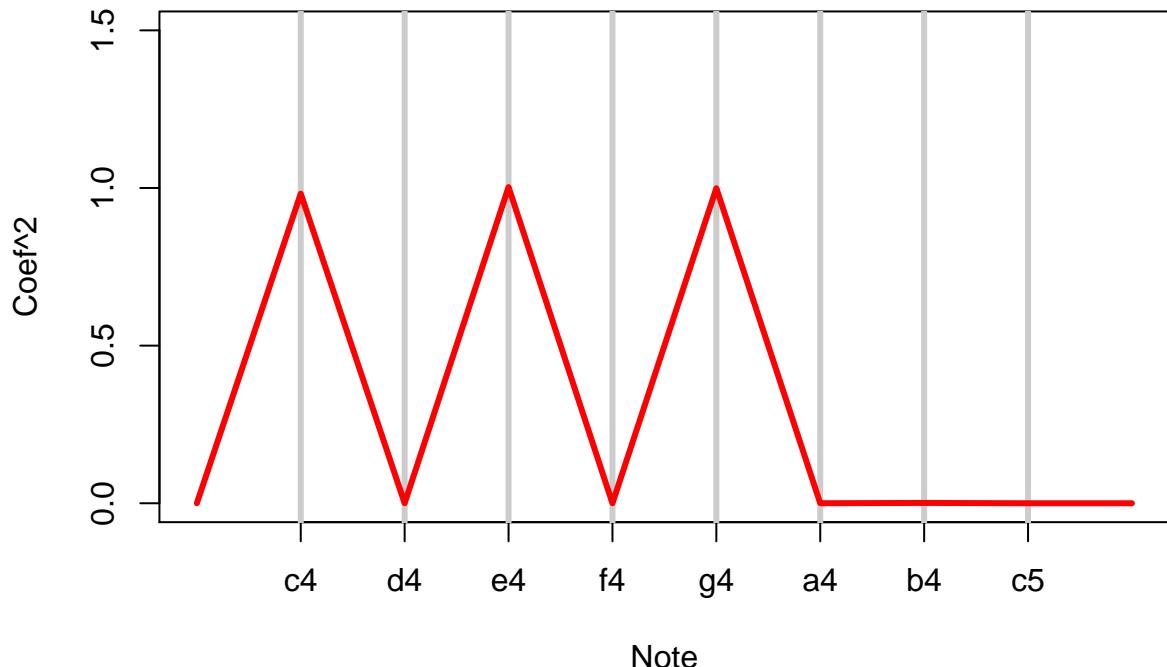
```

- The `savewav` package would let you save .wav files from code
- The model, `chord`, is predicting which note belongs to the chord, even with noise added

```

plot(c(0, 9), c(0, 1.5),
 xlab = "Note", ylab = "Coef^2",
 axes = FALSE, frame = TRUE, type = "n")
axis(2)
axis(1, at = 1:8, labels = c("c4", "d4", "e4", "f4", "g4", "a4", "b4", "c5"))
for (i in 1:8) abline(v = i, lwd = 3, col = grey(0.8))
lines(c(0:9), c(0, coef(fit)^2, 0), type = "l", lwd = 3, col = "#FF0000")

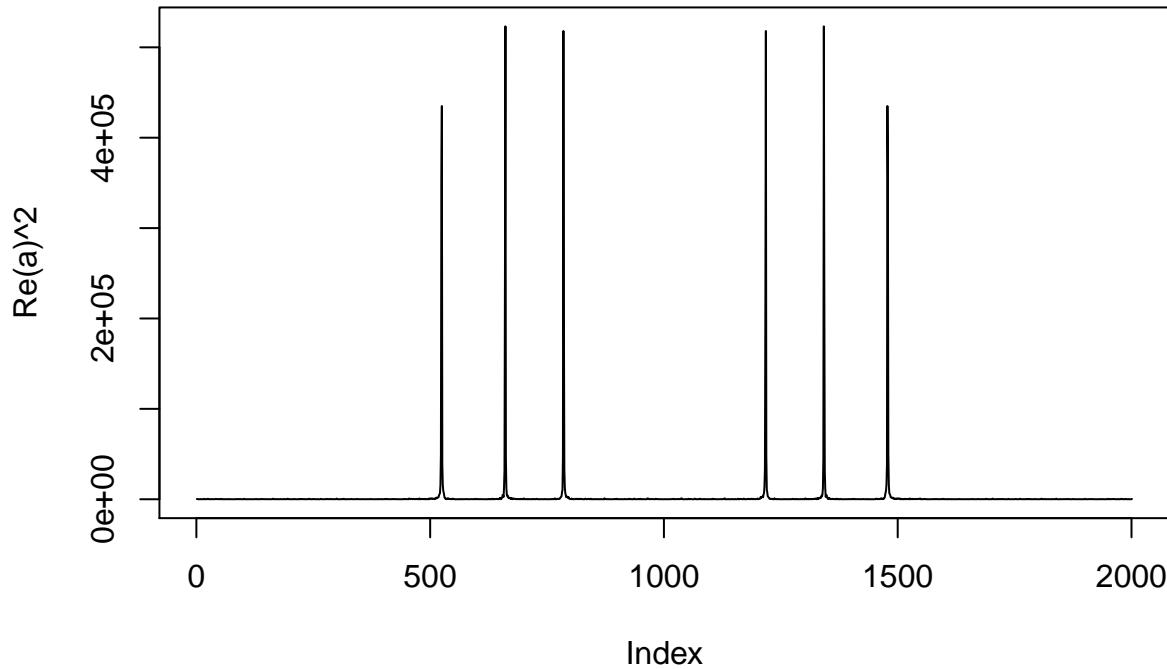
```



```

#(How one would really do it with fourier transforms)
a <- fft(chord)
plot(Re(a)^2, type = "l")

```



- `fft` computes the *discrete fourier transform* of an array using the “Fast Fourier Transform”

#### Quiz 4

1. Consider the space shuttle data, `shuttle` in the `MASS` library. Consider modeling the use of the autolander, `use`, as the outcome. Fit a logistic regression model with autolander use (`auto:1`) versus not (`noauto:0`) as predicted by wind sign, `wind`; `head:1` or `tail:2`. Give the estimated odds ratio for autolander use comparing head winds, labeled as “head” in the variable `headwind` (numerator) to tail winds (denominator).

```
library(MASS)
library(tidyverse)
data(shuttle)
str(shuttle)

'data.frame': 256 obs. of 7 variables:
$ stability: Factor w/ 2 levels "stab","xstab": 2 2 2 2 2 2 2 2 2 2 ...
$ error : Factor w/ 4 levels "LX","MM","SS",...: 1 1 1 1 1 1 1 1 1 1 ...
$ sign : Factor w/ 2 levels "nn","pp": 2 2 2 2 2 2 1 1 1 1 ...
$ wind : Factor w/ 2 levels "head","tail": 1 1 1 2 2 2 1 1 1 2 ...
$ magn : Factor w/ 4 levels "Light","Medium",...: 1 2 4 1 2 4 1 2 4 1 ...
$ vis : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
$ use : Factor w/ 2 levels "auto","noauto": 1 1 1 1 1 1 1 1 1 1 ...
```

```

#Use & wind has levels 1 or 2 so we have to make them binomial
dat <- mutate(shuttle, use = relevel(use, ref = "noauto"))
fit <- glm(as.numeric(use) - 1 ~ 0 + wind,
 family = "binomial", data = dat)
summary(fit)

##
Call:
glm(formula = (as.numeric(use) - 1) ~ 0 + wind, family = "binomial",
data = dat)
##
Deviance Residuals:
Min 1Q Median 3Q Max
-1.300 -1.286 1.060 1.073 1.073
##
Coefficients:
Estimate Std. Error z value Pr(>|z|)
windhead 0.2513 0.1782 1.410 0.158
windtail 0.2831 0.1786 1.586 0.113
##
(Dispersion parameter for binomial family taken to be 1)
##
Null deviance: 354.89 on 256 degrees of freedom
Residual deviance: 350.35 on 254 degrees of freedom
AIC: 354.35
##
Number of Fisher Scoring iterations: 4
exp(coef(fit))

windhead windtail
1.285714 1.327273
exp(coef(fit)[1])/exp(coef(fit)[2])

windhead
0.9686888

2. Consider the previous problem. Give the estimated odds ratio for autolander use comparing head winds (numerator) to tail winds (denominator) adjusting for wind strength from the variable magn.

fit <- glm(as.numeric(use) - 1 ~ 0 + wind + magn,
 family = "binomial", data = dat)
summary(fit)

##
Call:
glm(formula = (as.numeric(use) - 1) ~ 0 + wind + magn, family = "binomial",
data = dat)

```

```


Deviance Residuals:
Min 1Q Median 3Q Max
-1.349 -1.321 1.015 1.040 1.184

Coefficients:
Estimate Std. Error z value Pr(>|z|)
windhead 3.635e-01 2.841e-01 1.280 0.201
windtail 3.955e-01 2.844e-01 1.391 0.164
magnMedium 1.628e-16 3.599e-01 0.000 1.000
magnOut -3.795e-01 3.568e-01 -1.064 0.287
magnStrong -6.441e-02 3.590e-01 -0.179 0.858

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 354.89 on 256 degrees of freedom
Residual deviance: 348.78 on 251 degrees of freedom
AIC: 358.78

Number of Fisher Scoring iterations: 4
exp(coef(fit))

```

```

windhead windtail magnMedium magnOut magnStrong
1 1.4383682 1.4851533 1.0000000 0.6841941 0.9376181
exp(coef(fit)[1])/exp(coef(fit)[2])

```

```

windhead
0.9684981

```

3. If you fit a logistic regression model to a binary variable, for example use of the autolander, then fit a logistic regression model for one minus the outcome (not using the autolander) what happens to the coefficients?
  - The outcome is binary (1 for success or 0 for failure), doing  $1 - \text{outcome}$  will make it a 0 for a success or a 1 for a failure, which would effectively just change the sign of the coefficients.
4. Consider the insect spray data `InsectSprays`. Fit a Poisson model using spray as a factor level. Report the estimated relative rate comparing spray A (numerator) to spray B (denominator).

```

data("InsectSprays")
fit <- glm(count ~ 0+ spray, family = "poisson", data = InsectSprays)
summary(fit)

```

```


Call:
glm(formula = count ~ 0 + spray, family = "poisson", data = InsectSprays)

Deviance Residuals:

```

```

Min 1Q Median 3Q Max
-2.3852 -0.8876 -0.1482 0.6063 2.6922
##
Coefficients:
Estimate Std. Error z value Pr(>|z|)
sprayA 2.67415 0.07581 35.274 < 2e-16 ***
sprayB 2.73003 0.07372 37.032 < 2e-16 ***
sprayC 0.73397 0.20000 3.670 0.000243 ***
sprayD 1.59263 0.13019 12.233 < 2e-16 ***
sprayE 1.25276 0.15430 8.119 4.71e-16 ***
sprayF 2.81341 0.07071 39.788 < 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
(Dispersion parameter for poisson family taken to be 1)
##
Null deviance: 2264.808 on 72 degrees of freedom
Residual deviance: 98.329 on 66 degrees of freedom
AIC: 376.59
##
Number of Fisher Scoring iterations: 5
co <- coef(fit)
exp(co[1])/exp(co[2])

```

```

sprayA
0.9456522

```

5. Consider a Poisson glm with an offset, t. So, for example, a model of the form `glm(count ~ x + offset(t), family = poisson)` where x is a factor variable comparing a treatment (1) to a control (0) and t is the natural log of a monitoring time. What is impact of the coefficient for x if we fit the model `glm(count ~ x + offset(t2), family = poisson)` where `t2 <- log(10) + t`? In other words, what happens to the coefficients if we change the units of the offset variable. (Note, adding  $\log(10)$  on the log scale is multiplying by 10 on the original scale.)

- The `offset` function is just recentering the origin so this wouldn't affect the coefficient
6. Given:

```

x <- -5:5
y <- c(5.12, 3.93, 2.67, 1.87, 0.52, 0.08, 0.93, 2.05, 2.54, 3.87, 4.97)

```

Using a knot point at 0, fit a linear model that looks like a hockey stick with two lines meeting at  $x=0$ . Include an intercept term, x and the knot point term. What is the estimated slope of the line after 0?

```

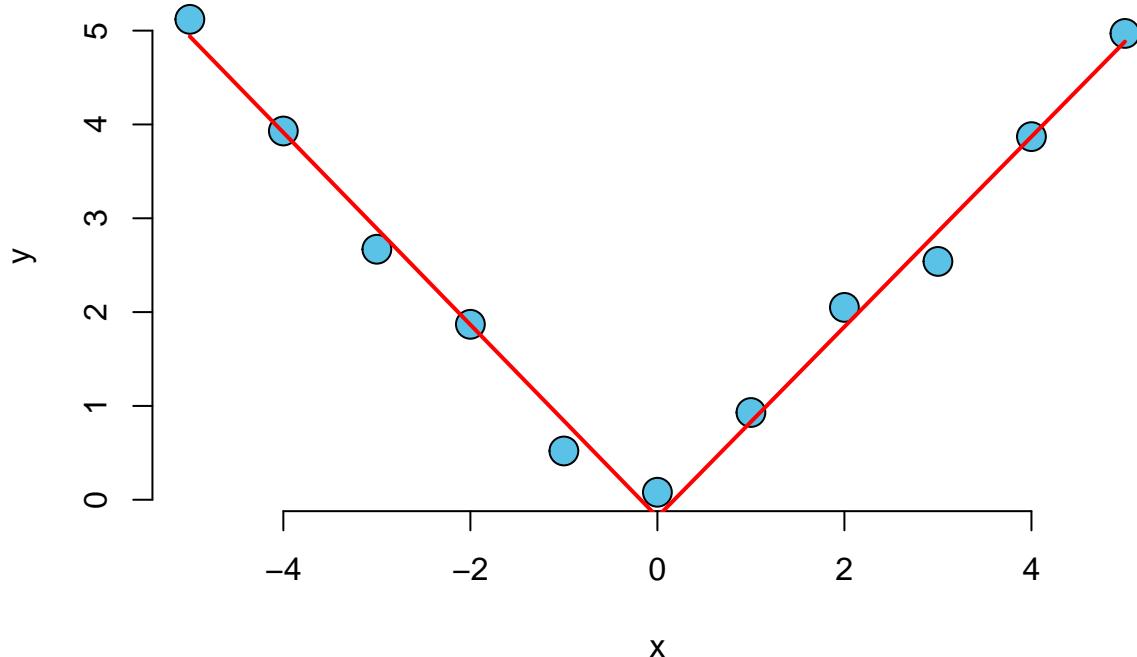
knot <- 0
splineTerms <- sapply(knot,
 function(knot) {(x > knot) * (x - knot)})
xmat <- cbind(1, x, splineTerms)

```

```

fit <- lm(y ~ xmat - 1)
yhat <- predict(fit) #int. is included in x
plot(x, y, frame = FALSE, pch = 21, bg = "#5BC2E7", cex = 2)
lines(x, yhat, col = "#FF0000", lwd = 2)

```



```

summary(fit)

##
Call:
lm(formula = y ~ xmat - 1)
##
Residuals:
Min 1Q Median 3Q Max
-0.32158 -0.10979 0.01595 0.14065 0.26258
##
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 2.03723 0.08575 23.759 1.05e-08 ***
xmat -0.18258 0.13558 -1.347 0.215
xmatx -1.02416 0.04805 -21.313 2.47e-08 ***
xmat 2.03723 0.08575 23.759 1.05e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##

```

```
Residual standard error: 0.2276 on 8 degrees of freedom
Multiple R-squared: 0.996, Adjusted R-squared: 0.9945
F-statistic: 665 on 3 and 8 DF, p-value: 6.253e-10
sum(fit$coefficients[2:3])
[1] 1.013067
```

Reminder to commit (S4) delete this line *AFTER* committing

## Course Project

Reminder to commit (P1) delete this line *BEFORE* committing